

## PRACTICAL N0 – 1

**Aim:** Arima used for acf and pacf

**Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error

#Load time Series data
#load built-in AirPassengers dataset
url="https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv"
df=pd.read_csv(url, parse_dates=['Month'],index_col='Month')
df_columns=['Passengers']
df.head()
```

**Output:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error

[ ] #Load time Series data
#load built-in AirPassengers dataset
url="https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv"
df=pd.read_csv(url, parse_dates=['Month'],index_col='Month')
df_columns=['Passengers']
df.head()



| Month      | Passengers |
|------------|------------|
| 1949-01-01 | 112        |
| 1949-02-01 | 118        |
| 1949-03-01 | 132        |
| 1949-04-01 | 129        |
| 1949-05-01 | 121        |


```

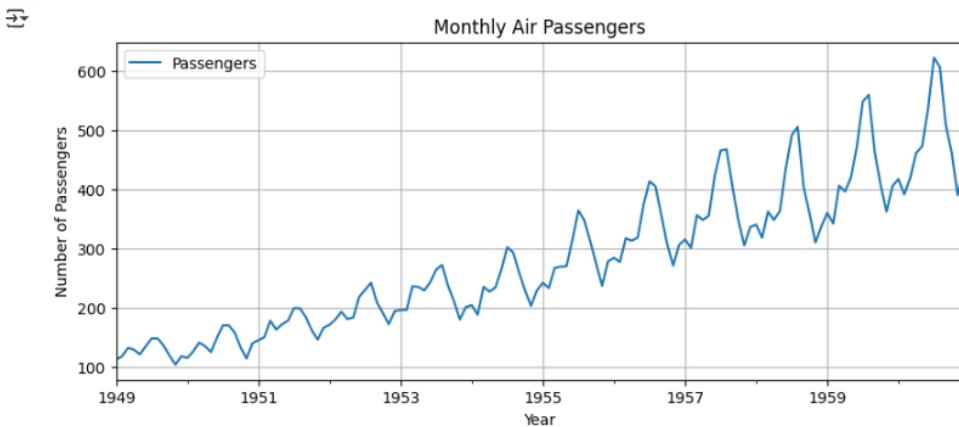
**Code:**

```
#plot the time series
df.plot(title='Monthly Air Passengers',figsize=(10,4))
plt.ylabel("Number of Passengers")
plt.xlabel("Year")
```

```
plt.grid()  
plt.show()
```

## Output:

```
▶ #plot the time series  
df.plot(title='Monthly Air Passengers', figsize=(10,4))  
plt.ylabel("Number of Passengers")  
plt.xlabel("Year")  
plt.grid()  
plt.show()
```



## Code:

```
#test for stationary(adf)  
def test_stationarity(series):  
    result = adfuller(series)  
    print(f'ADF statistics: {result[0]:.4f}')  
    print(f'p-value: {result[1]:.4f}')  
    if result[1] <= 0.05:  
        print('Stationary')  
    else:  
        print('Not Stationary')  
test_stationarity(df['Passengers'])
```

## Output:

```
[ ] #test for stationary(adf)  
def test_stationarity(series):  
    result = adfuller(series)  
    print(f'ADF statistics: {result[0]:.4f}')  
    print(f'p-value: {result[1]:.4f}')  
    if result[1] <= 0.05:  
        print('Stationary')  
    else:  
        print('Not Stationary')  
test_stationarity(df['Passengers'])
```

```
ADF statistics: 0.8154  
p-value: 0.9919  
Not Stationary
```

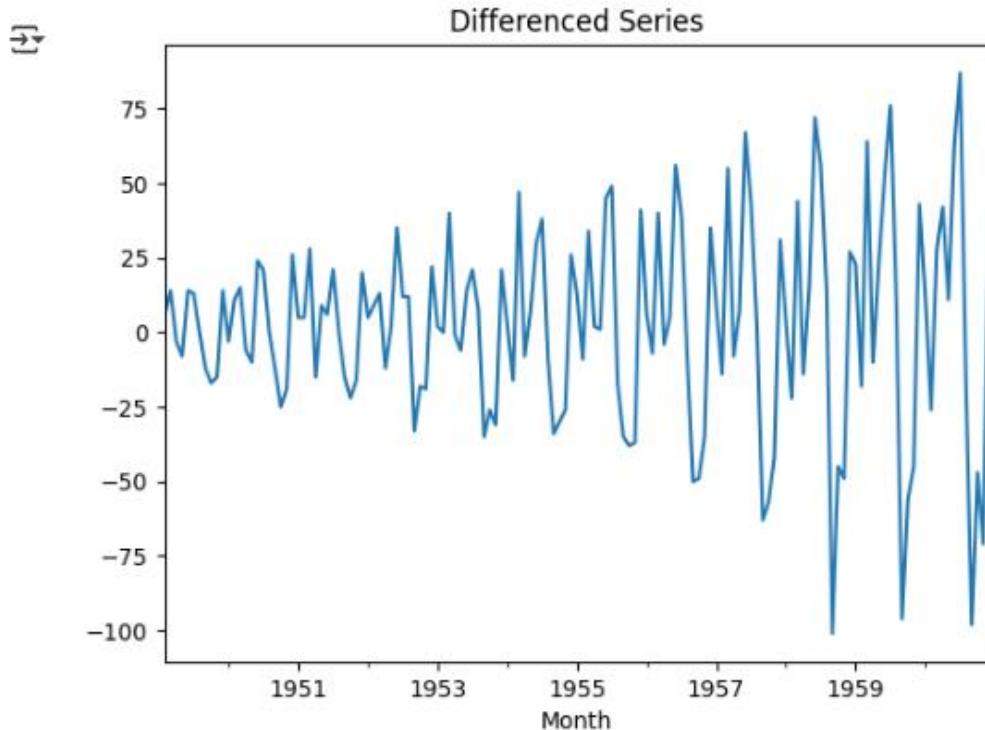
## Code:

```
#make stationary (Differencing)
```

```
df_diff = df['Passengers'].diff().dropna()
df_diff.plot(title='Differenced Series')
plt.show()

test_stationarity(df_diff)
```

### Output:



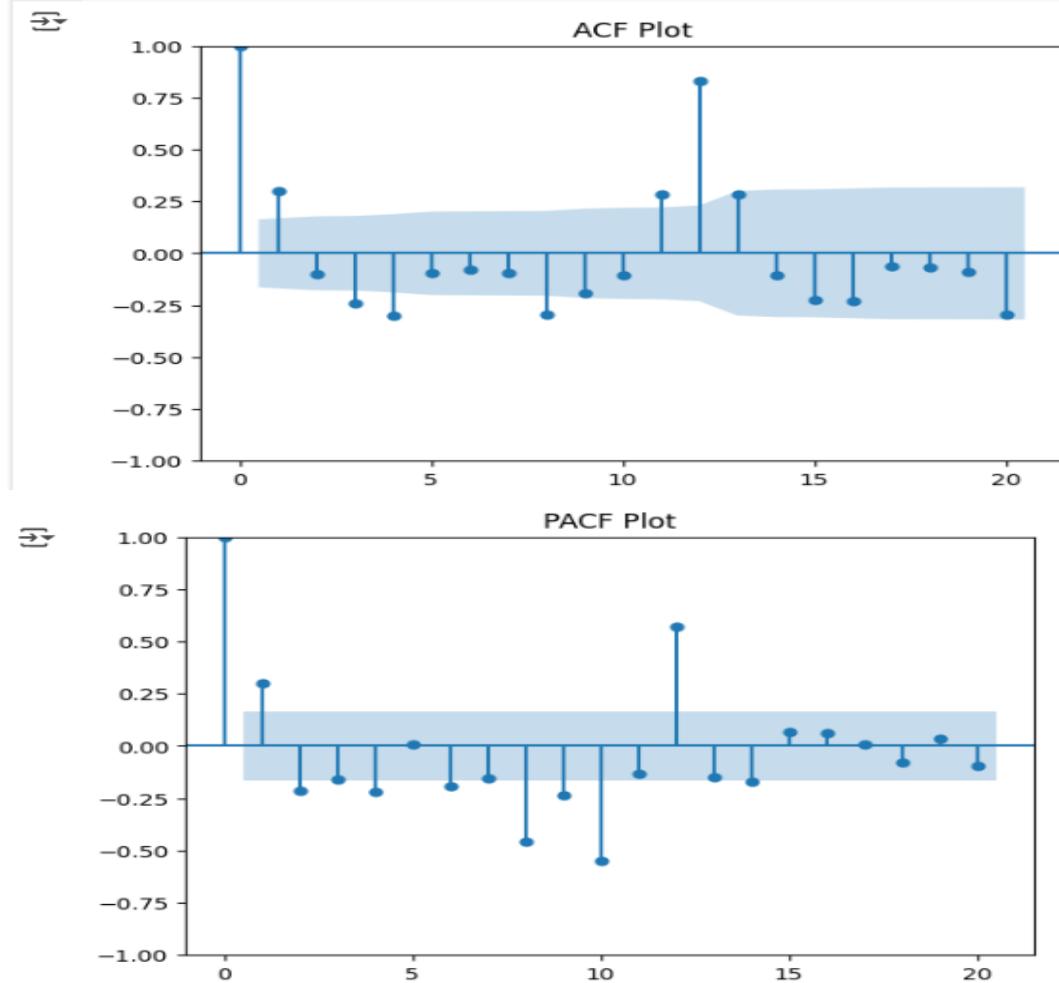
```
ADF statistics: -2.8293
p-value: 0.0542
Not Stationary
```

### Code:

```
#plot Acf and pacf
plot_acf(df_diff, lags=20)
plt.title("ACF Plot")
plt.show()

plot_pacf(df_diff, lags=20)
plt.title("PACF Plot")
plt.show()
```

### Output:



## Code:

```
#Arima
model=ARIMA(df['Passengers'],order=(2,1,2))
model_fit=model.fit()
model_fit.summary()
```

## Output:

```

In [1]: time series prac 1.ipynb
In [2]: %%time
        model_fit = sm.tsa.SARIMAX(df['Passengers'], order=(2, 1, 2),
                                    seasonal_order=(1, 1, 1, 12),
                                    enforce_stationarity=False,
                                    enforce_invertibility=False)
        model_fit.fit()
        model_fit.summary()

Out[2]:
<IPython.core.display.HTML>

```

```

model_fit.summary()

Dep. Variable: Passengers   No. Observations: 144
Model: ARIMA(2, 1, 2)   Log Likelihood: -671.673
Date: Sat, 12 Jul 2025   AIC: 1353.347
Time: 11:26:57   BIC: 1368.161
Sample: 01-01-1949   HQIC: 1359.366
                                                - 12-01-1960

Covariance Type: opg
            coef  std err      z  P>|z|  [0.025  0.975]
arL1    1.6850  0.020  83.060  0.000  1.645  1.725
arL2   -0.9548  0.017  -55.420  0.000  -0.989  -0.921
maL1   -1.8432  0.124  -14.814  0.000  -2.087  -1.599
maL2    0.9552  0.135   7.382  0.000  0.731  1.260
sigma2 665.9600 114.029 5.841  0.000 442.467 889.453
Ljung-Box (L1) (Q): 0.30 Jarque-Bera (JB): 1.84
Prob(Q): 0.59 Prob(JB): 0.40
Heteroskedasticity (H): 7.38 Skew: 0.27
Prob(H) (two-sided): 0.08 Kurtosis: 3.14

```

## Code:

```

#forecast and plot
forecast = model_fit.forecast(steps=24)
forecast.index = pd.date_range(start=df.index[-1] +
pd.DateOffset(months=1), periods=24, freq='MS')

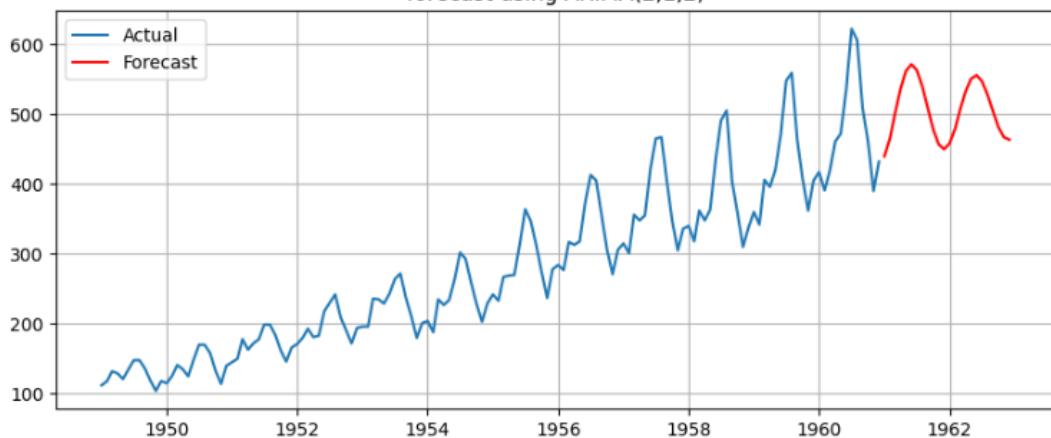
plt.figure(figsize=(10, 4))
plt.plot(df['Passengers'], label='Actual')
plt.plot(forecast, label='Forecast', color='red')
plt.title("forecast using ARIMA(2,1,2)")
plt.legend()
plt.grid()
plt.show()

```

## Output:



forecast using ARIMA(2,1,2)



## Code:

```
#evaluate forecast(train/test)

#train
train=df['Passengers'][:-24]
test=df['Passengers'][-24:]

model=ARIMA(train,order=(2,1,2))
model_fit=model.fit()
preds = model_fit.forecast(steps=24)

#evaluate
mse=mean_squared_error(test,preds)
print(f'Mean squared error: {mse:.2f}')
```

## Output:

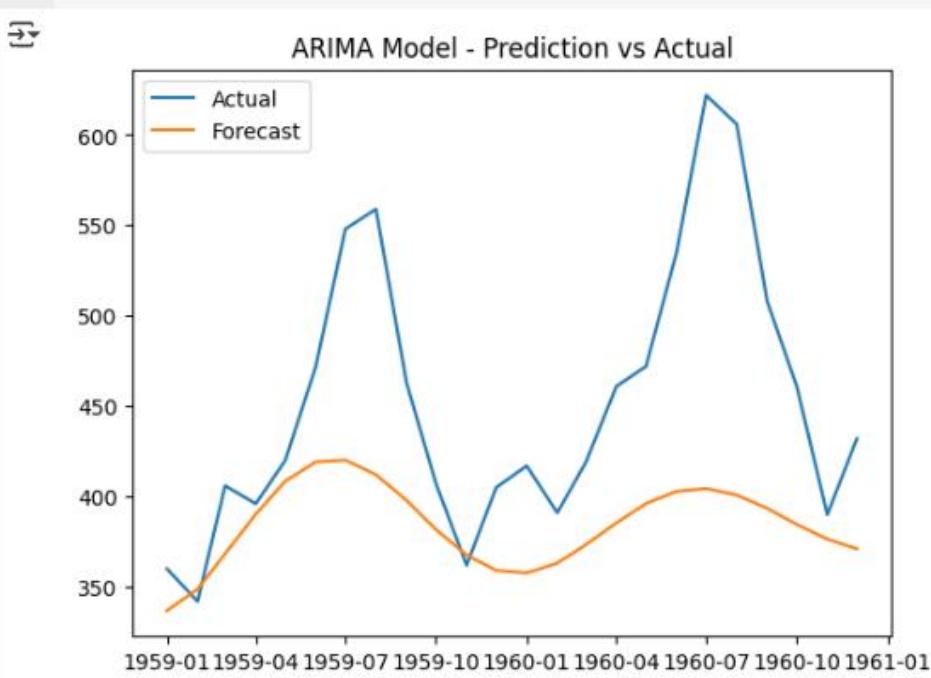
```
#evaluate
mse=mean_squared_error(test,preds)
print(f'Mean squared error: {mse:.2f}')

/usr/local/lib/python3.11/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
/usr/local/lib/python3.11/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
/usr/local/lib/python3.11/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
Mean squared error: 825.46
```

## Code:

```
plt.plot(test.index,test.values,label='Actual')
plt.plot(test.index,preds,label='Forecast')
plt.title("ARIMA Model - Prediction vs Actual")
plt.legend()
plt.show()
```

## Output:

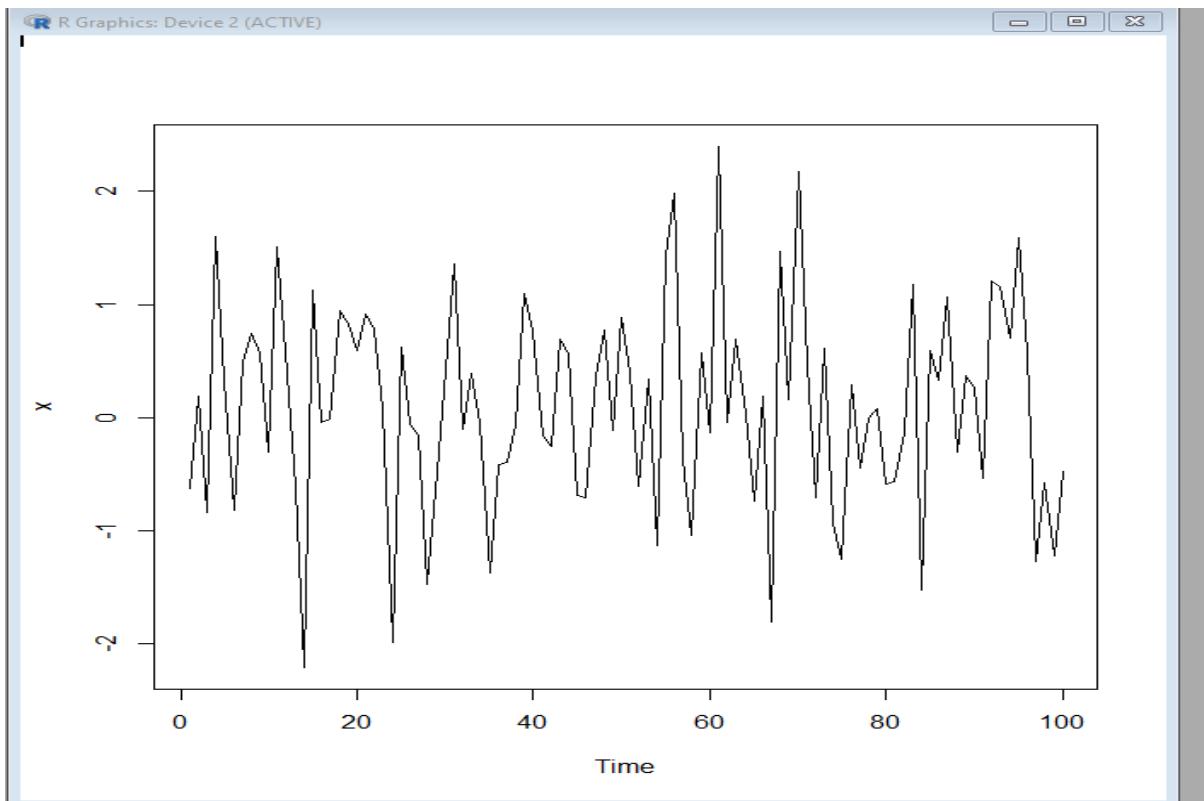


#Used in r studio

**Code:**

```
set.seed(1)  
x<-rnorm(100,mean=0,sd=1)  
plot(x,type='l',xlab="Time")
```

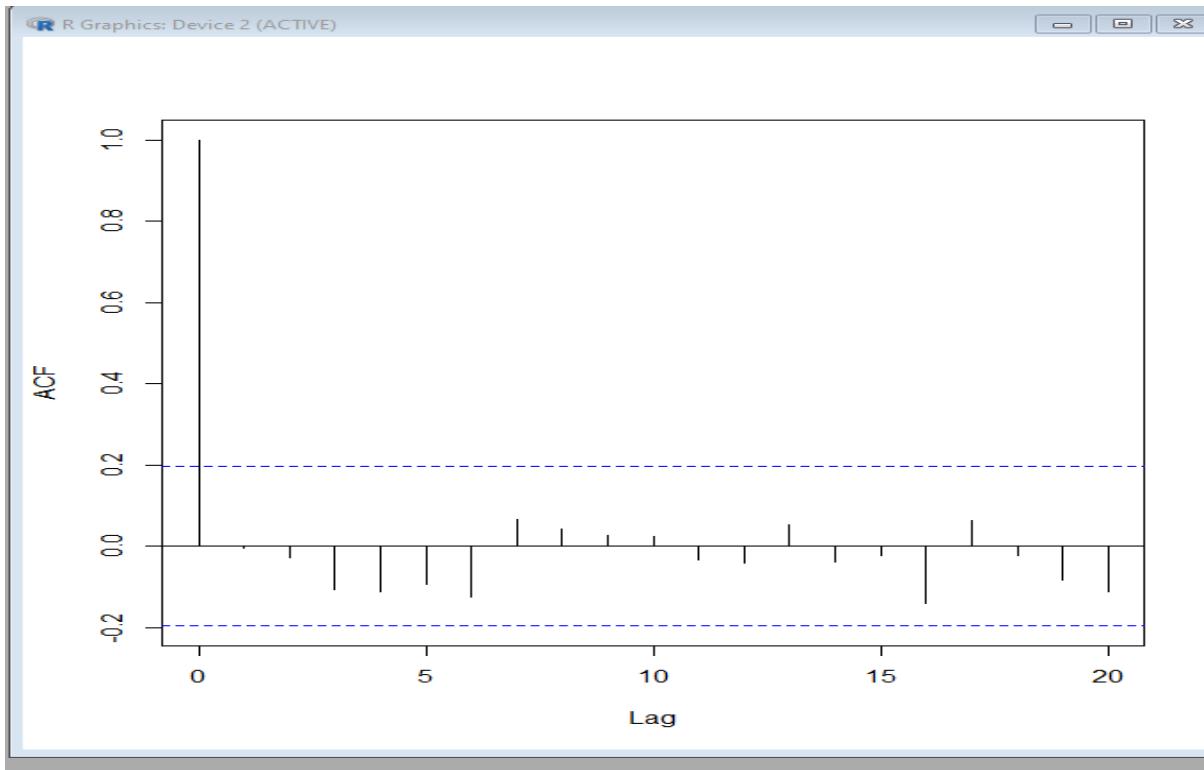
**Output:**



**Code:**

```
acf(x,main="")
```

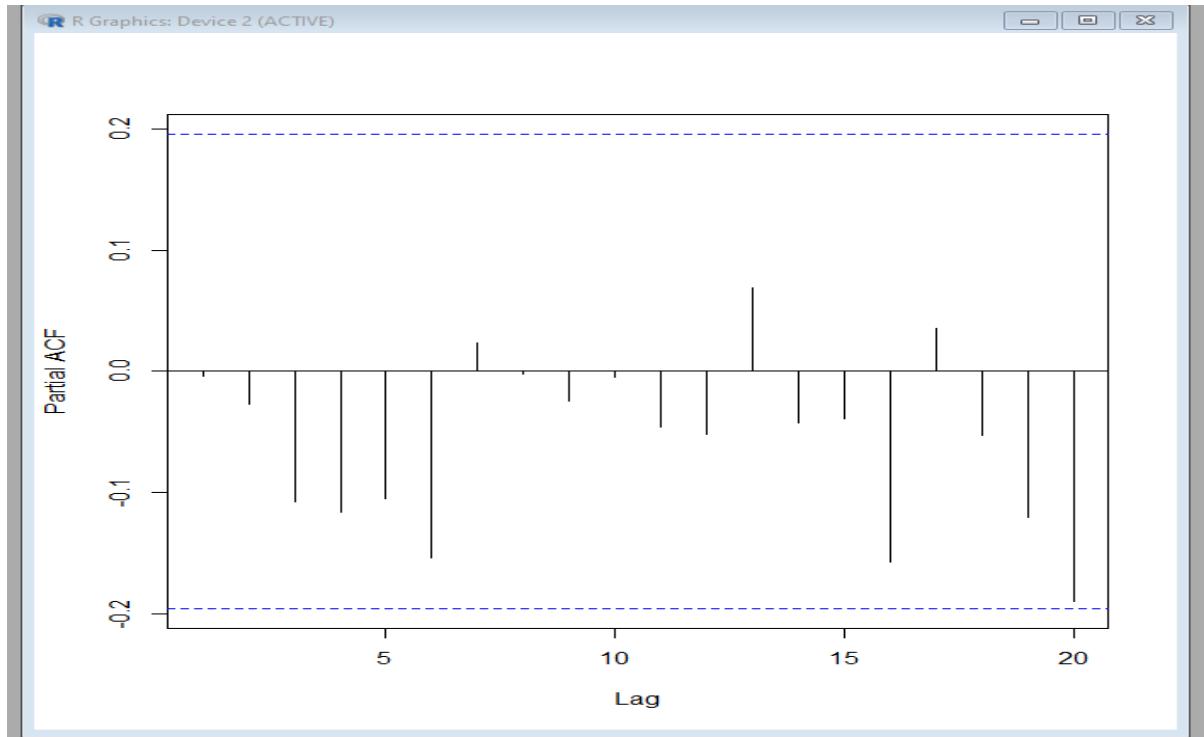
**Output:**



**Code:**

```
pacf(x,main="")
```

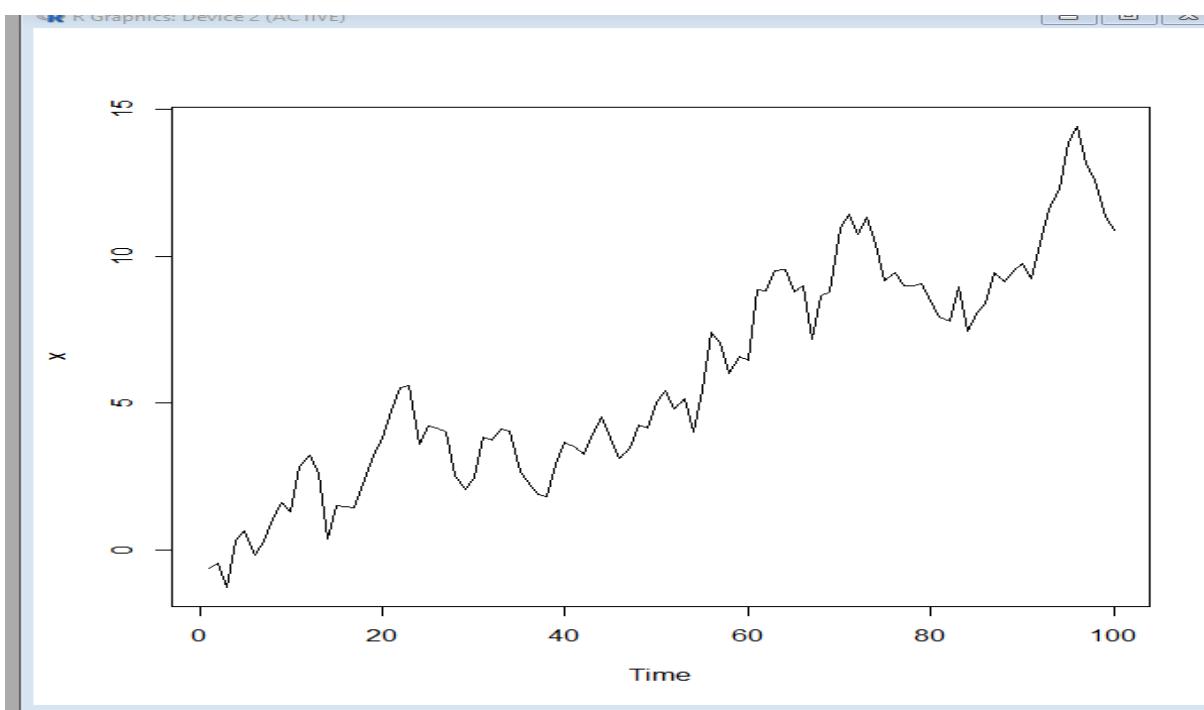
**Output:**



**Code:**

```
set.seed(1)  
x<-w<-rnorm(100,mean=0,sd=1)  
for(t in 2:100)x[t]<-x[t-1]+w[t]  
plot(x,type='l',xlab="Time")
```

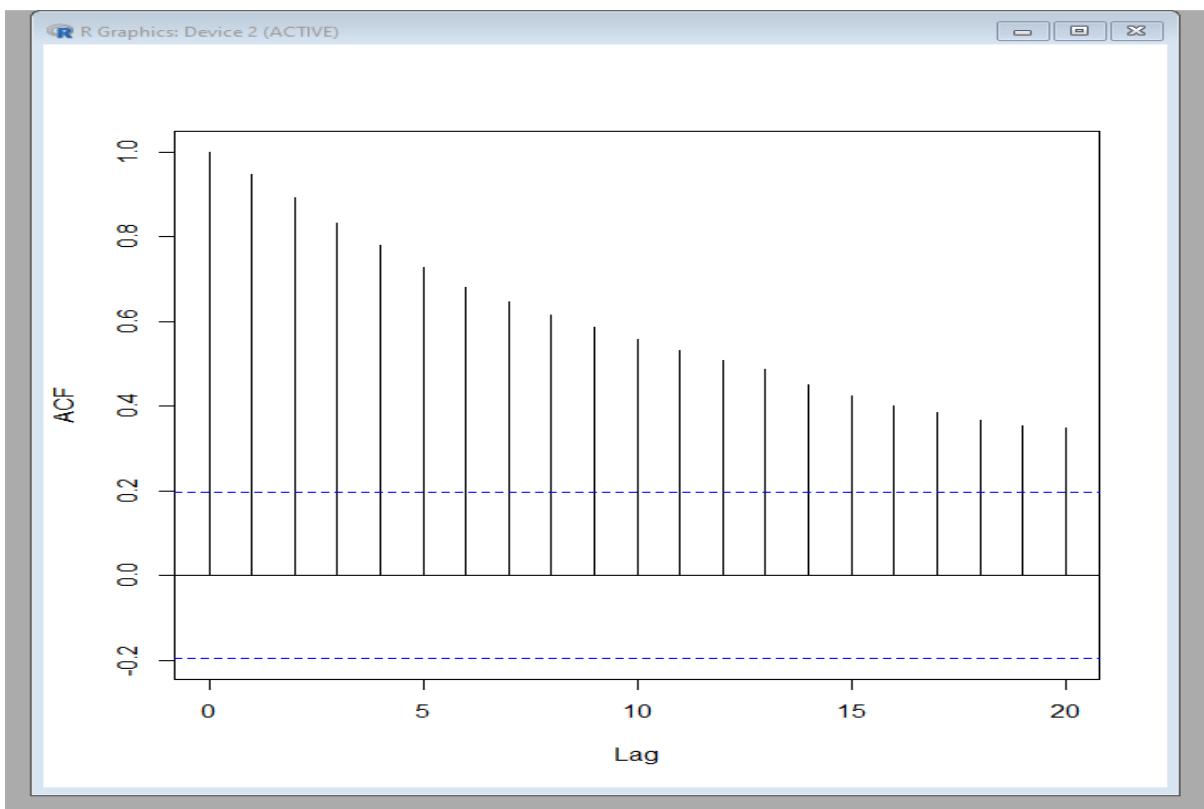
**Output:**



**Code:**

```
acf(x,main="")
```

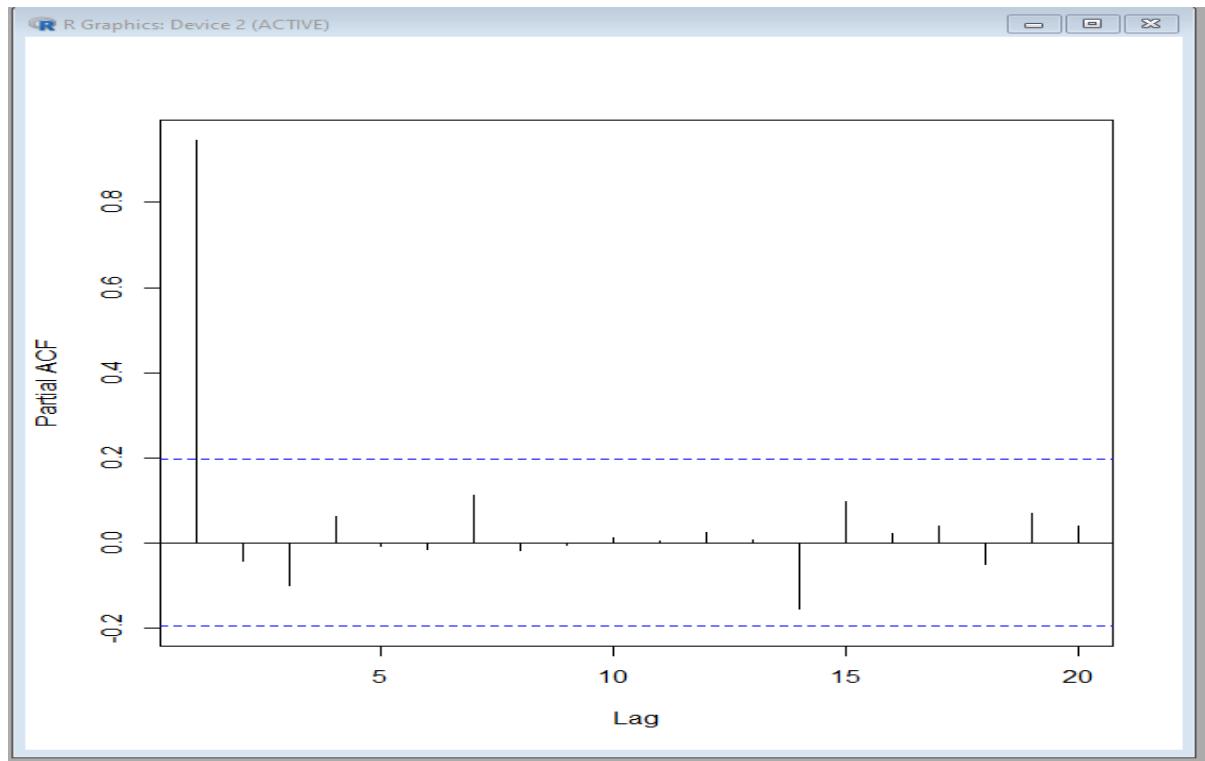
**Output:**



**Code:**

```
pacf(x,main="")
```

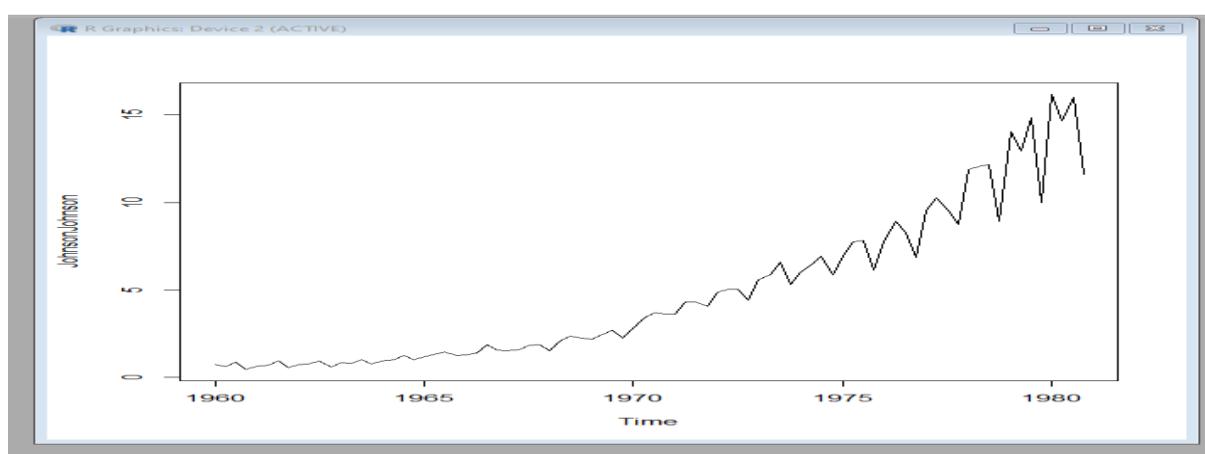
**Output:**



**Code:**

```
library(datasets)  
plot(JohnsonJohnson)
```

**Output:**



## PRACTICAL NO – 2

**Aim:** Linear regression

**Code:**

```
import numpy as nm
import matplotlib.pyplot as mtp
import pandas as pd

data_set = pd.read_csv('Salary_Data.csv')
print(data_set)
```

**Output:**

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
..	..	..

**Code:**

```
x=data_set.iloc[:, :-1].values
y=data_set.iloc[:, 1].values
print(y)

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=1/3,random_state=0)
from sklearn.linear_model import LinearRegression
regressor=LinearRegression()
regressor.fit(x_train,y_train)
```

**Output:**

```
x=data_set.iloc[:, :-1].values
y=data_set.iloc[:, 1].values
print(y)

[ 39343.  46205.  37731.  43525.  39891.  56642.  60150.  54445.  64445.
 57189.  63218.  55794.  56957.  57081.  61111.  67938.  66029.  83088.
 81363.  93940.  91738.  98273.  101302.  113812.  109431.  105582.  116969.
 112635.  122391.  121872.]
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=1/3,random_state=0)
```

```
[ ] from sklearn.linear_model import LinearRegression
regressor=LinearRegression()
regressor.fit(x_train,y_train)
```

```
LinearRegression()
```

## Code:

```
_pred=regressor.predict(x_test)
x_pred=regressor.predict(x_train)
print(y_pred)
print(x_pred)
```

## Output:

```
y_pred=regressor.predict(x_test)
x_pred=regressor.predict(x_train)
print(y_pred)
print(x_pred)

[ 40835.10590871 123079.39940819  65134.55626083  63265.36777221
 115602.64545369 108125.8914992  116537.23969801  64199.96201652
 76349.68719258 100649.1375447 ]
[ 53919.42532909 74480.49870396  56723.20806202  68872.93323808
 103452.92027763 90368.60085726  38965.91742009 124948.58789682
 54854.0195734 47377.2656189   81957.25265845  82891.84690277
 61396.17928358 56723.20806202 110929.67423213  45508.07713028
 37096.72893147 93172.3835902   72611.31021533  64199.96201652]
```

## Code:

```
mtp.scatter(x_train,y_train,color="green")
mtp.plot(x_train,x_pred,color="red")
mtp.title("Salary vs Experience(Training DataSet)")
mtp.xlabel("Years of Experience")
mtp.ylabel("Salary(In Rupees)")
mtp.show()
```

## Output:



## PRACTICAL N0 – 3

**Aim:** Exponential Smoothing using fit holt winter additive

**Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
#load or Define data
#load and prepare data
sales=[210,220,230,300,310,305,280,290,295,310,320,330,220,225,235,310,
315,310,285,295,300,315,325,335]
dates=pd.date_range(start='2022-01-01',periods=len(sales),freq='MS')
df = pd.Series(sales,index=dates)
df.plot(title='Monthly Retail Sales($)',figsize=(10,4))
plt.ylabel("Sales")
plt.grid()
plt.show()
```

**Output:**



**Code:**

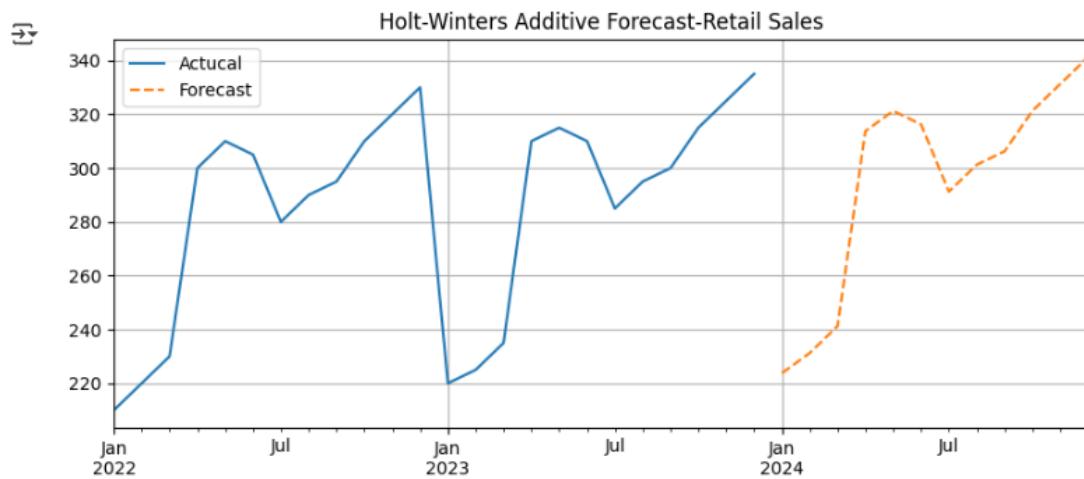
```
#Create models
#fit holt-winters Additive
model =
ExponentialSmoothing(df,trend='add',seasonal='add',seasonal_periods=12)
fit=model.fit()
#forecast next 12 month
forecast=fit.forecast(12)
#plot actucal & forecasted values
#plot actual + forecast
plt.figure(figsize=(10,4))
df.plot(label='Actucal')
forecast.plot(label='Forecast',linestyle='--')
plt.title("Holt-Winters Additive Forecast-Retail Sales")
```

```

plt.legend()
plt.grid()
plt.show()

```

## Output:



#using different method

## Code:

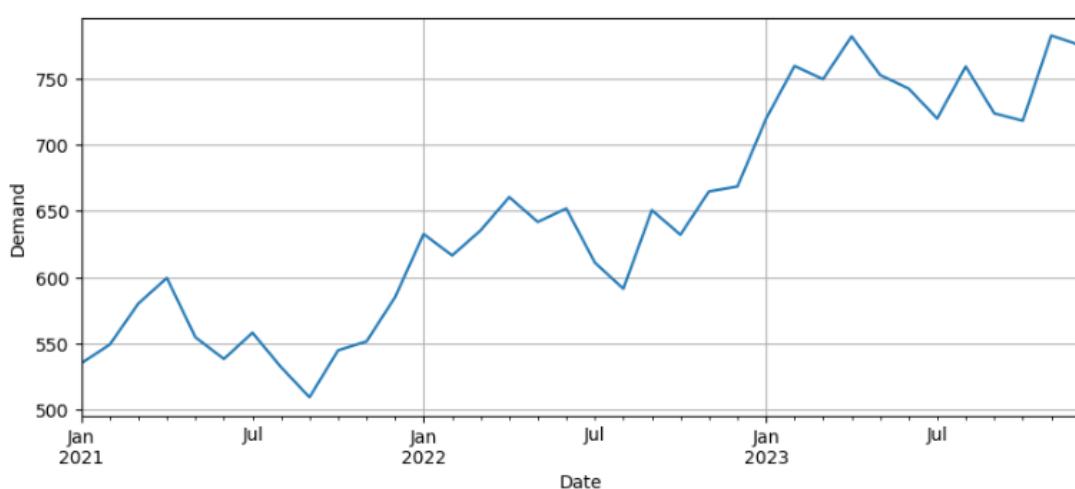
```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
#Simulated Electronics demand data
#monthly demand over 3 year(36month)
np.random.seed(42)
months=pd.date_range(start='2021-01-01',periods=36,freq='MS')
#Trend+seasonality+noise
demand=
    500+#base demand
    np.linspace(0,300,36)+ #upward trend
    50*np.sin(2*np.pi*months.month/12)+ #seasonality(1-yearcycle)
    np.random.normal(0,20,36) #random noise
)
df=pd.Series(demand,index=months)
df.name='Monthly Electronics Demand'
plt.figure(figsize=(10,4))
df.plot()
df=pd.Series(demand,index=months)
df.name = ('Monthly Electronics Demand(units)')
plt.xlabel('Date')
plt.ylabel('Demand')
plt.grid()
plt.show()

```

## Output:

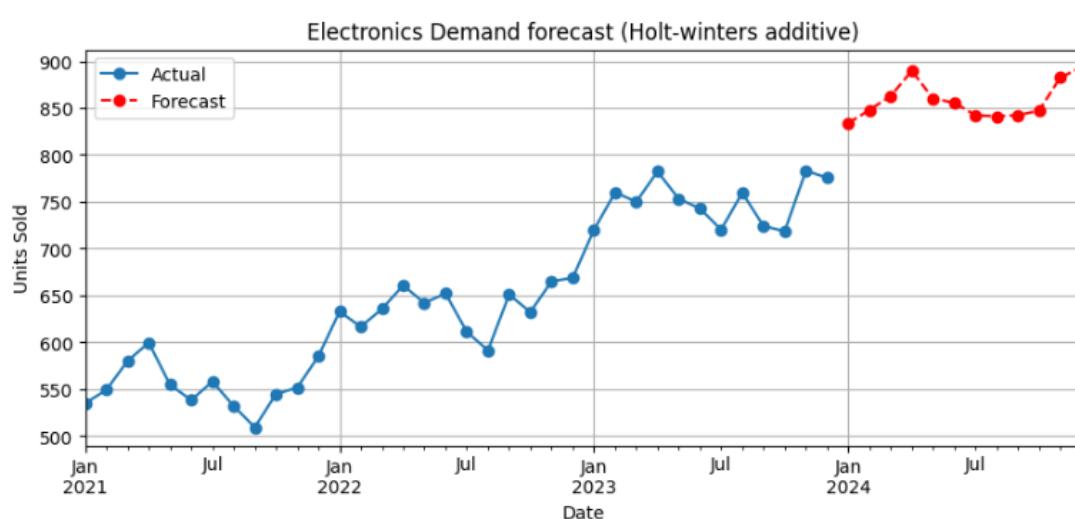
```
plt.show()
```



## Code:

```
#apply holt-winter forecasting (additive)
model=ExponentialSmoothing(df,trend='add',seasonal='add',seasonal_periods=12)
fit=model.fit()
#forecast for 12 months
forecast=fit.forecast(12)
#plotting actucal and forecasting demand
plt.figure(figsize=(10,4))
df.plot(label='Actual',marker='o')
forecast.plot(label='Forecast',marker='o',linestyle='--',color='red')
plt.title('Electronics Demand forecast (Holt-winters additive)')
plt.xlabel('Date')
plt.ylabel("Units Sold")
plt.legend()
plt.grid()
plt.show()
```

## Output:

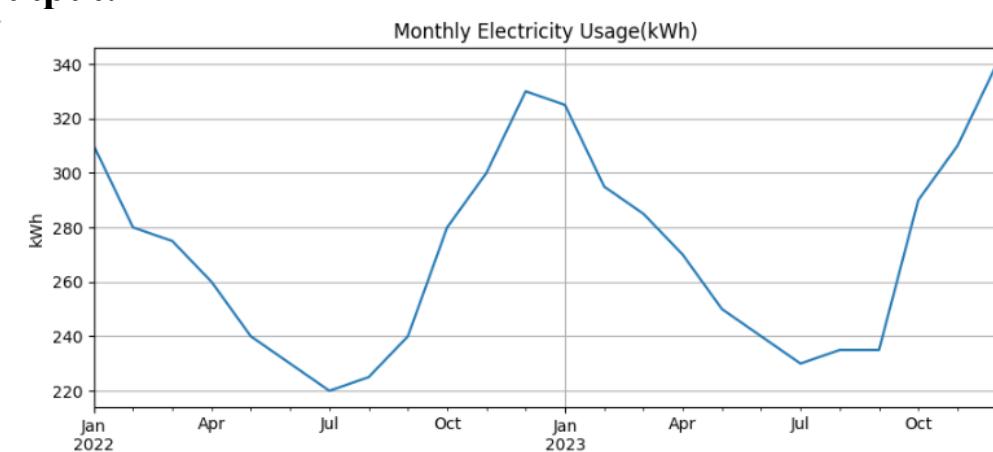


## #using multiplicative

### Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
#energy consumption(multiplicative)
#predict monthly electricity usage
#synthetic but realistic kwh values for monthly electricity
energy=[310,280,275,260,240,230,220,225,240,280,300,330,325,295,285,270
,250,240,230,235,235,290,310,340]
dates=pd.date_range(start='2022-01-01',periods=len(energy),freq='MS')
df2=pd.Series(energy,index=dates)
df2.plot(title='Monthly Electricity Usage(kWh)',figsize=(10,4))
plt.ylabel('kWh')
plt.grid()
plt.show()
```

### Output:

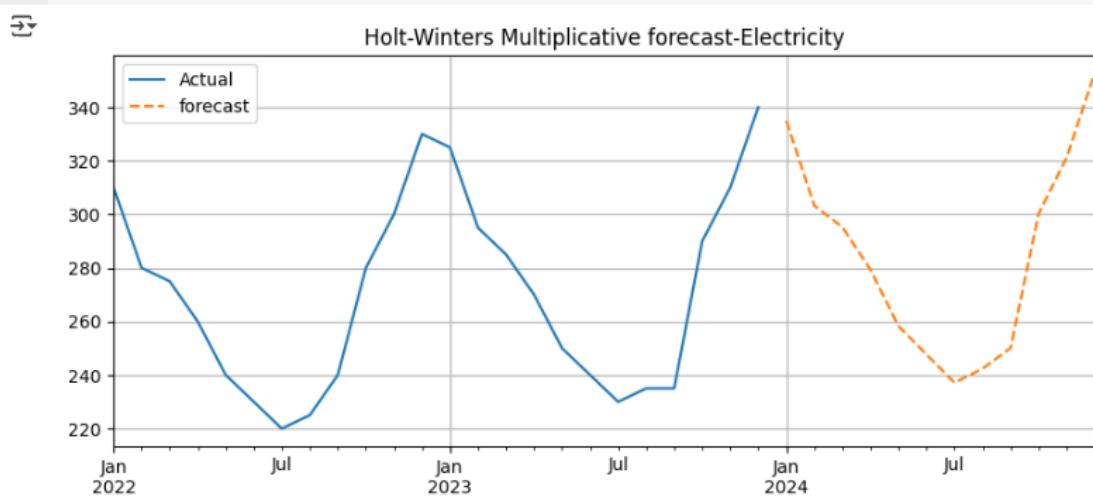


### Code:

```
#Holt Winters Multiplicative
model2=ExponentialSmoothing(df2,trend='add',seasonal='mul',seasonal_periods=12)
fit2=model2.fit()
#forecast
forecast2=fit2.forecast(12)

#plot electricity usage
plt.figure(figsize=(10,4))
df2.plot(label='Actual')
forecast2.plot(label='forecast',linestyle='--')
plt.title("Holt-Winters Multiplicative forecast-Electricity")
plt.grid()
plt.legend()
plt.show()
```

## Output:



#using different dataset

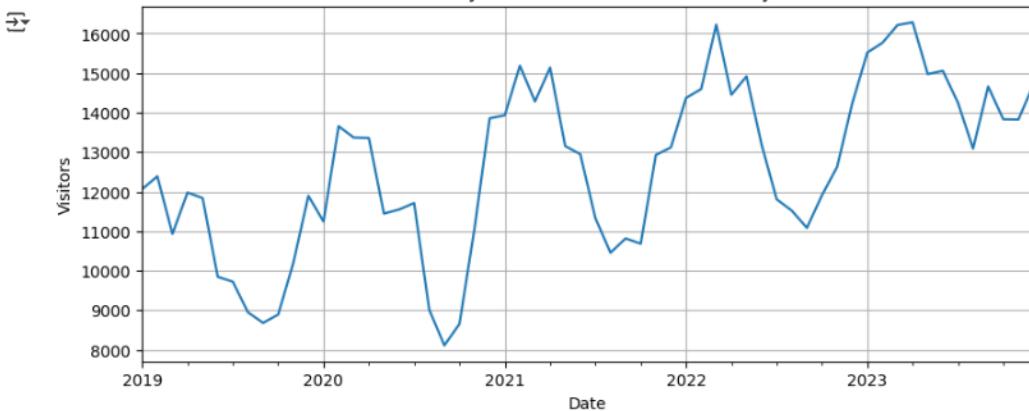
Code:

```
#prac 7
#holt-winter on kashmir Tourism data
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import ExponentialSmoothing
#monthly tourist data
np.random.seed(10)
months=pd.date_range(start='2019-01-01',periods=60,freq='MS')
#base+trend+seasonal variation(e.g high in may-july and dec-jan)
base=10000
trend=np.linspace(0,5000,60)
seasonal_pattern=np.sin(2*np.pi*months.month/12)*2000 #seasonal effect

noise=np.random.normal(0,800,60) #random variation
visitors = base+trend+seasonal_pattern+noise
df = pd.Series(visitors, index=months)
df.name='Monthly Visitors(Kashmir)'
#plot actual data
plt.figure(figsize=(10,4))
df.plot()
plt.title("Monthly Tourist Visitors-Kashmir Valley")
plt.xlabel("Date")
plt.ylabel('Visitors')
plt.grid()
plt.show()
```

## Output:

Monthly Tourist Visitors-Kashmir Valley



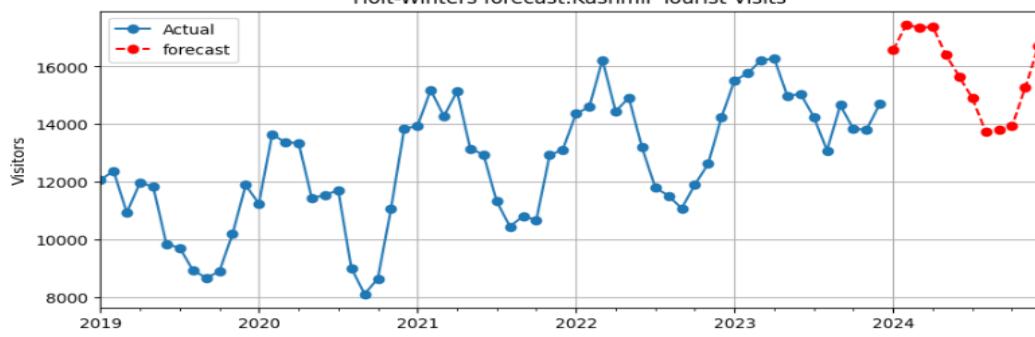
```
#apply holt-Winters (additive seasonality)
```

## Code:

```
model=ExponentialSmoothing(df,trend='add',seasonal='add',seasonal_perio  
ds=12)  
#12 months per seasonal cycle  
fit=model.fit()  
#forecast next 12 months tourist  
forecast=fit.forecast(12)  
#plot actual vs forecast data  
plt.figure(figsize=(10,4))  
df.plot(label='Actual',marker='o')  
forecast.plot(label='forecast',marker='o',linestyle='--',color='red')  
plt.title("Holt-Winters forecast:Kashmir Tourist Visits")  
plt.xlabel("Date")  
plt.ylabel("Visitors")  
plt.grid()  
plt.legend()  
plt.show()  
print("Forecasted Tourist Visits(Next 12 Months):")  
print(forecast.round(0).astype(int))
```

## Output:

Holt-Winters forecast:Kashmir Tourist Visits



## **PRACTICAL N0 – 4**

**Aim:** Simple Exponential Smoothing

**Code:**

```
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import SimpleExpSmoothing,Holt
#create or load dataset
#monthly sales data(in thousand)
sales_data=[200,210,215,220,230,225,235,240,250,255,260,265]

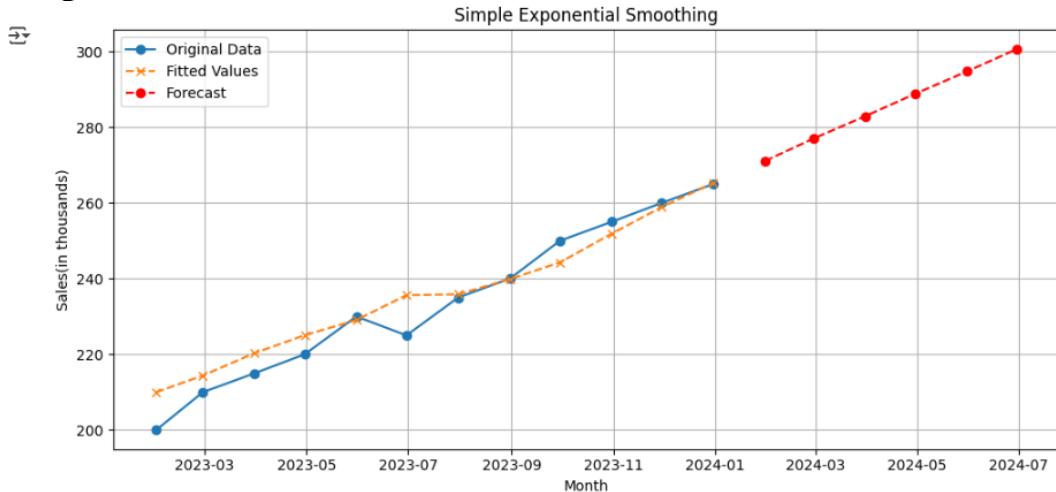
#create a time index
dates=pd.date_range(start='2023-01-
01',periods=len(sales_data),freq='M')

#create a time series object
sales_series=pd.Series(sales_data,index=dates)
#apply ses
#create the ses model
model=Holt(sales_series)
#fit the model alpha=0.4,0.3,0.2
fit=model.fit(smoothing_level=0.4,optimized='False')

# forecast for next 6 month
forecast=fit.forecast(6)
#visualize result
#plotting the original data,fitted values and forecast
plt.figure(figsize=(10,5))
plt.plot(sales_series,label='Original Data',marker='o')
plt.plot(fit.fittedvalues,label='Fitted Values',linestyle='--',
'marker=x')
plt.plot(forecast,label='Forecast',linestyle='--',
'marker=o',color='red')

plt.title('Simple Exponential Smoothing')
plt.xlabel('Month')
plt.ylabel('Sales(in thousands)')
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()
```

## Output:



## Code:

```
#print forecast values
print("Forecasted sales(next 6 month)")
print(forecast.round(2))
```

## Output:

```
▶ #print forecast values
  print("Forecasted sales(next 6 month)")
  print(forecast.round(2))
```

```
→ Forecasted sales(next 6 month)
  2024-01-31    271.10
  2024-02-29    277.01
  2024-03-31    282.92
  2024-04-30    288.84
  2024-05-31    294.75
  2024-06-30    300.67
  Freq: ME, dtype: float64
```

## #using csv file

### Code:

```
#program
#csv files
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
#load dataset
df=pd.read_csv('electricity_usage.csv',parse_dates=['Date'],index_col='Date')
print(df.head())
#checking optional values
print(df.isnull().sum())
#create time series object
series = df['Consumption']
#apply ses
model=Holt(series)
```

```

fit=model.fit(smoothing_level=0.4,optimized='False')
#forecast next 6 month
forecast=fit.forecast(6)

#print forecast values
print("Next 6 month forecast:")
print(forecast)

```

## Output:

```

▶ #forecast next 6 month
forecast=fit.forecast(6)

#print forecast values
print("Next 6 month forecast:")
print(forecast)

```

→ Next 6 month forecast:  
 2022-01-05 321.749235  
 2022-01-06 322.811622  
 2022-01-07 323.874010  
 2022-01-08 324.936397  
 2022-01-09 325.998784  
 2022-01-10 327.061171  
 Freq: D, dtype: float64

## Code:

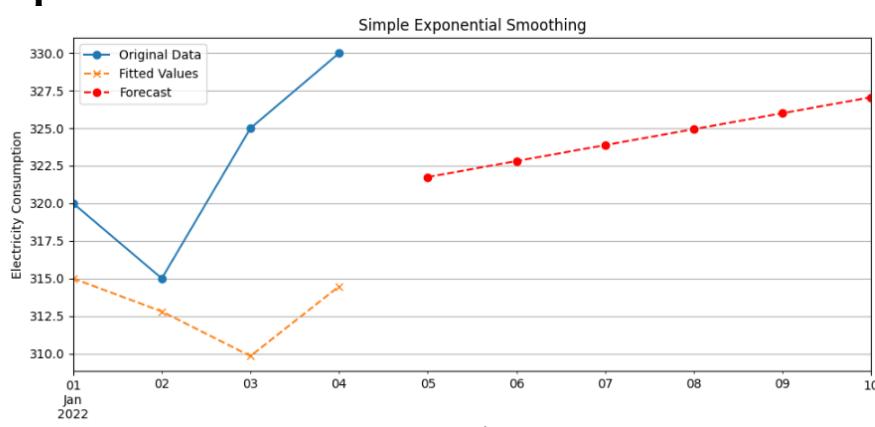
```

#plot the result
plt.figure(figsize=(10,5))
series.plot(label='Original Data',marker='o')
fit.fittedvalues.plot(label='Fitted Values',linestyle='--',marker='x')
forecast.plot(label='Forecast',linestyle='--',marker='o',color='red')

plt.title('Simple Exponential Smoothing')
plt.xlabel('Month')
plt.ylabel('Electricity Consumption')
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()

```

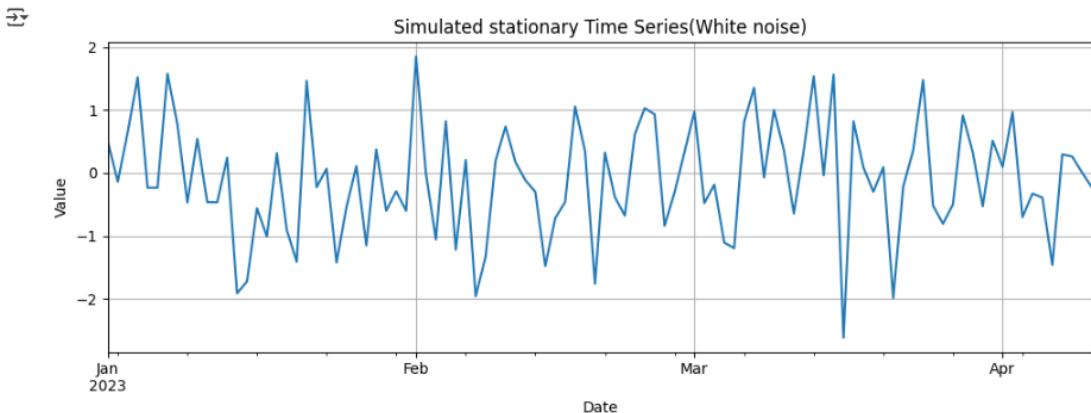
## Output:



## Code:

```
#stationary time series example
# import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
#set seed for reproducibility
np.random.seed(42)
#build model
data = np.random.normal(loc=0,scale=1,size=100)
dates=pd.date_range(start='2023-01-01',periods=100,freq='D')
series=pd.Series(data,index=dates)
#plot stationary data series
plt.figure(figsize=(10,4))
series.plot(title='Simulated stationary Time Series(White noise)')
plt.xlabel('Date')
plt.ylabel('Value')
plt.grid()
plt.tight_layout()
plt.show()
```

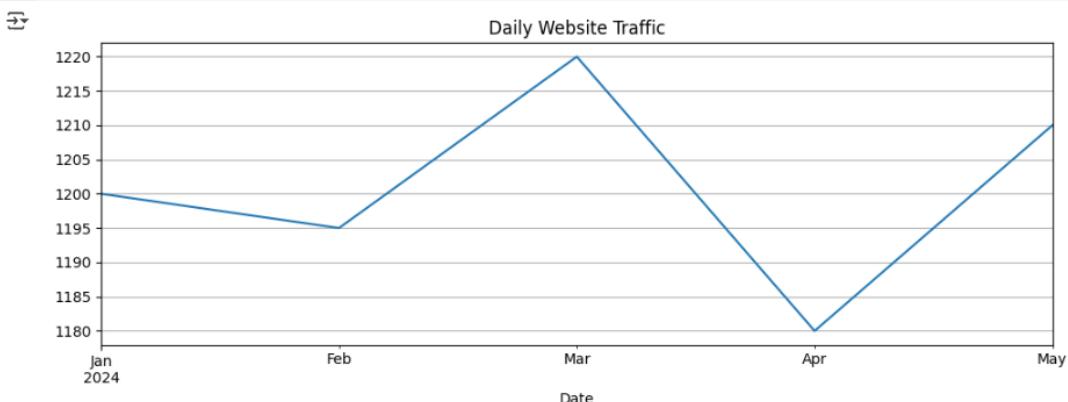
## Output:



## Code:

```
#check whether stationary or non stationary
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
df=pd.read_csv('website_traffic.csv',parse_dates=['Date'],index_col='Date')
df['Visits'].plot(figsize=(10,4),title="Daily Website Traffic")
plt.grid()
plt.tight_layout()
plt.show()
```

## Output:



## Code:

```
#perform augmented Dickey-Fuller test
result=adfuller(df['Visits'])
print("Augmented Dickey-Fuller Test:",result[0])
print("p-value:",result[1])
```

## Output:

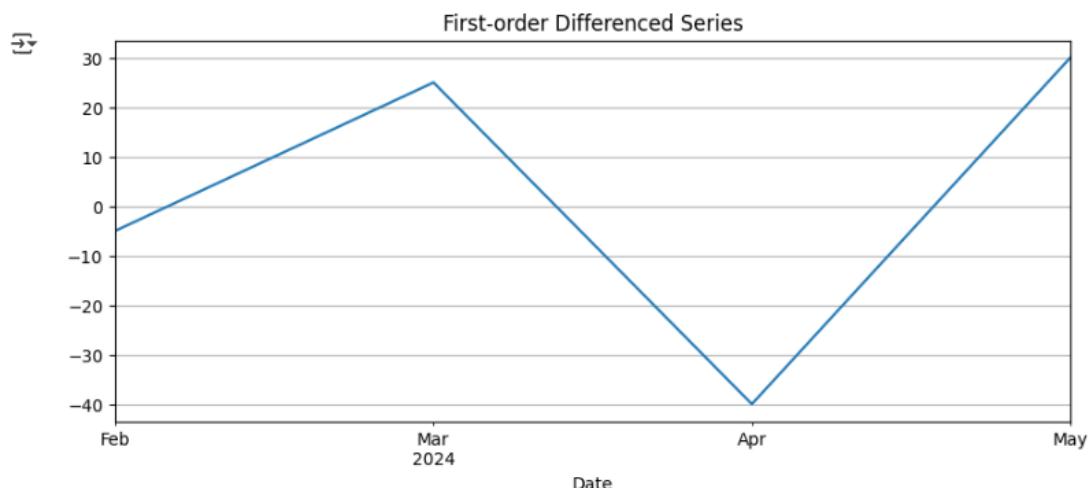
```
▶ #perform augmented Dickey-Fuller test
  result=adfuller(df['Visits'])
  print("Augmented Dickey-Fuller Test:",result[0])
  print("p-value:",result[1])

→ Augmented Dickey-Fuller Test: -4.109375445408512
  p-value: 0.0009345359045134847
```

## Code:

```
if result[1]<0.05:
    print("The series is stationary")
else:
    print("The series is non-stationary.Applying differencing...")
#first order differencing
df_diff=df['Visits'].diff().dropna()
#plot differenced series
df_diff.plot(figsize=(10,4),title="First-order Differenced Series")
plt.grid()
plt.show()
```

## Output:



## Code:

```
#re-run adf test on differenced series
result_diff= adfuller(df_diff)
print("Augmented Dickey-Fuller Test:",result_diff[0])
print("p-value after decreasing:",result_diff[1])
if result_diff[1]<0.05:
    print("The differenced series is now stationary")
else:
    print("The differenced series is now non-stationary")
```

## Output:

```
▶ #re-run adf test on differenced series
  result_diff= adfuller(df_diff)
  print("Augmented Dickey-Fuller Test:",result_diff[0])
  print("p-value after decreasing:",result_diff[1])

▶ Augmented Dickey-Fuller Test: -3.5388175323270006
  p-value after decreasing: 0.007044379192708437

[ ] if result_diff[1]<0.05:
  · print("The differenced series is now stationary")
  else:
  · print("The differenced series is now non-stationary")

▶ The differenced series is now stationary
```

## PRACTICAL N0 – 5

**Aim:** Simple Exponential Smoothing holt(in rstudio) & Exponential and moving average smoothing-Holts forecasting

### **Code:**

```
install.packages("forecast")
```

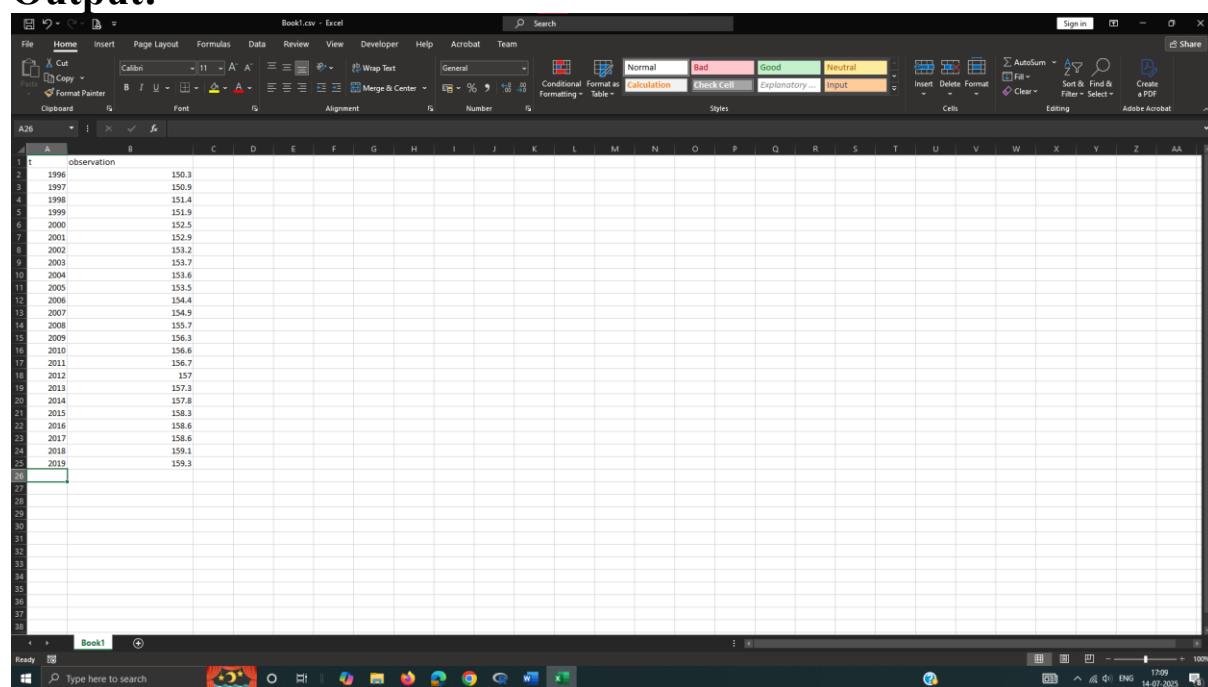
### **Output:**

```
> install.packages("forecast")
Warning in install.packages("forecast") :
  'lib' = "C:/Program Files/R/R-4.5.1/library" is not writable
--- Please select a CRAN mirror for use in this session ---
also installing the dependencies 'utf8', 'farver', 'labeling', 'R6', 'RColorBrewer', 'viridisLite', 'pillar', 'pkgconfig', 'xts', 'TTR', 'curl',
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/utf8_1.2.6.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/farver_2.1.2.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/labeling_0.4.3.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/R6_2.6.1.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/RColorBrewer_1.1-3.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/viridisLite_0.4.2.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/pillaz_1.11.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/pkgconfig_2.0.3.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/xts_0.14.1.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/TTR_0.24.4.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/curl_6.4.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/clm_3.6.5.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/glue_1.8.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/gtable_0.3.6.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/isoband_0.2.7.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/lifecycle_1.0.4.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/xlang_1.1.6.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/scales_1.4.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/tibble_3.3.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/vctrs_0.6.5.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/quadprog_1.5-8.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/quantmod_0.4.28.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/jsonlite_2.0.0.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/colorspace_2.1-1.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/fracdiff_1.5-3.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/generics_0.1.4.zip'
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/ggplot2_3.5.2.zip'
```

### **Code:**

```
library(forecast)
```

### **Output:**



	observation
t	
1	1996
2	150.3
3	150.9
4	151.4
5	151.9
6	152.5
7	152.9
8	153.2
9	153.7
10	153.6
11	153.5
12	154.4
13	154.9
14	155.7
15	156.3
16	156.6
17	156.7
18	157.
19	157.3
20	157.8
21	158.3
22	158.6
23	158.6
24	159.1
25	159.3
26	
27	
28	
29	
30	
31	
32	
33	
34	
35	
36	
37	
38	

	A	B	C	D	E	F	G	H	I
1	t	observation							
2	1996	150.3							
3	1997	150.9							
4	1998	151.4							
5	1999	151.9							
6	2000	152.5							
7	2001	152.9							
8	2002	153.2							
9	2003	153.7							
0	2004	153.6							
1	2005	153.5							
2	2006	154.4							
3	2007	154.9							
4	2008	155.7							
5	2009	156.3							
6	2010	156.6							
7	2011	156.7							
8	2012	157							
9	2013	157.3							
0	2014	157.8							
1	2015	158.3							
2	2016	158.6							
3	2017	158.6							
4	2018	159.1							
5	2019	159.3							
6									
7									
8									

### Code:

```
data<-read.csv('C:/Users/Admin/Desktop/Book1.csv')
```

data

### Output:

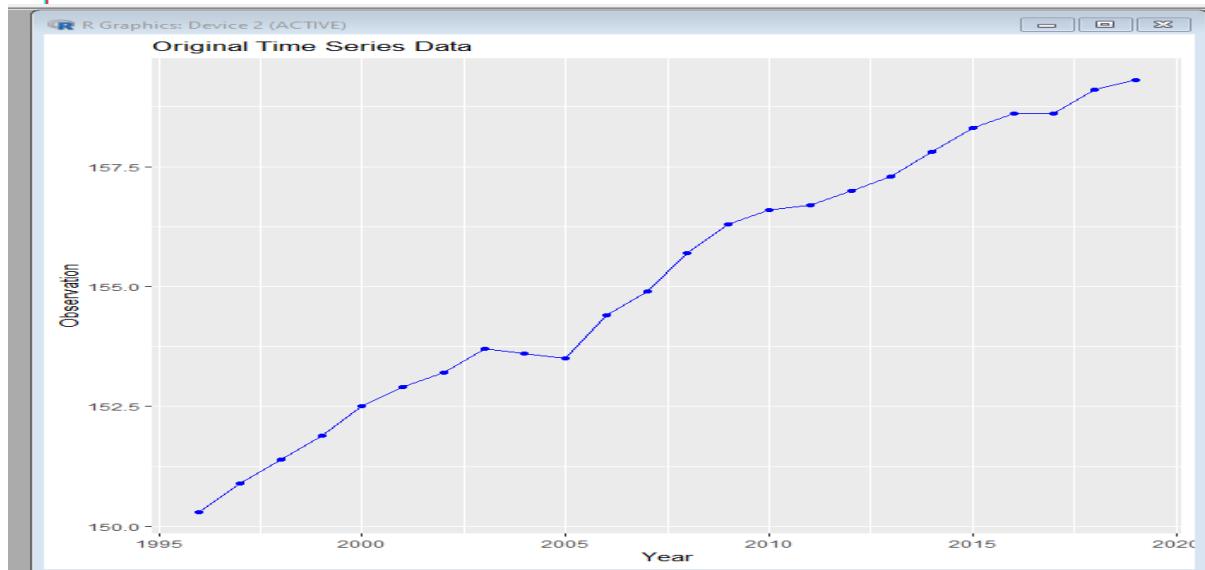
```
> data<-read.csv('C:/Users/Admin/Desktop/Book1.csv')
> data
   t observation
1 1996     150.3
2 1997     150.9
3 1998     151.4
4 1999     151.9
5 2000     152.5
6 2001     152.9
7 2002     153.2
8 2003     153.7
9 2004     153.6
10 2005    153.5
11 2006    154.4
12 2007    154.9
13 2008    155.7
14 2009    156.3
15 2010    156.6
16 2011    156.7
17 2012    157.0
18 2013    157.3
19 2014    157.8
20 2015    158.3
21 2016    158.6
22 2017    158.6
23 2018    159.1
24 2019    159.3
```

**Code:**

```
library(ggplot2)
ggplot(data,aes(x=t,y=observation))+
  geom_line(color="blue")+
  geom_point(color="blue")+
  ggtitle("Original Time Series Data")+
  xlab("Year")+
  ylab("Observation")
```

**Output:**

```
> ggplot(data,aes(x=t,y=observation))+  
+   geom_line(color="blue") +  
+   geom_point(color="blue") +  
+   ggtitle("Original Time Series Data") +  
+   xlab("Year") +  
+   ylab("Observation")
```

**Code:**

```
ses_model<-ses(data$observation, alpha=0.3,initial="simple",h=1)
ses_fitted<-fitted(ses_model)
ses_fitted
```

**Output:**

```
> ses_model<-ses(data$observation, alpha=0.3,initial="simple",h=1)
> ses_fitted<-fitted(ses_model)
> ses_fitted
Time Series:
Start = 1
End = 24
Frequency = 1
[1] 150.3000 150.3000 150.4800 150.7560 151.0992 151.5194 151.9336 152.3135
[9] 152.7295 152.9906 153.1434 153.5204 153.9343 154.4640 155.0148 155.4904
[17] 155.8533 156.1973 156.5281 156.9097 157.3268 157.7087 157.9761 158.3133
```

**Code:**

```
holt_model<-
holt(data$observation,alpha=0.3,beta=0.2,initial="simple",h=1)
holt_fitted<-fitted(holt_model)
holt_fitted
```

**Output:**

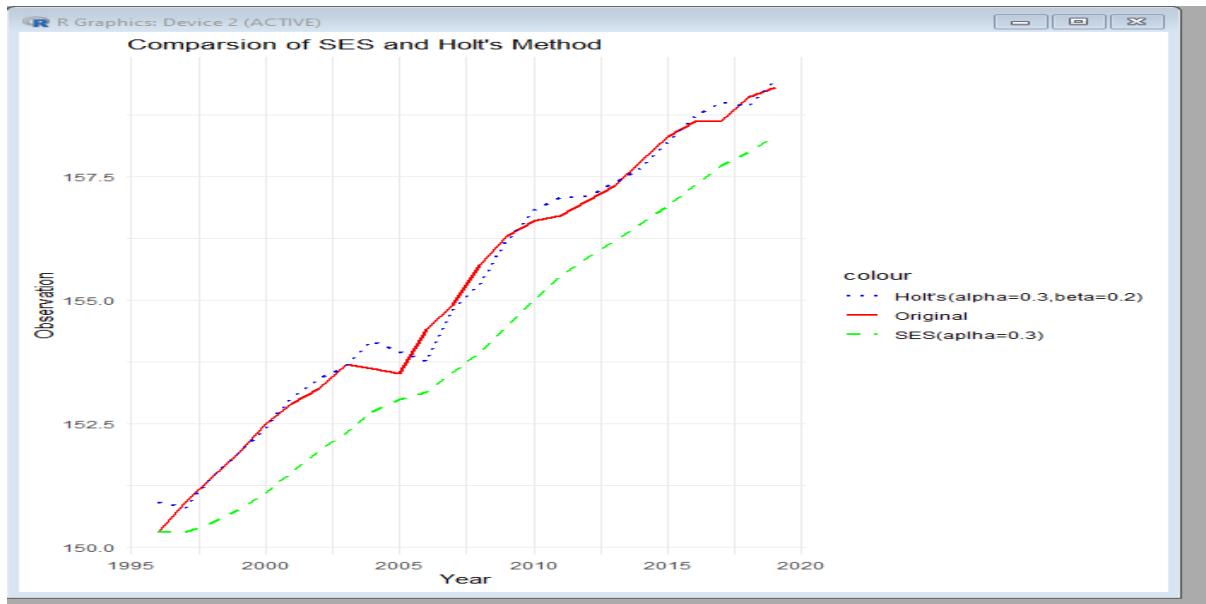
```
> holt_model<-holt(data$observation,alpha=0.3,beta=0.2,initial="simple",h=1)
> holt_fitted<-fitted(holt_model)
> holt_fitted
Time Series:
Start = 1
End = 24
Frequency = 1
[1] 150.9000 150.7800 151.4040 151.9032 152.4026 153.0220 153.3976 153.6581
[9] 154.1665 153.9532 153.7626 154.7900 155.3120 156.1896 156.8117 157.0694
[17] 157.0955 157.3764 157.6611 158.1889 158.7111 158.9889 158.9111 159.4489
```

**Code:**

```
ggplot()+
  geom_line(aes(x=data$t,
y=data$observation,color="Original"),size=1)+
  geom_line(aes(x=data$t, y=ses_fitted,
color="SES(alpha=0.3)"),linetype="dashed",size=1)+
  geom_line(aes(x=data$t,
y=holt_fitted,color="Holt's(alpha=0.3,beta=0.2)"),linetype="dotted",
size=1)+
  ggtitle("Comparsion of SES and Holt's Method")+
  xlab("Year")+
  ylab("Observation")+
  scale_color_manual(values=c("blue","red","green"))+
  theme_minimal()
```

**Output:**

```
> ggplot()+
+   geom_line(aes(x=data$t, y=data$observation,color="Original"),size=1)+
+   geom_line(aes(x=data$t, y=ses_fitted, color="SES(alpha=0.3)"),linetype="dashed",size=1)+
+   geom_line(aes(x=data$t, y=holt_fitted,color="Holt's(alpha=0.3,beta=0.2)"),linetype="dotted",
+   ggtitle("Comparsion of SES and Holt's Method")+
+   xlab("Year")+
+   ylab("Observation")+
+   scale_color_manual(values=c("blue","red","green"))+
+   theme_minimal()
```



## #Different dataset

### Code:

```

years<-1991:2003

co2_concentration<-
c(355.62,356.36,357.1,358.86,360.9,362.58,363.84,366.58,368.3,369
.47,371.03,373.61,357.61)

ts_co2<-ts(co2_concentration,start=min(years),frequency=1)

ts_co2

```

### Output:

```

> years<-1991:2003
> co2_concentration<-c(355.62,356.36,357.1,358.86,360.9,362.58,363.84,366.58,368.3,369.47,371.03,373.61,357.61)
> ts_co2
Time Series:
Start = 1991
End = 2003
Frequency = 1
[1] 355.62 356.36 357.10 358.86 360.90 362.58 363.84 366.58 368.30 369.47
[11] 371.03 373.61 357.61

```

### Code:

```

moving_avg_manual<-function(data,window_size){

n<-length(data)

ma<-rep(NA,n)

for(i in seq(window_size,n)){

```

```

ma[i]<-mean(data[(i>window_size+1):i])
}
return(ma)
}

window_size<-3
ma_co2<-moving_avg_manual(co2_concentration,window_size)
ma_co2

```

**Output:**

```

> moving_avg_manual<-function(data,window_size) {
+ n<-length(data)
+ ma<-rep(NA,n)
+ for(i in seq(window_size,n)){
+ ma[i]<-mean(data[(i>window_size+1):i])
+ }
+ return(ma)
+ }
> window_size<-3
> ma_co2<-moving_avg_manual(co2_concentration,window_size)
> ma_co2
[1]      NA      NA 356.3600 357.4400 358.9533 360.7800 362.4400 364.3333
[9] 366.2400 368.1167 369.6000 371.3700 367.4167
> |

```

**Code:**

```

forecast_2004<-tail(ma_co2,1)
cat("Forecasted co2 concentration for 2004 using 3-year Moving
Average:",forecast_2004,"\n")

```

**Output:**

```

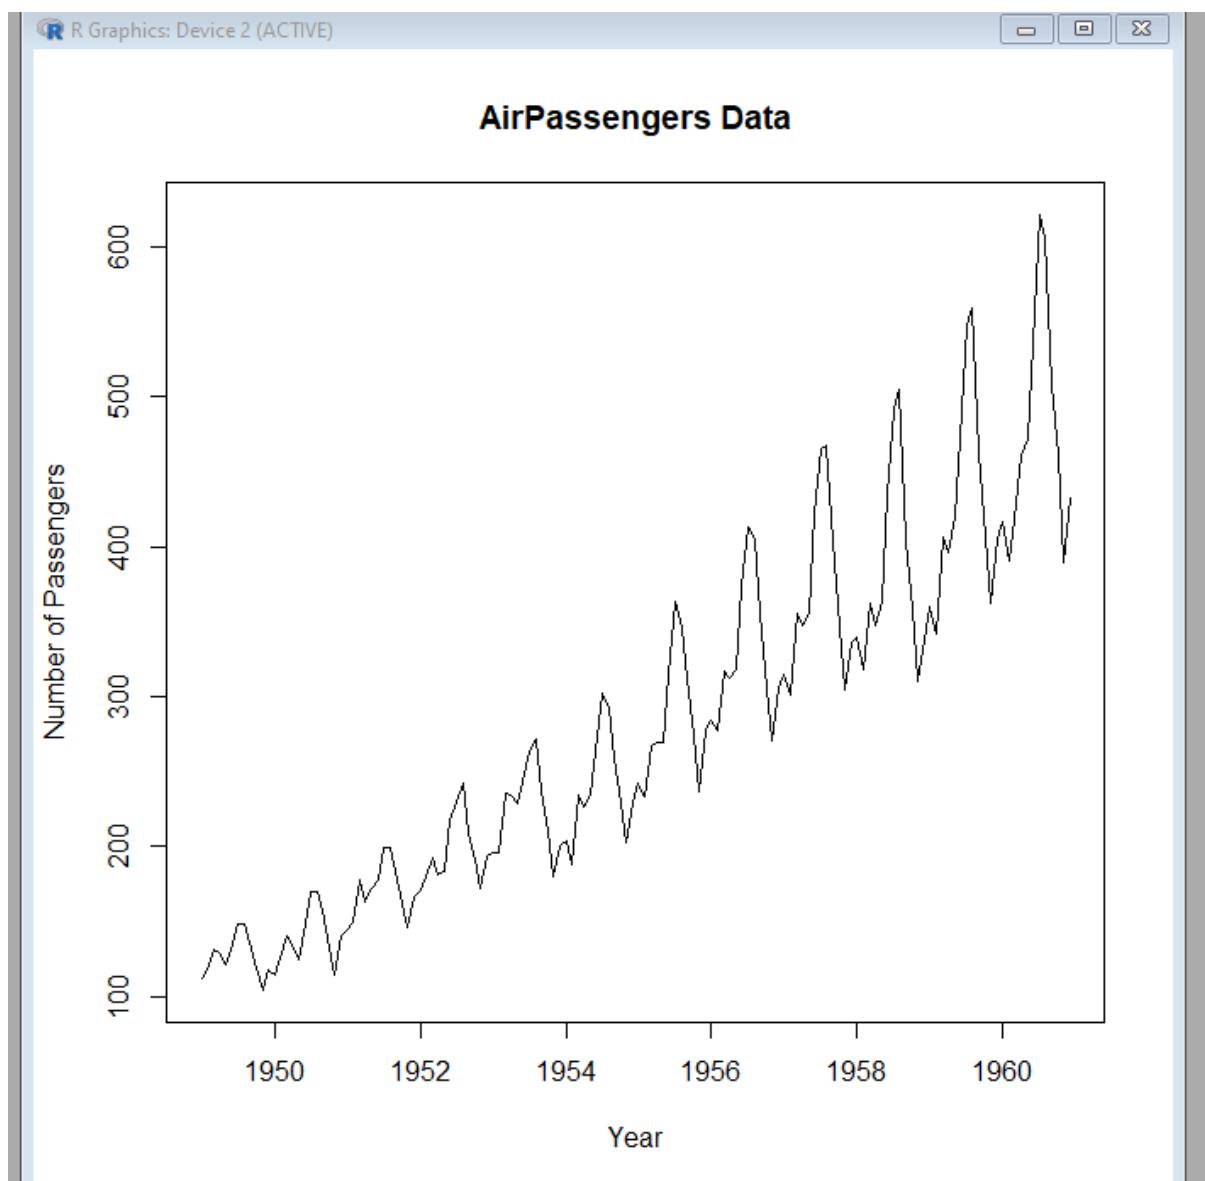
> forecast_2004<-tail(ma_co2,1)
> cat("Forecasted co2 concentration for 2004 using 3-year Moving Average:",forecast_2004,"\n")
Forecasted co2 concentration for 2004 using 3-year Moving Average: 367.4167

```

**Code:**

```
install.packages("forecast")
library(forecast)
data("AirPassengers")
plot(AirPassengers,main="AirPassengers Data",ylab="Number of
Passengers",xlab="Year")
```

**Output:**



#Exponential and moving average smoothing-Holts forecasting

**Code:**

```
data<-
c(150.3,150.9,151.4,151.9,152.5,152.9,153.2,153.7,153.6,153.5,154.
4,154.9,155.7,156.3,156.6,156.7,157,157.3,157.8,158.3,158.6,158.6,
159.1,159.3)

time_series_data<-ts(data,start=1996,frequency=1)

#A)simple exponential smoothing aplha=0.3

#load necessary package

library(forecast)
```

**Output:**

```
> #Exponential and moving average smoothing-Holts forecasting
> data<-c(150.3,150.9,151.4,151.9,152.5,152.9,153.2,153.7,153.6,153.5,154.4,154.9,155.7,156.3,156.6,156.7,157,157.3,157.8,158.3,158.6,158.6,159.1,159.3)
> time_series_data<-ts(data,start=1996,frequency=1)
> #A)simple exponential smoothing aplha=0.3
> #load necessary package
> library(forecast)
Registered S3 method overwritten by 'quantmod':
  method           from
  as.zoo.data.frame zoo
  ....
```

**Code:**

```
ses_model<-ses(time_series_data,aplha=0.3,h=1)

summary(ses_model)
```

**Output:**

```
> ses_model<-ses(time_series_data,aplha=0.3,h=1)
> summary(ses_model)

Forecast method: Simple exponential smoothing

Model Information:
Simple exponential smoothing [REDACTED]

Call:
ses(y = time_series_data, h = 1, aplha = 0.3)

Smoothing parameters:
  aplha = 0.9999

Initial states:
  l = 150.3001

sigma:  0.4749

      AIC      AICc      BIC
44.43717 45.63717 47.97134

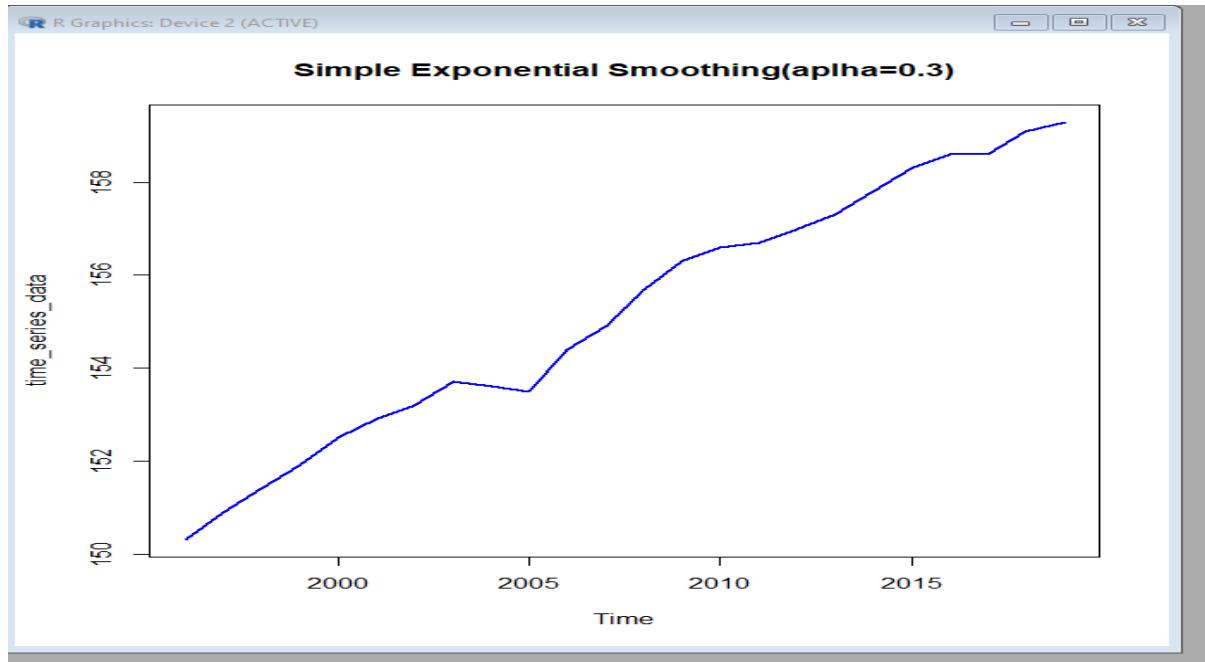
Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE
Training set 0.3750322 0.4546386 0.3917047 0.2419044 0.2527626 0.9584264
          ACF1
Training set 0.0627605

Forecasts:
      Point Forecast    Lo 80     Hi 80    Lo 95     Hi 95
2020       159.3 158.6914 159.9085 158.3693 160.2307
```

### Code:

```
#plot the ses model  
  
plot(time_series_data,main="Simple Exponential  
Smoothing(alpha=0.3)",col="blue",lwd=2)
```

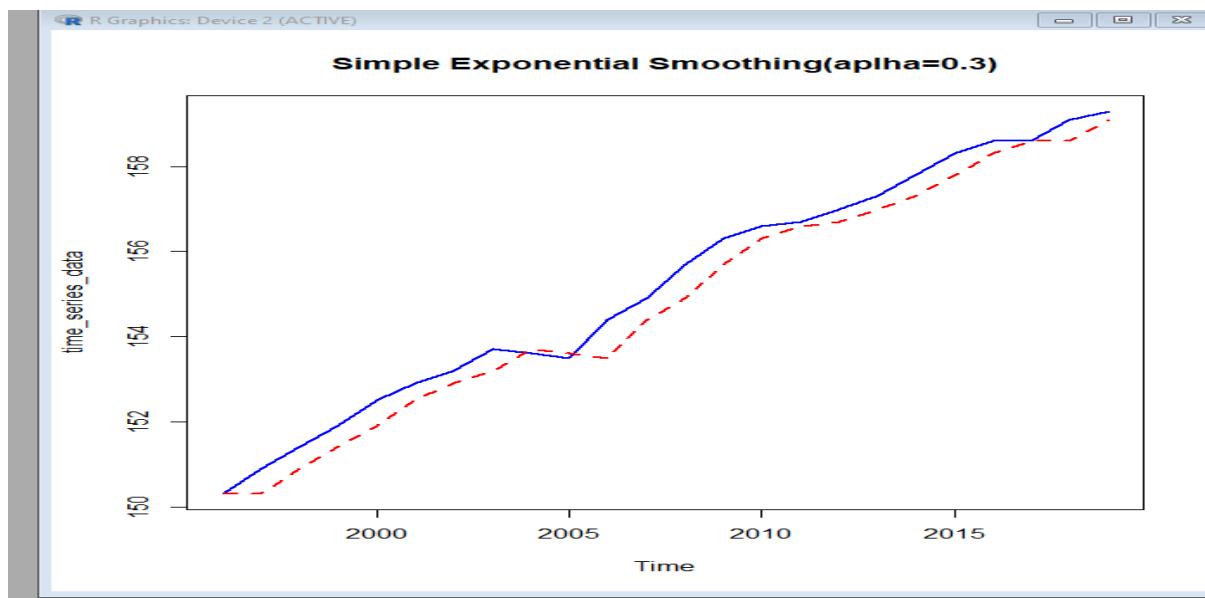
### Output:



### Code:

```
lines(fitted(ses_model),col="red",lty=2,lwd=2)
```

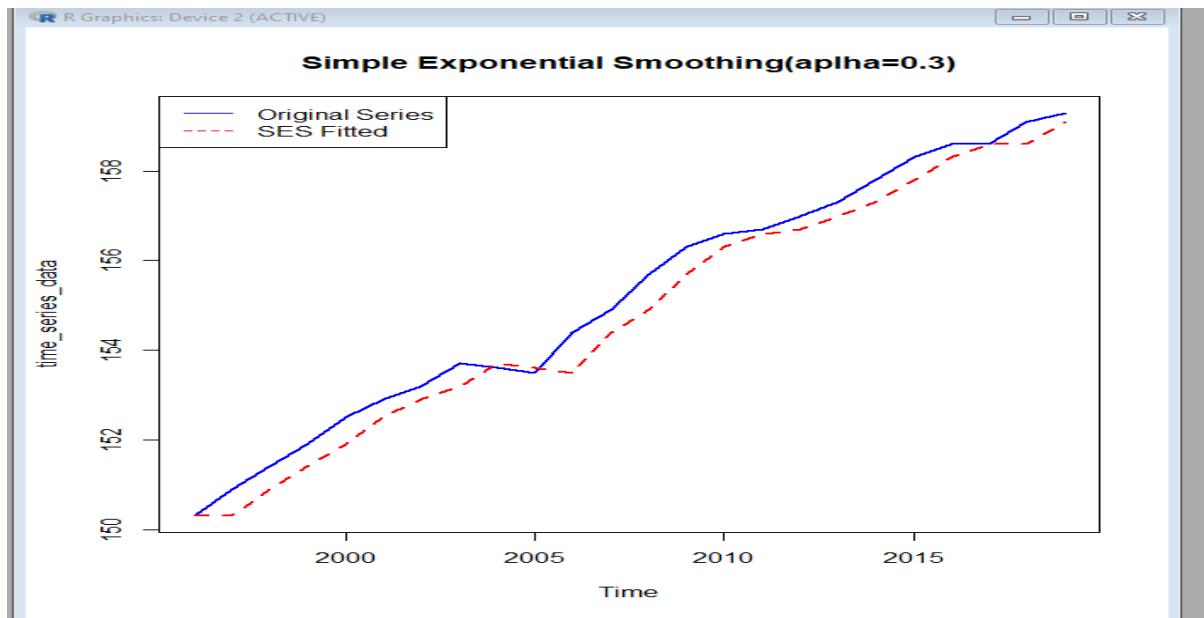
### Output:



## Code:

```
legend("topleft",legend=c("Original Series","SES  
Fitted"),col=c("blue","red"),lty=1:2)
```

## Output:

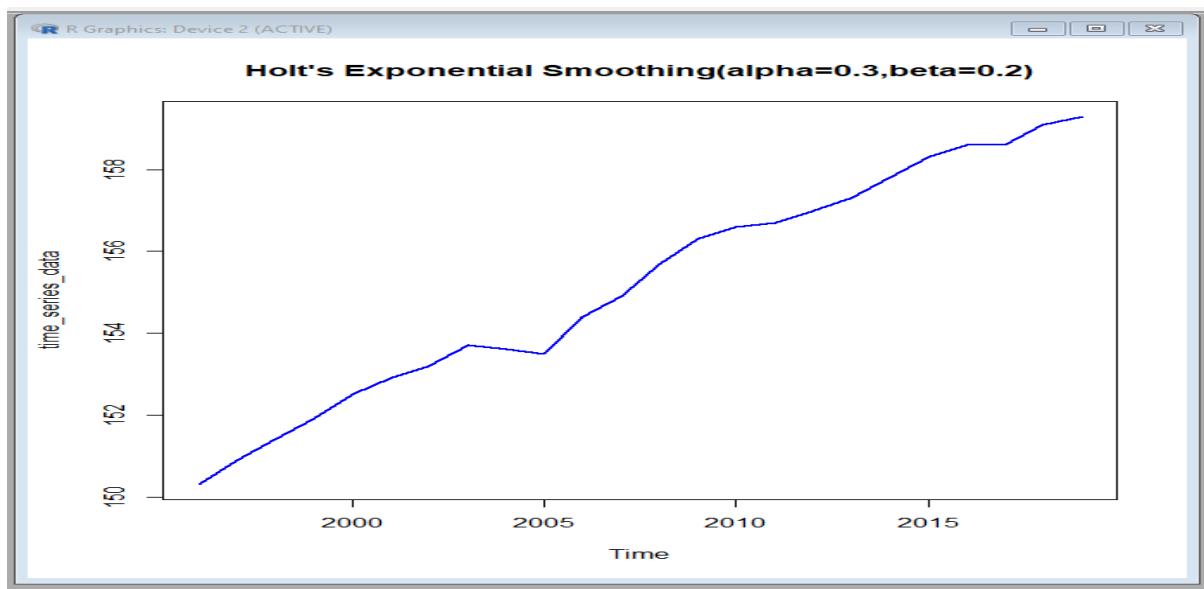


## Code:

```
#plot the holt's model
```

```
plot(time_series_data,main="Holt's Exponential  
Smoothing(alpha=0.3,beta=0.2)",col="blue",lwd=2)
```

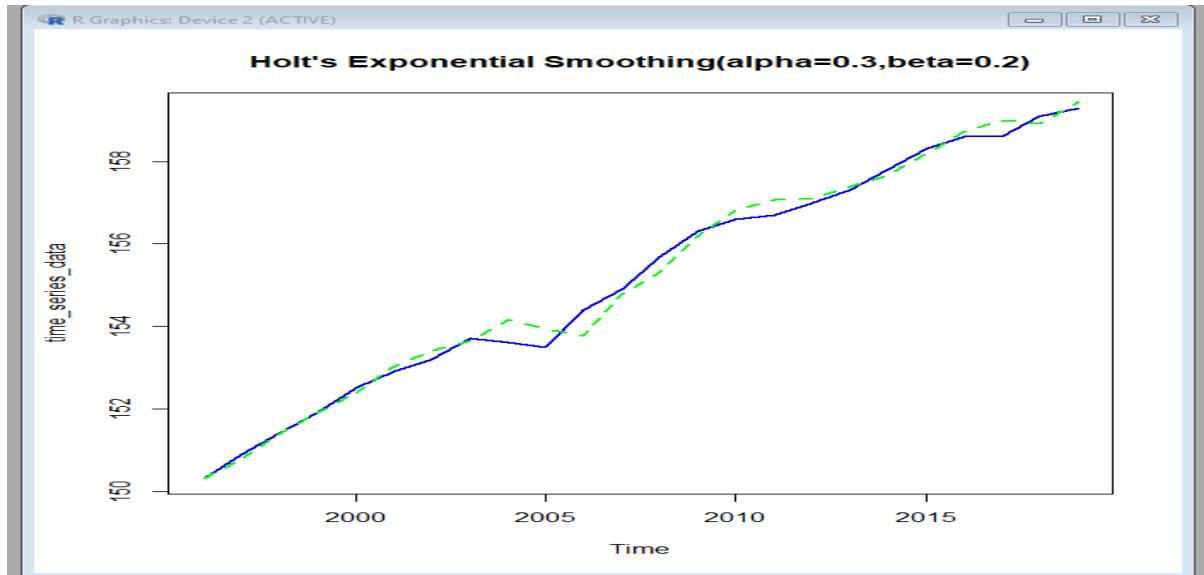
## Output:



**Code:**

```
lines(fitted(holt_model),col="green",lty=2,lwd=2)
```

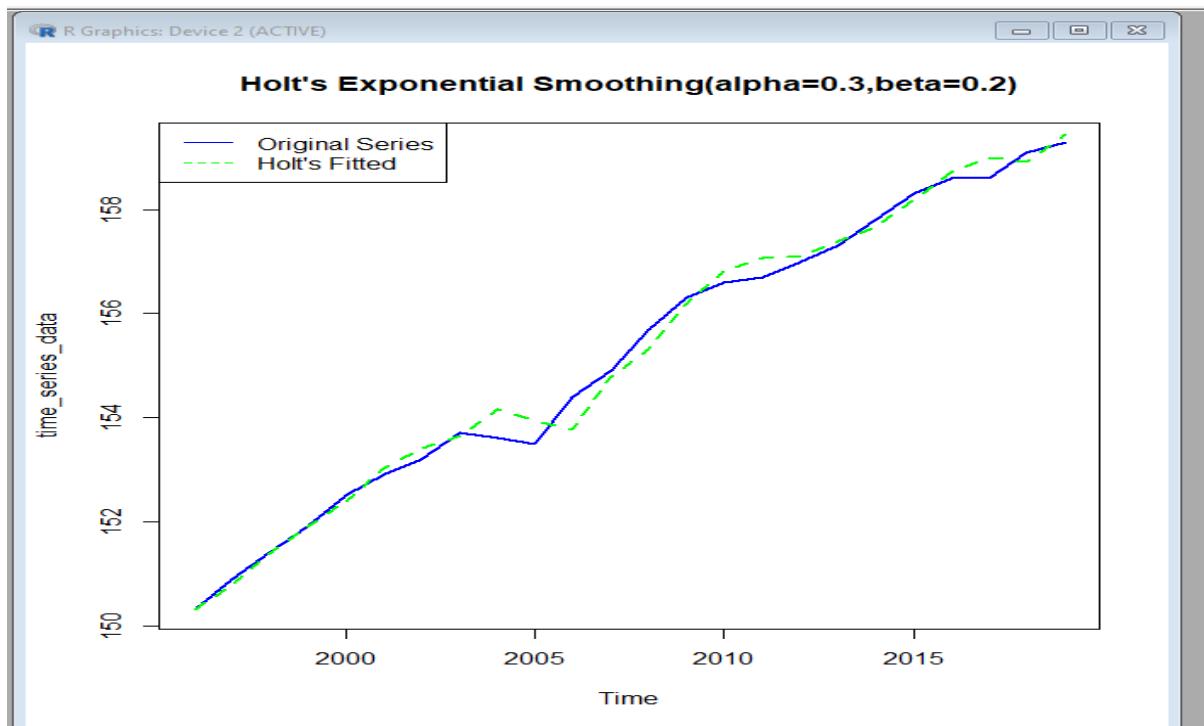
**Output:**



**Code:**

```
legend("topleft",legend=c("Original Series","Holt's Fitted"),col=c("blue","green"),lty=1:2)
```

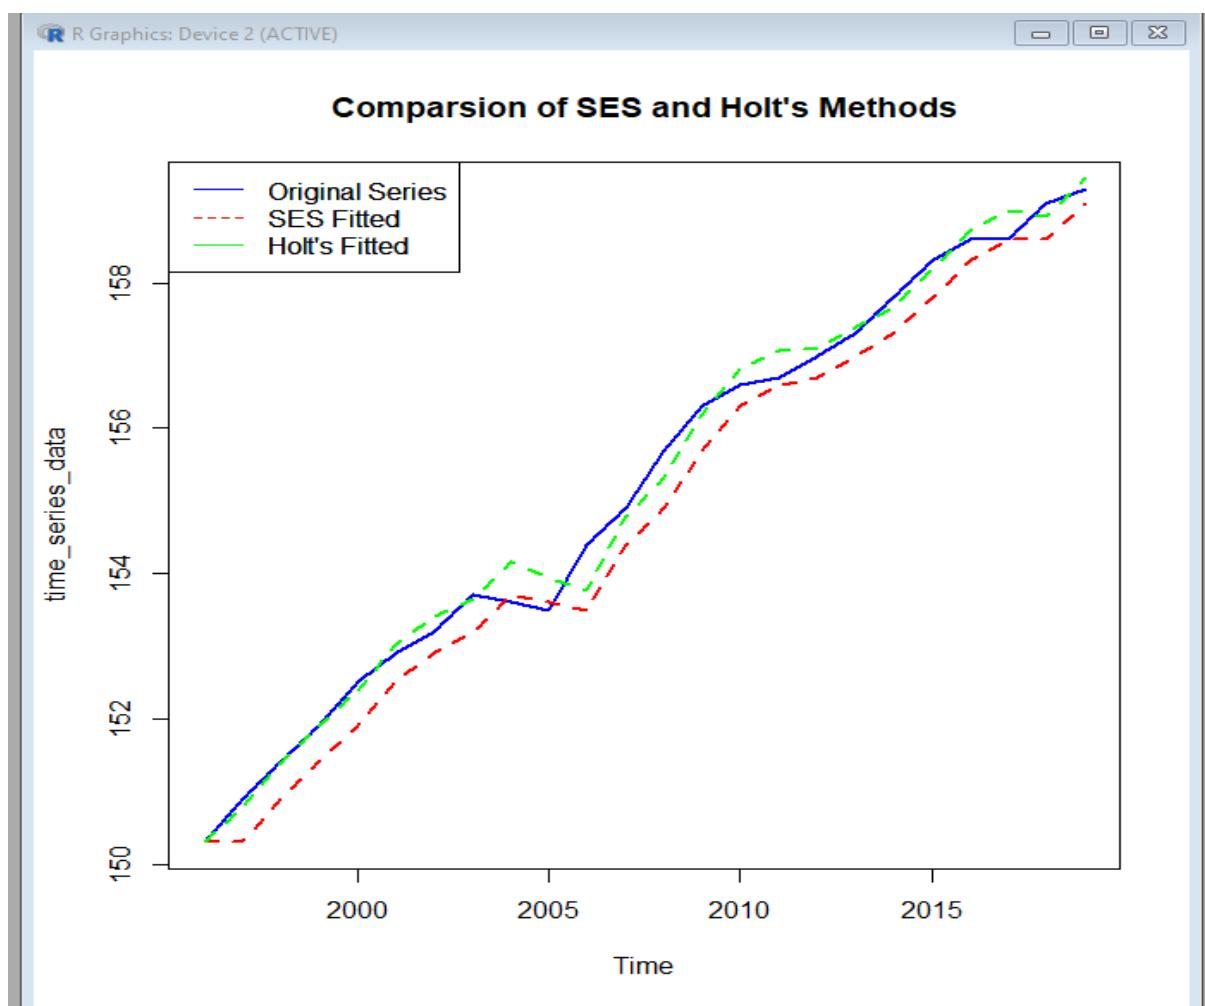
**Output:**



## Code:

```
#plotting both model  
  
plot(time_series_data,main="Comparsion of SES and Holt's  
Methods",col="blue",lwd=2)  
  
lines(fitted(ses_model),col="red",lty=2,lwd=2)  
  
lines(fitted(holt_model),col="green",lty=2,lwd=2)  
  
legend("topleft",legend=c("Original Series","SES Fitted","Holt's  
Fitted"),col=c("blue","red","green"),lty=1:2)
```

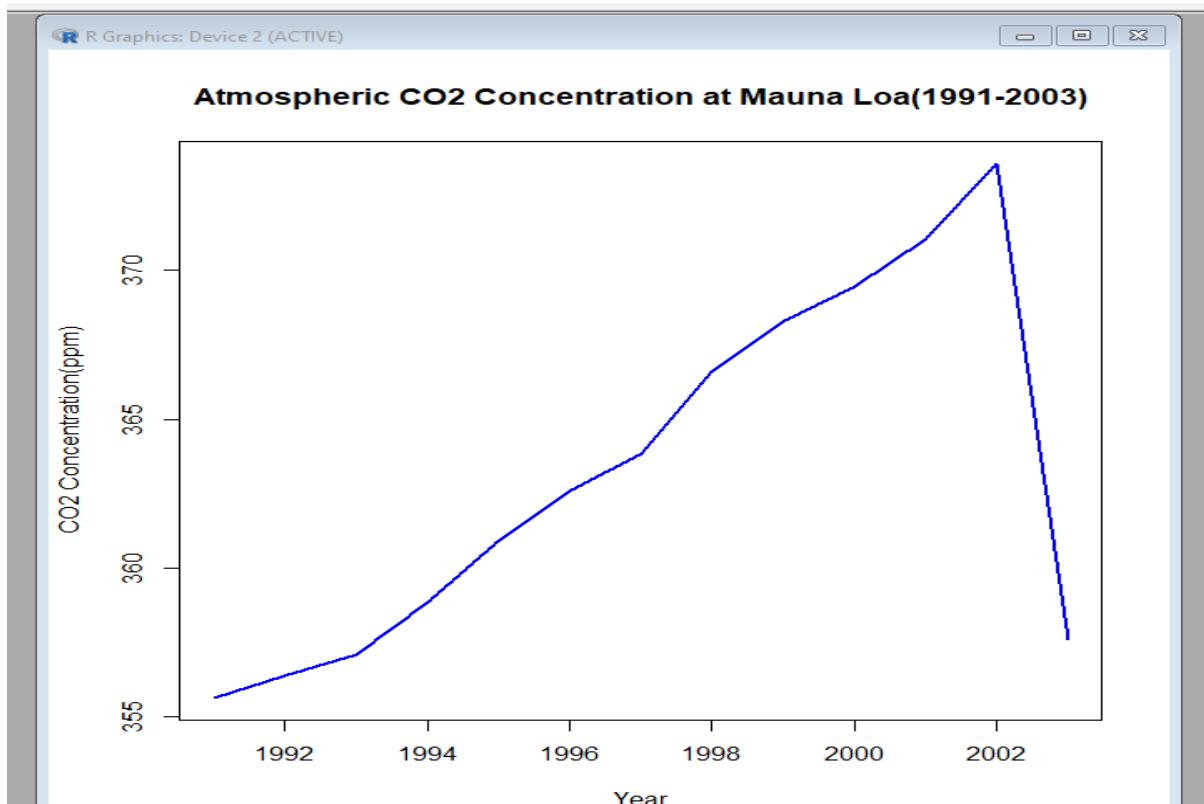
## Output:



## Code:

```
#Exponential and moving average smoothing 3-years  
#given data  
co2_concentration<-  
c(355.62,356.36,357.1,358.86,360.9,362.58,363.84,366.58,368.3,369  
.47,371.03,373.61,357.61)  
#create a time series object for the co2 concentration data  
co2_ts<-ts(co2_concentration,start=1991,frequency=1)  
#plotting the time series  
plot(co2_ts,main="Atmospheric CO2 Concentration at Mauna  
Loa(1991-2003)",xlab="Year",ylab="CO2  
Concentration(ppm)",col="blue",lwd=2)
```

## Output:



**Code:**

```
#Calculate 3-year moving average  
moving_average_3<-filter(co2_ts,rep(1/3,3),sides=2)  
#forecast the value for 2004  
#using the last three values(2001,2002,2003)to forecast for 2004  
forecast_2004<-mean(co2_concentration[11:13])  
forecast_2004
```

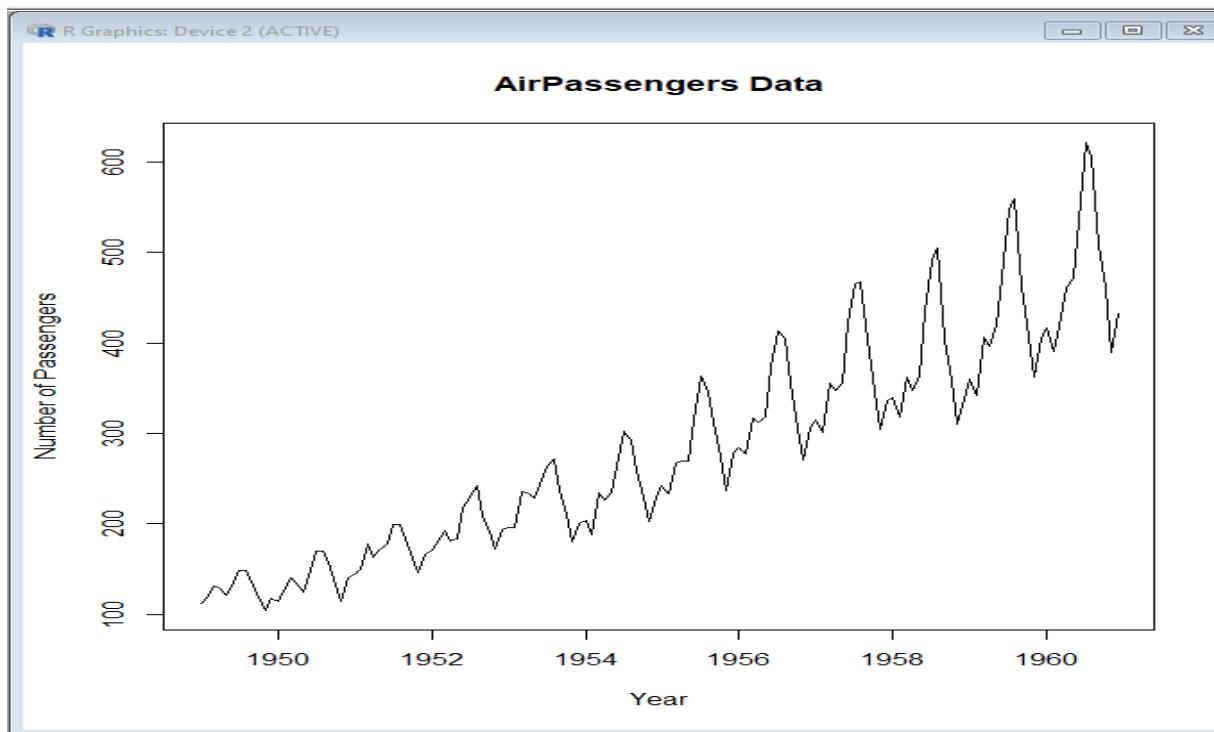
**Output:**

```
> #Calculate 3-year moving average  
> moving_average_3<-filter(co2_ts,rep(1/3,3),sides=2)  
> #forecast the value for 2004  
> #using the last three values(2001,2002,2003)to forecast for 2004  
> forecast_2004<-mean(co2_concentration[11:13])  
> forecast_2004  
[1] 367.4167  
`|
```

**Code:**

```
#Exponential and moving average smoothing for holt forecasting  
library(stats)  
data("AirPassengers")  
#plot the time series  
plot(AirPassengers,main="AirPassengers  
Data",xlab="Year",ylab="Number of Passengers")
```

**Output:**



### Code:

```
#apply holt-winters model (multiplicative)
hw_model<-HoltWinters(AirPassengers,seasonal="multiplicative")
#summary
summary(hw_model)
```

### Output:

```
> summary(hw_model)
      Length Class  Mode
fitted     528   mts   numeric
x         144    ts   numeric
alpha       1   -none- numeric
beta        1   -none- numeric
gamma       1   -none- numeric
coefficients 14   -none- numeric
seasonal     1   -none- character
SSE         1   -none- numeric
call        3   -none- call
```

### Code:

```
#forecast for the next 12 months
forecast<-predict(hw_model,n.ahead=12)
forecast
```

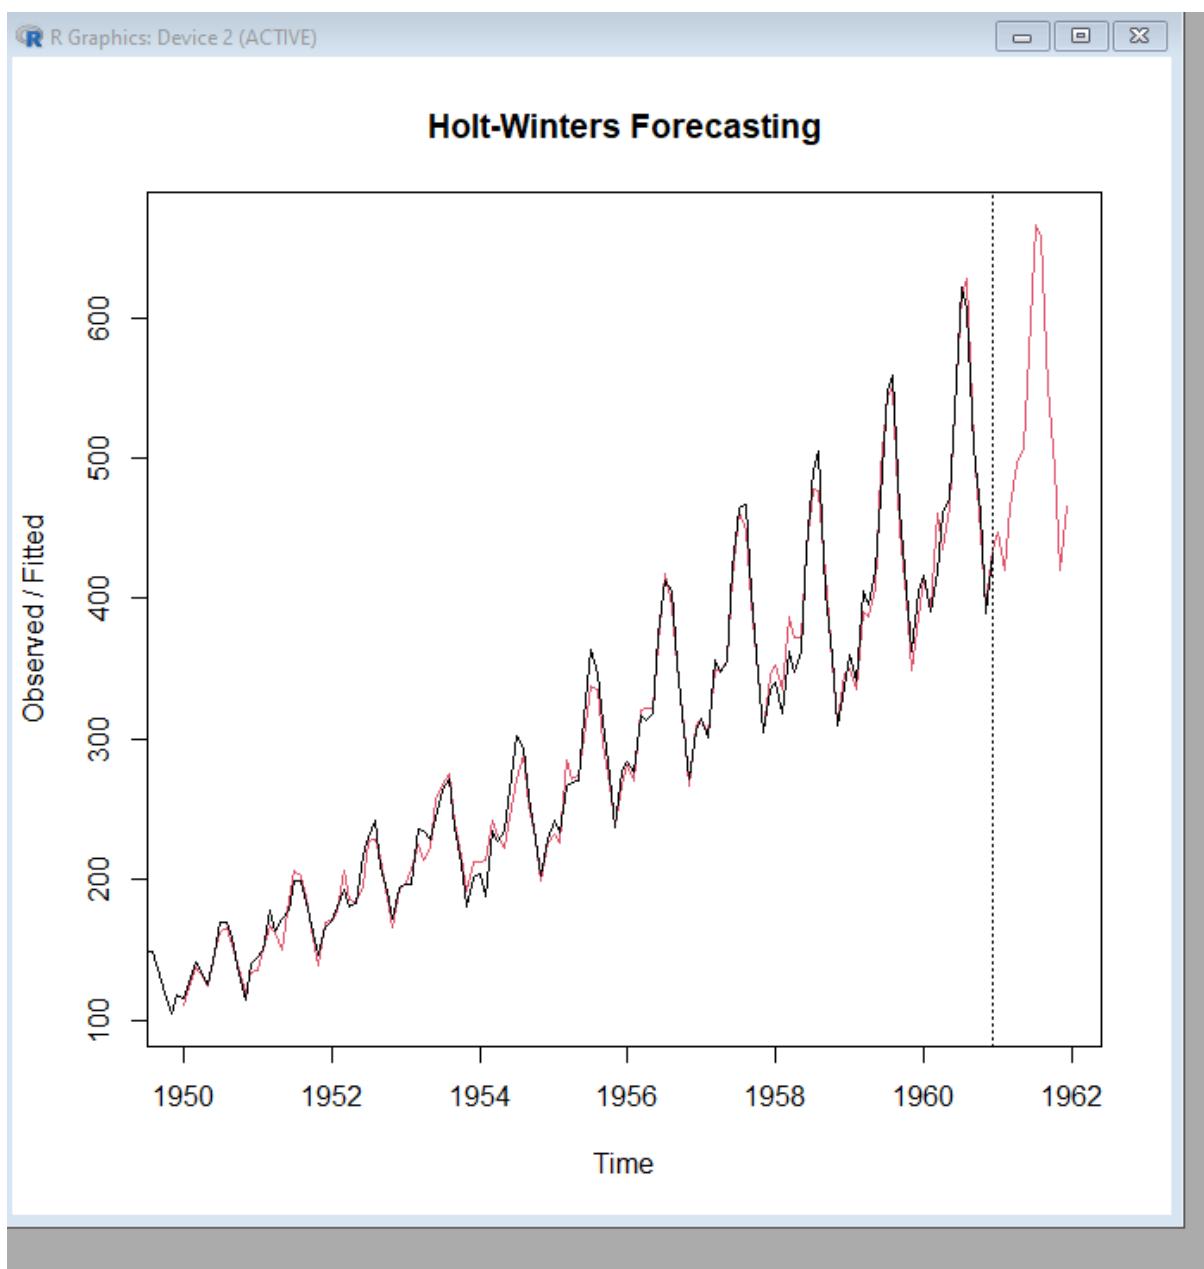
## Output:

```
> #forecast for the next 12 months
> forecast<-predict(hw_model,n.ahead=12)
> forecast
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug
1961 447.0559 419.7123 464.8671 496.0839 507.5326 575.4509 666.5923 657.9137
      Sep      Oct      Nov      Dec
1961 550.3088 492.9853 420.2073 465.6345
> |
```

## Code:

```
#plot the original series and forecast
```

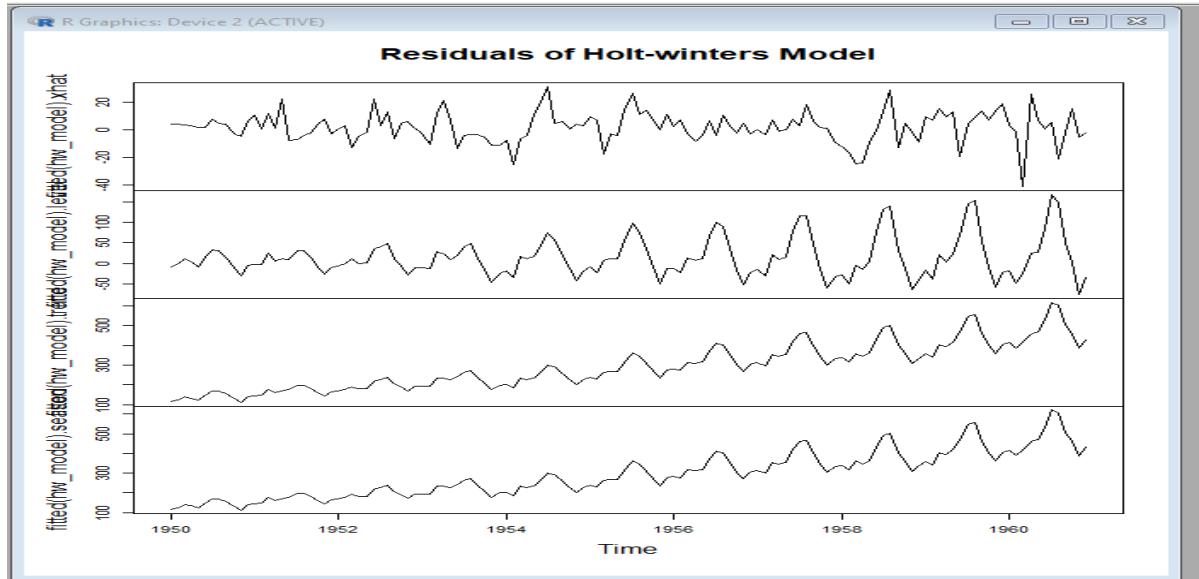
```
plot(hw_model,forecast,main="Holt-Winters Forecasting")
```



## Code:

```
#plot the original series and forecast  
residuals<-AirPassengers-fitted(hw_model)  
plot(residuals,main="Residuals of Holt-winters Model")
```

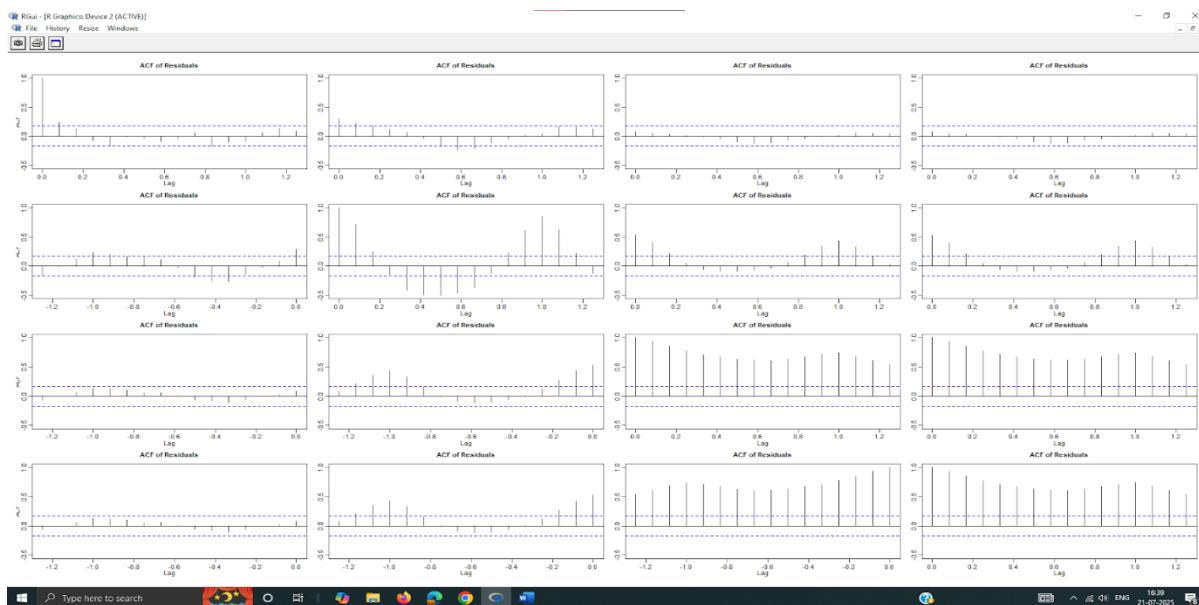
## Output:



## Code:

```
#check for autocorrelation  
acf(residuals,main="ACF of Residuals")
```

## Output:



## PRACTICAL N0 – 6

**Aim:** Autocorrelation & Autoregression using Arima

**Code:**

```
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA
data=[100,105,102,107,110,112,115,117]
s = pd.Series(data)
#model
model = ARIMA(s,order=(1,0,1))
result = model.fit()
#summary
print(result.summary())
```

**Output:**

```
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA

data=[100,105,102,107,110,112,115,117]
s = pd.Series(data)

#model
model = ARIMA(s,order=(1,0,1))
result = model.fit()

#summary
print(result.summary())

SARIMAX Results
=====
Dep. Variable: y No. Observations: 8
Model: ARIMA(1, 0, 1) Log Likelihood: -21.936
Date: Sat, 26 Jul 2025 AIC: 51.872
Time: 10:46:48 BIC: 52.190
Sample: 0 HQIC: 49.729
Covariance Type: opg
=====
coef std err z P>|z| [0.025 0.975]
-----
```

**Code:**

```
#forecast
future=result.forecast(steps=3)
print("Next 3 prediction:",future.tolist())
```

**Output:**

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
#forecast
future=result.forecast(steps=3)
print("Next 3 prediction:",future.tolist())

Next 3 prediction: [115.95972016043233, 114.96894031537057, 114.10923549133534]

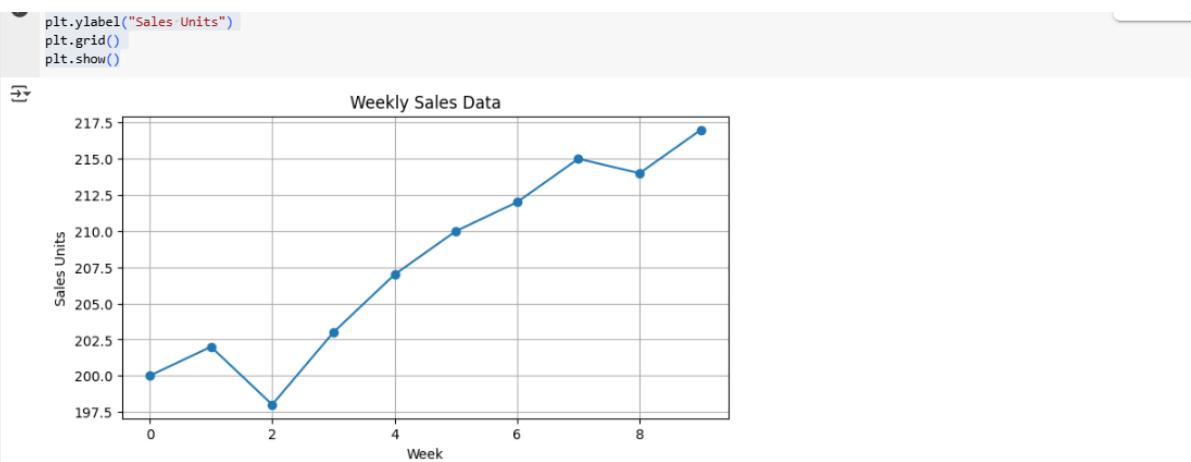
#next datasets
```

## #Using different dataset

### Code:

```
#next datasets
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
#load data
data=[200,202,198,203,207,210,212,215,214,217]
series=pd.Series(data)
#plot time series
plt.figure(figsize=(8,4))
plt.plot(series,marker='o')
plt.title("Weekly Sales Data")
plt.xlabel("Week")
plt.ylabel("Sales Units")
plt.grid()
plt.show()
```

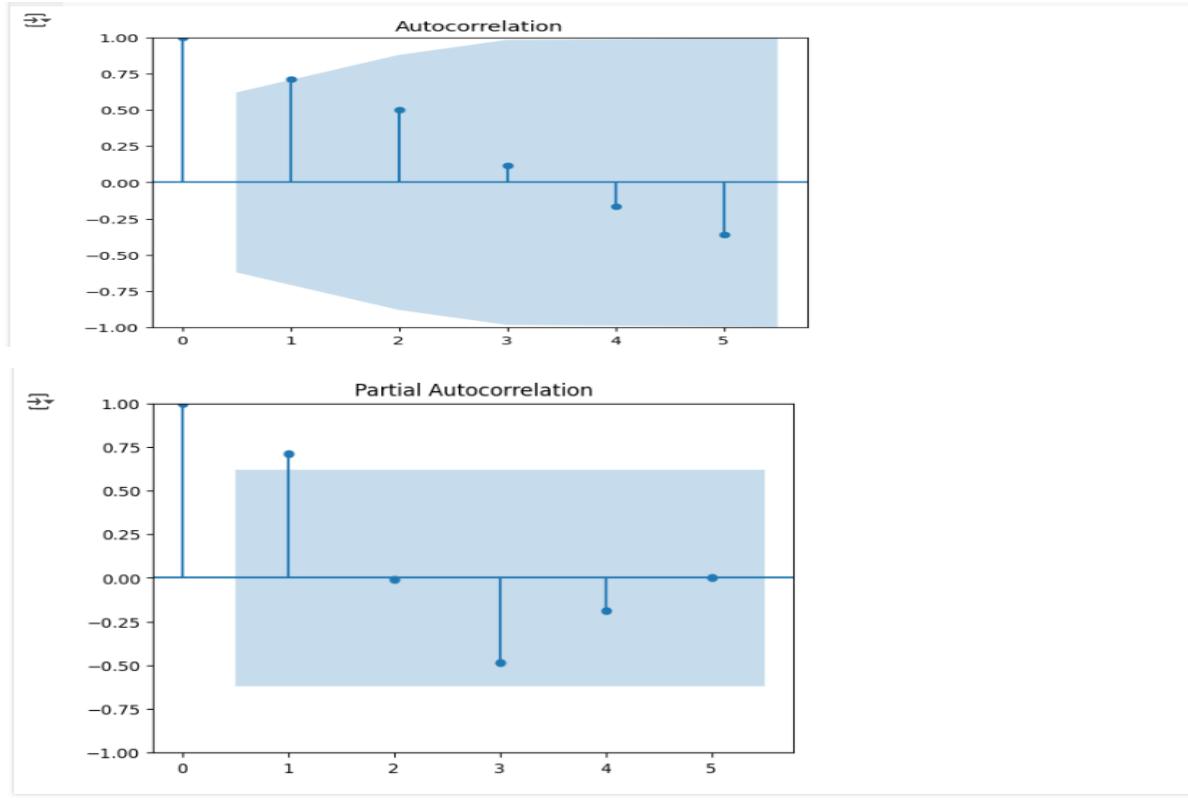
### Output:



### Code:

```
#ACF & PACF
plot_acf(series, lags=5)
plt.show()
plot_pacf(series, lags=5)
plt.show()
```

## Output:



```
[ ] #arima  
model=ARIMA(series,order=(1,0,1))
```

## Code:

```
#arima  
model=ARIMA(series,order=(1,0,1))  
result=model.fit()  
print(result.summary())
```

## Output:

SARIMAX Results

Dep. Variable:	y	No. Observations:	10
Model:	ARIMA(1, 0, 1)	Log Likelihood:	-26.529
Date:	Sat, 26 Jul 2025	AIC:	61.058
Time:	11:04:43	BIC:	62.269
Sample:	- 10	HQIC:	59.731
Covariance Type:	opg		

coef	std err	z	P> z	[0.025	0.975]
const	208.2588	6.325	32.926	0.000	195.862
ar.L1	0.8934	0.208	4.297	0.000	0.486
ma.L1	0.0806	0.443	0.182	0.856	-0.789
sigma2	9.9080	11.073	0.895	0.371	-11.796

Ljung-Box (L1) (Q): 0.42 Jarque-Bera (JB): 1.59  
Prob(Q): 0.52 Prob(JB): 0.45  
Heteroskedasticity (H): 0.67 Skew: -0.94  
Prob(H) (two-sided): 0.75 Kurtosis: 2.48

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Code:

```
#Forecast  
forecast=result.forecast(steps=3)  
print("\nForecast for next 3 weeks:")  
print(forecast)
```

## Output:

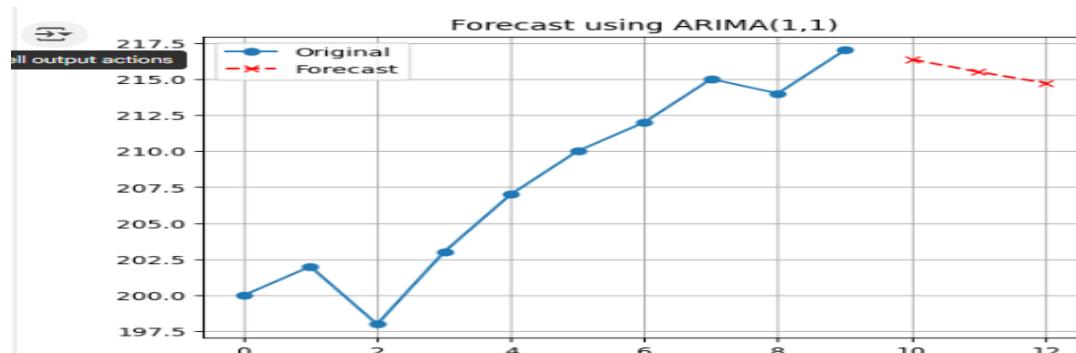
```
#Forecast  
forecast=result.forecast(steps=3)  
print("\nForecast for next 3 weeks:")  
print(forecast)
```

Forecast for next 3 weeks:  
10 216.362951  
11 215.499264  
12 214.727623  
Name: predicted\_mean, dtype: float64

## Code:

```
#plot forecast  
plt.plot(series,label='Original',marker='o')  
forecast_index=range(len(series),len(series)+3)  
plt.plot(forecast_index,forecast,label='Forecast',marker='x',linestyle='--',color='red')  
plt.legend()  
plt.title("Forecast using ARIMA(1,1)")  
plt.grid()  
plt.show()  
print(result.summary())
```

## Output:



```
[ ] print(result.summary())

```

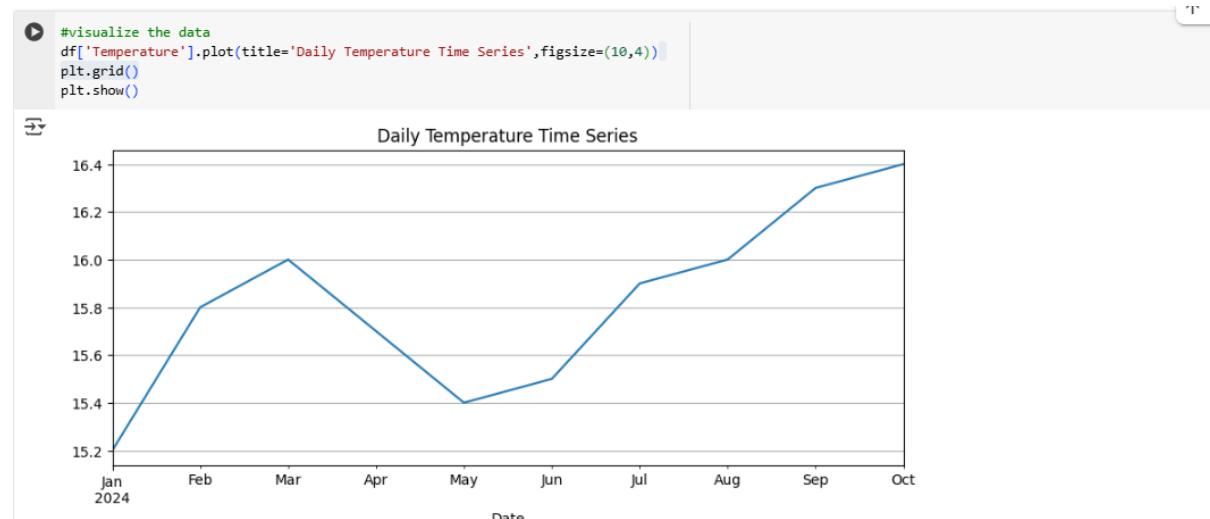
SARIMAX Results				
Dep. Variable:	y	No. Observations:	10	
Model:	ARIMA(1, 0, 1)	Log Likelihood:	-26.529	
Date:	Sat, 26 Jul 2025	AIC:	61.058	
Time:	11:22:53	BIC:	62.269	
Sample:	0 - 10	HQIC:	59.731	

## #next dataset from csv file

### Code:

```
##next dataset from csv file
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from statsmodels.tsa.ar_model import AutoReg
from statsmodels.tsa.stattools import acf
#load dataset
df=pd.read_csv("temperature_data.csv",parse_dates=['Date'],index_col='Date')
#visualize the data
df['Temperature'].plot(title='Daily Temperature Time Series',figsize=(10,4))
plt.grid()
plt.show()
```

### Output:



## Code:

```
#calculate and plot autocovariance manually
def autocovariance(x,lag):
    x_mean=np.mean(x)
    n=len(x)
    return np.sum((x[:n-lag]-x_mean)*(x[lag:]-x_mean))/n
#print autocovariance for lags 1 to 5
for lag in range(1,6):
    cov=autocovariance(df['Temperature'].values,lag)
    print(f"Autocovariance for lag {lag}: {cov:.4f}")
```

## Output:

```
[ ] #calculate and plot autocovariance manually
def autocovariance(x,lag):
    x_mean=np.mean(x)
    n=len(x)
    return np.sum((x[:n-lag]-x_mean)*(x[lag:]-x_mean))/n

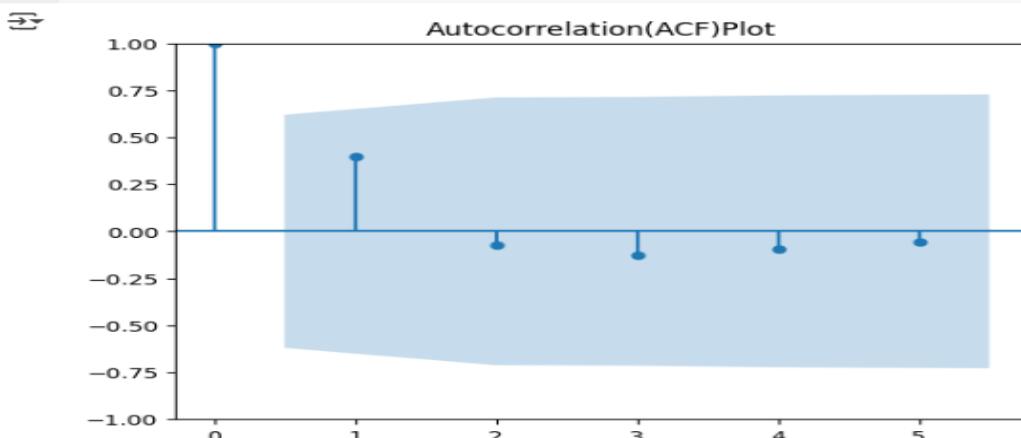
▶ #print autocovariance for lags 1 to 5
for lag in range(1,6):
    cov=autocovariance(df['Temperature'].values,lag)
    print(f"Autocovariance for lag {lag}: {cov:.4f}")

→ Autocovariance for lag 1: 0.052560
Autocovariance for lag 2: -0.009480
Autocovariance for lag 3: -0.016720
Autocovariance for lag 4: -0.012760
Autocovariance for lag 5: -0.007200
```

## Code:

```
#plot autocorrelation function (ACF)
from statsmodels.graphics.tsaplots import plot_acf
#ACF
plot_acf(df['Temperature'],lags=5)
plt.title("Autocorrelation(ACF) Plot")
plt.show()
```

## Output:



## Code:

```
#build and train on autoregression model
ar_model=AutoReg(df['Temperature'],lags=3)
ar_result=ar_model.fit()
#View coefficients
print(ar_result.summary())
```

## Output:

```
[ ] #build and train on autoregression model
ar_model=AutoReg(df['Temperature'],lags=3)
ar_result=ar_model.fit()

/usr/local/lib/python3.11/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)

#View coefficients
print(ar_result.summary())

AutoReg Model Results
=====
Dep. Variable: Temperature No. Observations: 10
Model: AutoReg(3) Log Likelihood: 2.177
Method: Conditional MLE S.D. of innovations 0.177
Date: Sat, 26 Jul 2025 AIC: 5.645
Time: 11:52:11 BIC: 5.375
Sample: 04-01-2024 HQIC: 2.303
- 10-01-2024
=====
coef std err z P>|z| [0.025 0.975]
-----
const 0.5402 6.676 0.081 0.936 -12.545 13.626
Temperature.L1 1.4472 0.325 4.451 0.000 0.810 2.084
Temperature.L2 -1.0229 0.413 -2.474 0.013 -1.833 -0.213
Temperature.L3 0.5470 0.307 1.781 0.075 -0.055 1.149
Roots
Real Imaginary Modulus Frequency
```

## Code:

```
#make prediction and plot
#for next 10 values
forecast=ar_result.predict(start=len(df),end=len(df)+9,dynamic=False)
print(forecast)
```

## Output:

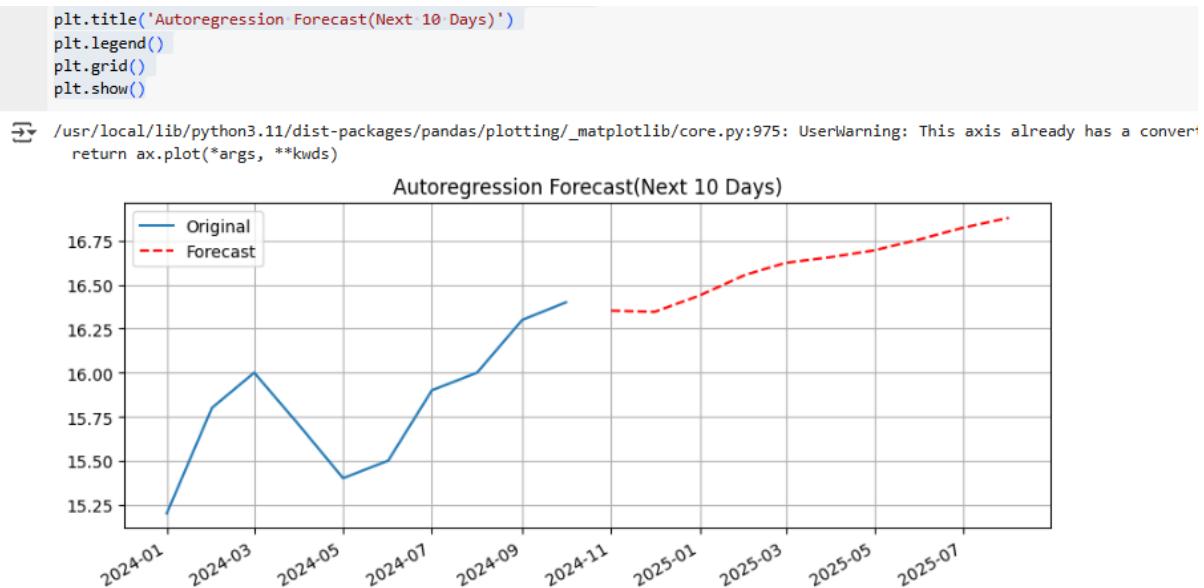
```
[ ] #make prediction and plot
#for next 10 values
forecast=ar_result.predict(start=len(df),end=len(df)+9,dynamic=False)
print(forecast)

2024-11-01 16.352744
2024-12-01 16.346167
2025-01-01 16.439687
2025-02-01 16.555907
2025-03-01 16.624839
2025-04-01 16.656870
2025-05-01 16.696286
2025-06-01 16.758271
2025-07-01 16.825176
2025-08-01 16.880156
Freq: MS, dtype: float64
/usr/local/lib/python3.11/dist-packages/statsmodels/tsa/deterministic.py:308: UserWarning: Only PeriodIndexes, DatetimeIndexes w
  fcst_index = self._extend_index(index, steps, forecast_index)
```

## Code:

```
#plot original data and forecast
plt.figure(figsize=(10, 4))
plt.plot(df['Temperature'], label='Original')
forecast.plot(label='Forecast', linestyle='--', color='red')
plt.title('Autoregression Forecast(Next 10 Days)')
plt.legend()
plt.grid()
plt.show()
```

## Output:



```
#program 1 project next2 years big billon days sales of flipkart using
ARMA(1)
```

## Code:

```
#Hypothetical dataset(2014-2023)
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

#load data
df=pd.read_csv("Hypodataset.csv")
#Create time series
sales = df['BBB Sales(in Crores)']
years = df['Year']
sales_series=pd.Series(data=sales.values,index=years)
print("Sales Data:\n",sales_series)
```

## Output:

```
[ ] #program 1 project next2 years big billon days sales of flipkart using ARMA(1)
#Hypothetical dataset(2014-2023)
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsplots import plot_acf, plot_pacf

#load data
df=pd.read_csv("Hypodataset.csv")

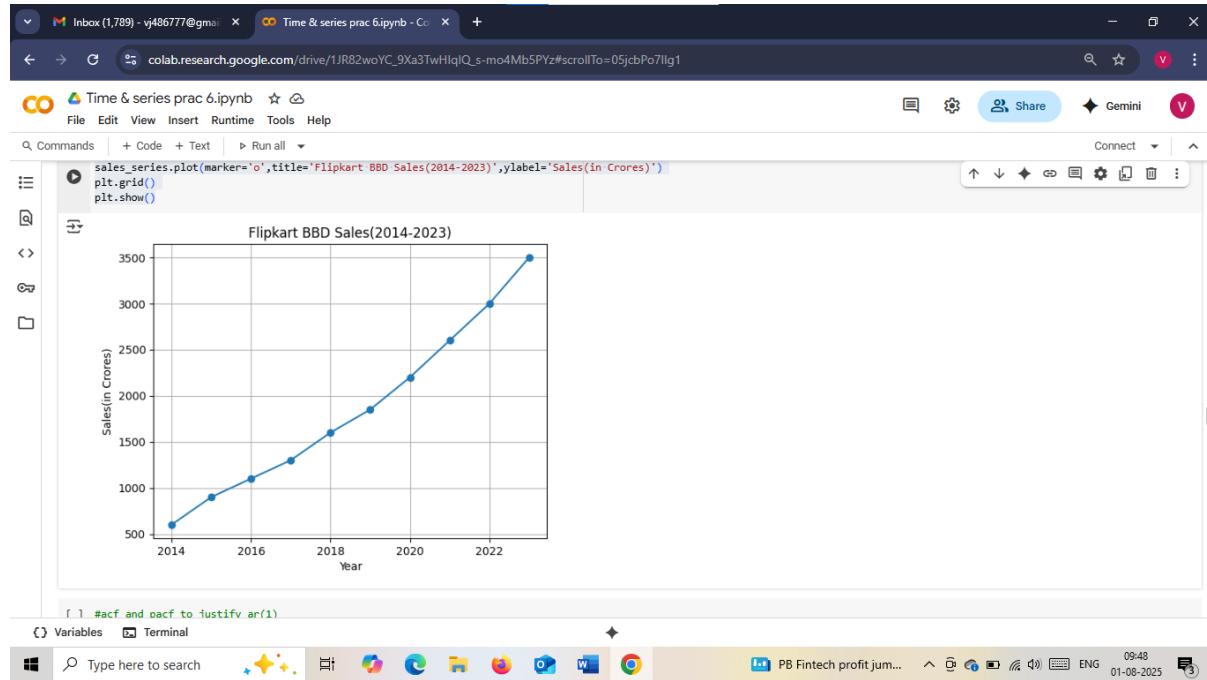
[ ] #Create time series
sales = df['BBD Sales(in Crores)']
years = df['Year']
sales_series=pd.Series(data=sales.values,index=years)
print("Sales Data:\n",sales_series)

[ ] Sales Data:
[ ] Year
[ ] 2014      600
[ ] 2015      900
[ ] 2016     1100
[ ] 2017     1300
[ ] 2018     1600
[ ] 2019     1850
[ ] 2020     2200
[ ] 2021     2600
```

## Code:

```
#plot the time series
sales_series.plot(marker='o',title='Flipkart BBD Sales(2014-2023)',ylabel='Sales(in Crores)')
plt.grid()
plt.show()
```

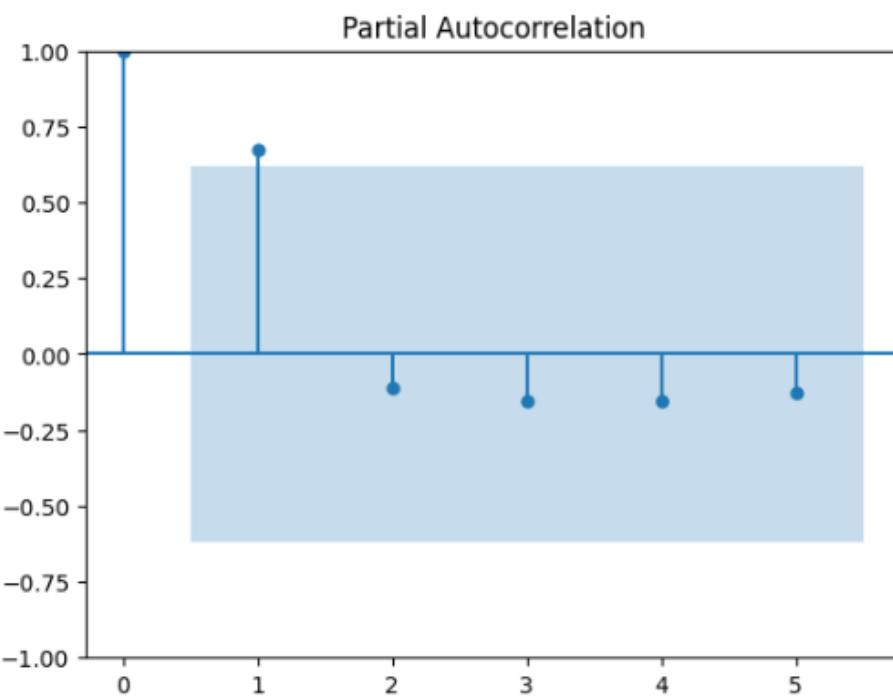
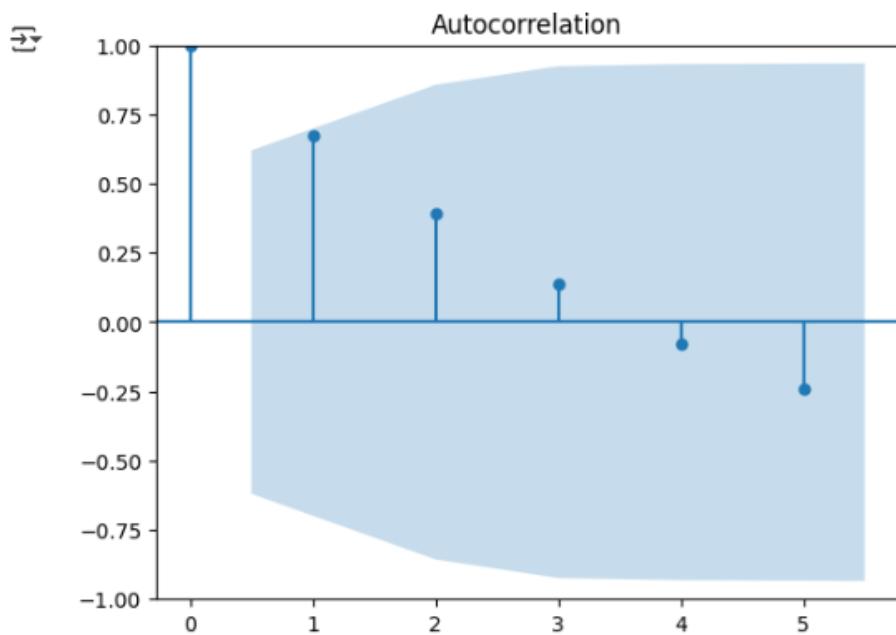
## Output:



## Code:

```
#acf and pacf to justify ar(1)
plot_acf(sales_series)
plt.show()
plot_pacf(sales_series)
plt.show()
```

## Output:



## Code:

```
#fit arima(1,0) (Arima with d=0, ma=0)
model=ARIMA(sales_series,order=(1,0,0))
result=model.fit()

#model summary
print(result.summary())
```

## Output:

The screenshot shows a Google Colab notebook titled "Time & series prac 6.ipynb". The code cell contains the provided Python script. The output cell displays the results of the model fitting and summary printing. It includes several warning messages from the statsmodels library about non-stationary data and starting parameters. The summary table shows the model details and coefficient statistics.

```
#fit arima(1,0) (Arima with d=0, ma=0)
model=ARIMA(sales_series,order=(1,0,0))
result=model.fit()

#model summary
print(result.summary())

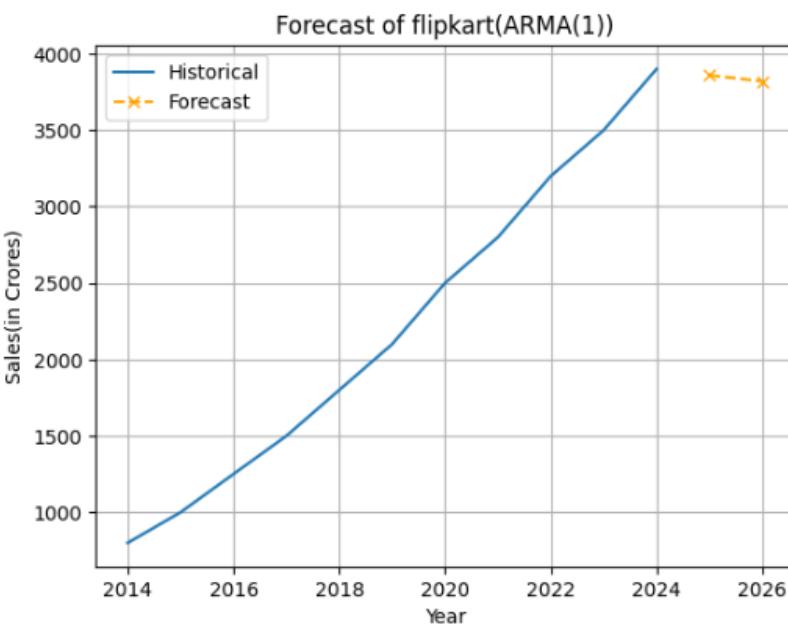
/.../statsmodels/tsa/base/tsa_model.py:473: ValueWarning: An unsupported index was provided. As a result, forecasts cannot be generated. To use the self._init_dates(dates, freq)
/.../statsmodels/tsa/base/tsa_model.py:473: ValueWarning: An unsupported index was provided. As a result, forecasts cannot be generated. To use the self._init_dates(dates, freq)
/.../statsmodels/tsa/base/tsa_model.py:473: ValueWarning: An unsupported index was provided. As a result, forecasts cannot be generated. To use the self._init_dates(dates, freq)
/.../statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters
SARIMAX Results
-----
Dep. Variable: y No. Observations: 10
Model: ARIMA(1, 0, 0) Log Likelihood -73.709
Date: Mon, 28 Jul 2025 AIC 153.419
Time: 10:38:29 BIC 154.326
Sample: 0 HQIC 152.423
- 10
Covariance Type: opg
-----
coef std err z P>|z| [0.025 0.975]
-----
```

## Code:

```
forecast=result.forecast(steps=2)
forecast_years=pd.date_range(start='2024',periods=2,freq='YE')
plt.plot(sales_series,label='Historical')
plt.plot(forecast_years,forecast,label='Forecast',marker='x',linestyle='--',color='orange')
plt.title('Forecast of flipkart(ARMA(1))')
plt.ylabel('Sales(in Crores)')
plt.xlabel('Year')
plt.legend()
plt.grid()
plt.show()

print("Forecasted Sales for 2024:", forecast.iloc[0])
print("Forecasted Sales for 2025:", forecast.iloc[1])
```

## Output:



```
▶ print("Forecasted Sales for 2024:", forecast.iloc[0])
print("Forecasted Sales for 2025:", forecast.iloc[1])
```

```
→ Forecasted Sales for 2024: 3859.5327484533027
Forecasted Sales for 2025: 3820.0366142882617
```

## #Next dataset

### Code:

```
#next dataset
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
#create the time series
years=pd.date_range(start='2013',periods=11,freq='Y')
sales=[800,1000,1250,1500,1800,2100,2500,2800,3200,3500,3900]
sales_series=pd.Series(data=sales,index=years)
print("Sales Data:\n",sales_series)
```

## Output:

The screenshot shows a Google Colab notebook interface. The code cell contains the following Python script:

```
#next: dataset
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

#create the time series

```
years=pd.date_range(start='2013',periods=11,freq='Y')
sales=[800,1000,1250,1500,1800,2100,2500,2800,3200,3500,3900]
sales_series=pd.Series(data=sales,index=years)
print("Sales Data:\n",sales_series)
```

The output cell displays the generated time series data:

Date	Sales
2013-12-31	800
2014-12-31	1000
2015-12-31	1250
2016-12-31	1500
2017-12-31	1800
2018-12-31	2100
2019-12-31	2500
2020-12-31	2800
2021-12-31	3200
2022-12-31	3500
2023-12-31	3900

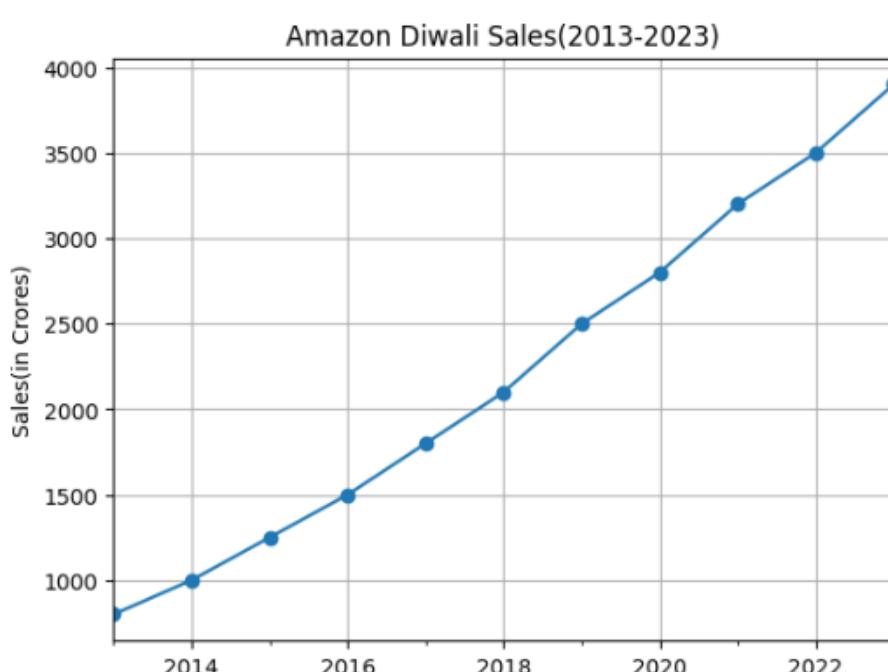
Freq: YE-DEC, dtype: int64

The status bar at the bottom indicates "Very humid" and the date "01-08-2025".

## Code:

```
#plot the original sales data
sales_series.plot(marker='o',title='Amazon Diwali Sales(2013-2023)',ylabel='Sales(in Crores)')
plt.grid()
plt.show()
```

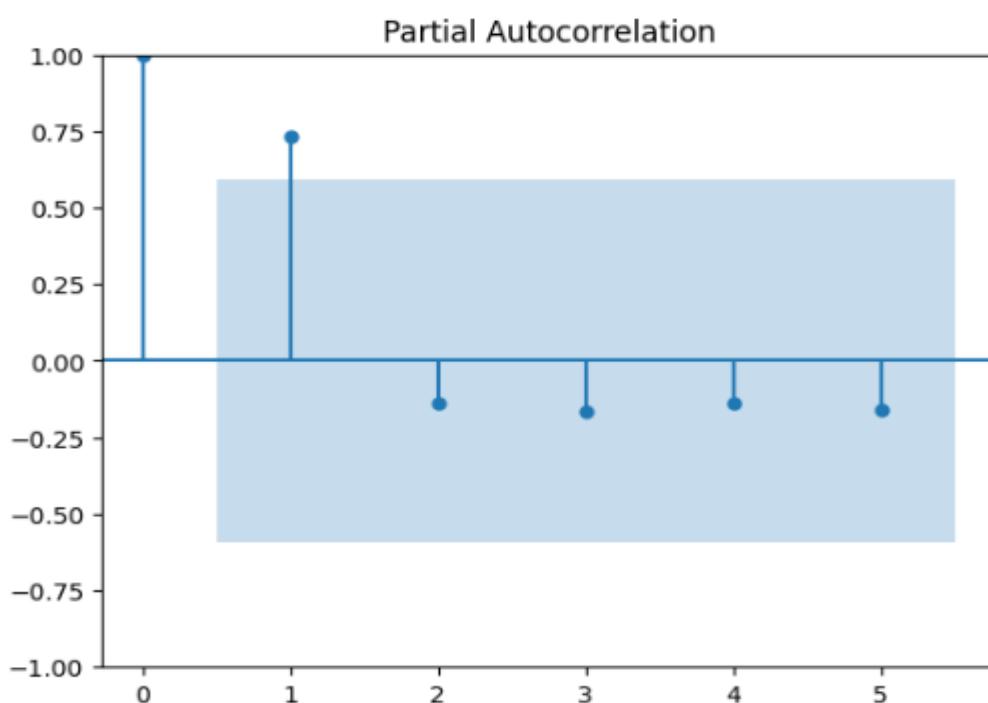
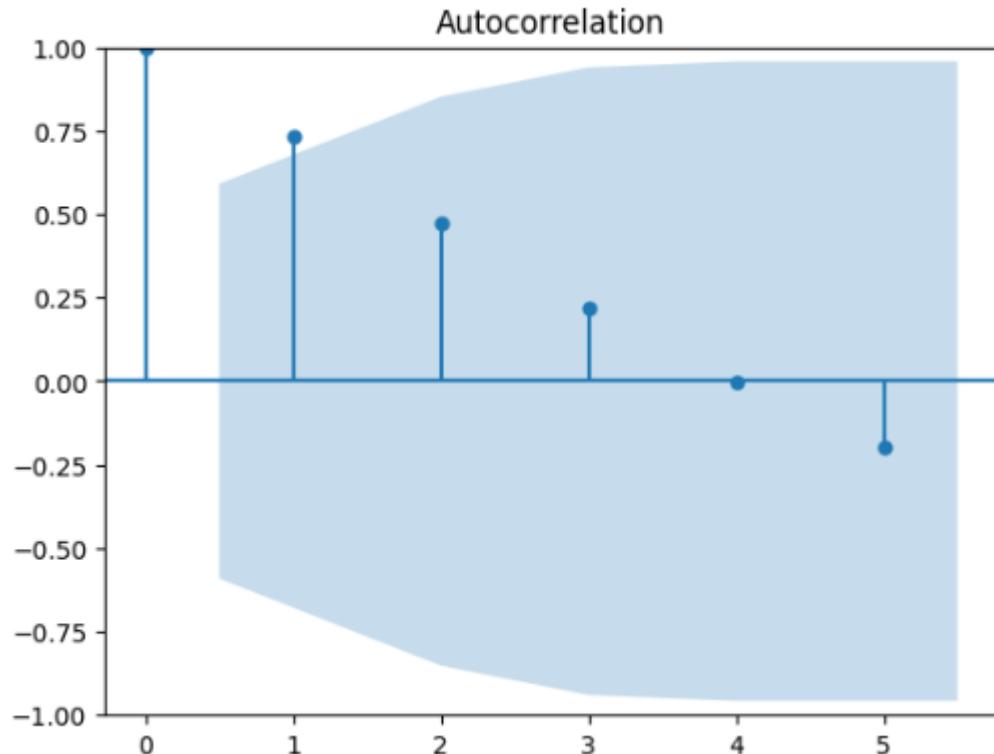
## Output:



## Code:

```
#acf and pacf  
plot_acf(sales_series)  
plot_pacf(sales_series)  
plt.show()
```

## Output:



## Code:

```
#fit arma(1,0)
model=ARIMA(sales_series,order=(1,0,0))
result=model.fit()
#model summary
print(result.summary())
```

## Output:

```
[ ] #fit arma(1,0)
model=ARIMA(sales_series,order=(1,0,0))
result=model.fit()

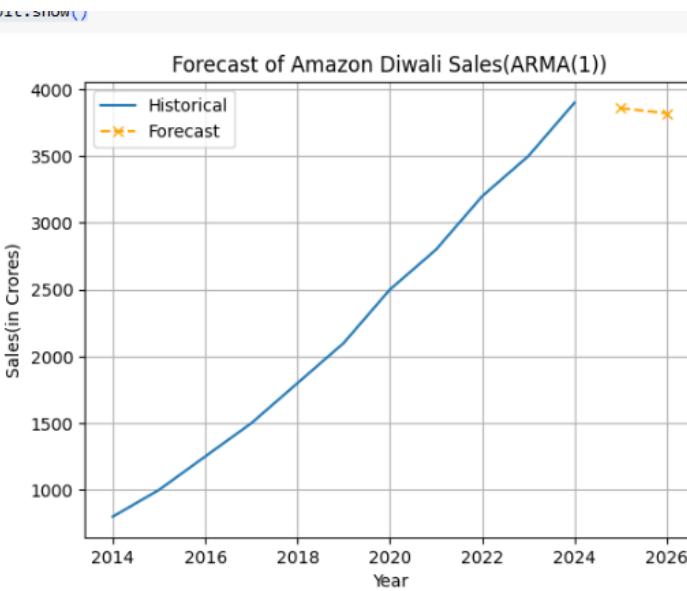
▶ #model summary
print(result.summary())

SARIMAX Results
=====
Dep. Variable:                      y   No. Observations:                 11
Model:                          ARIMA(1, 0, 0)   Log Likelihood:            -80.457
Date:                Mon, 28 Jul 2025   AIC:                         166.914
Time:                      11:08:04      BIC:                         168.108
Sample:                12-31-2013   HQIC:                        166.161
                           - 12-31-2023
Covariance Type:                  opg
=====
            coef    std err        z     P>|z|      [0.025    0.975]
-----
const    2213.6967   1445.975     1.531     0.126    -620.362    5047.756
ar.L1      0.9760     0.201     4.863     0.000      0.583     1.369
sigma2    1.003e+05   2.06e+05     0.487     0.626   -3.03e+05    5.04e+05
=====
```

## Code:

```
#forecast next 2 years
forecast=result.forecast(steps=2)
forecast_years=pd.date_range(start='2024',periods=2,freq='YE')
#plot historical and forecasted values
plt.plot(sales_series,label='Historical')
plt.plot(forecast_years,forecast,label='Forecast',marker='x',linestyle='--',color='orange')
plt.title('Forecast of Amazon Diwali Sales (ARMA(1))')
plt.ylabel('Sales(in Crores)')
plt.xlabel('Year')
plt.legend()
plt.grid()
plt.show()
print("Forecasted Amazon Diwali Sales:")
for year,value in zip(forecast_years,forecast):
    print(f'{year}: {value:.2f}Crores')
```

## Output:



```
print("Forecasted Amazon Diwali Sales:")
for year,value in zip(forecast_years,forecast):
    print(f'{year}: {value:.2f}Crores')
```

Forecasted Amazon Diwali Sales:  
2024-12-31 00:00:00: 3859.53Crores  
2025-12-31 00:00:00: 3820.04Crores

## #Same program in r-studio

### Code:

```
#load required libraries in r
library(forecast)
library(tseries)

#create time series object(from 2013,frequency=1 sice yearly data)
sales_ts<-
ts(c(750,1000,1300,1650,1900,2200,2550,2900,3300,3700,4100),start
=2013,frequency=1)
#plot the time series
plot(sales_ts,main="Amazon Great Indian Festival Sales(2013-
2023)",ylab="Sales(Crores)",xlab="Year",type="o",col="blue")
```

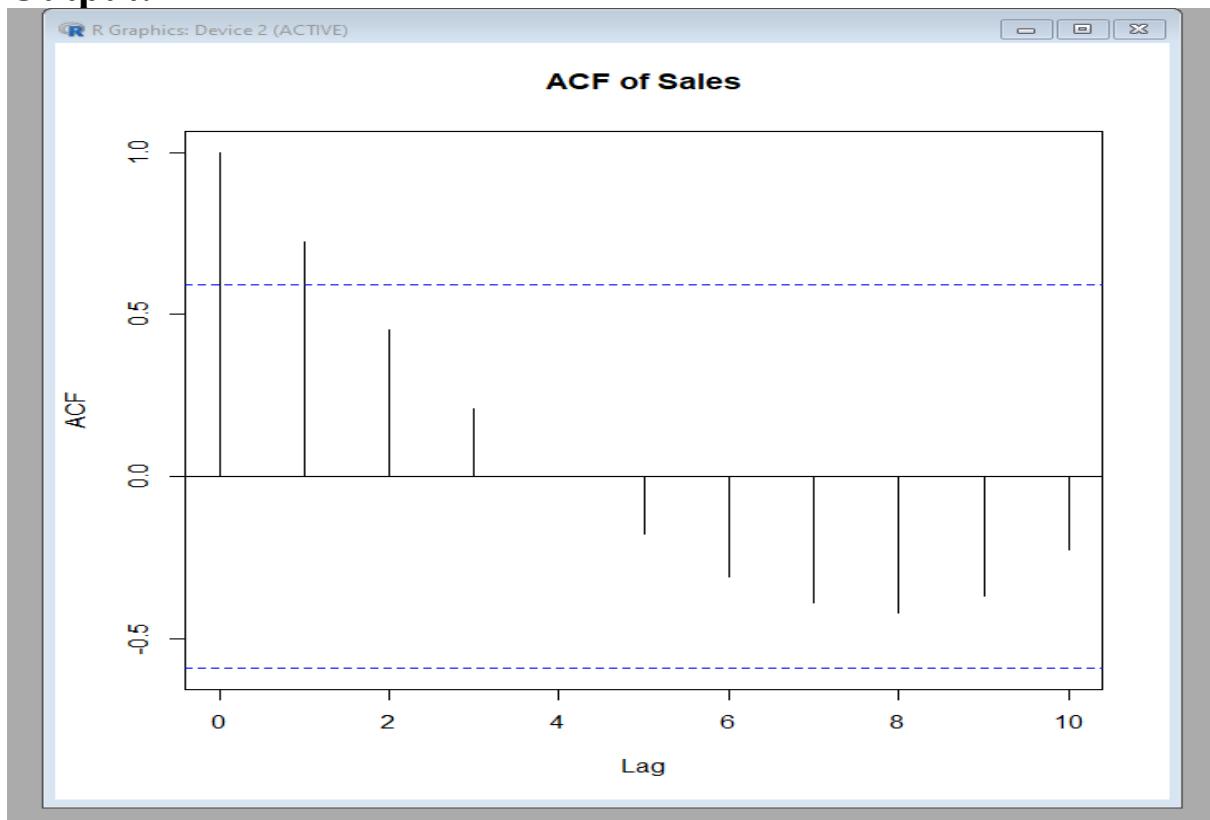
## Output:



## Code:

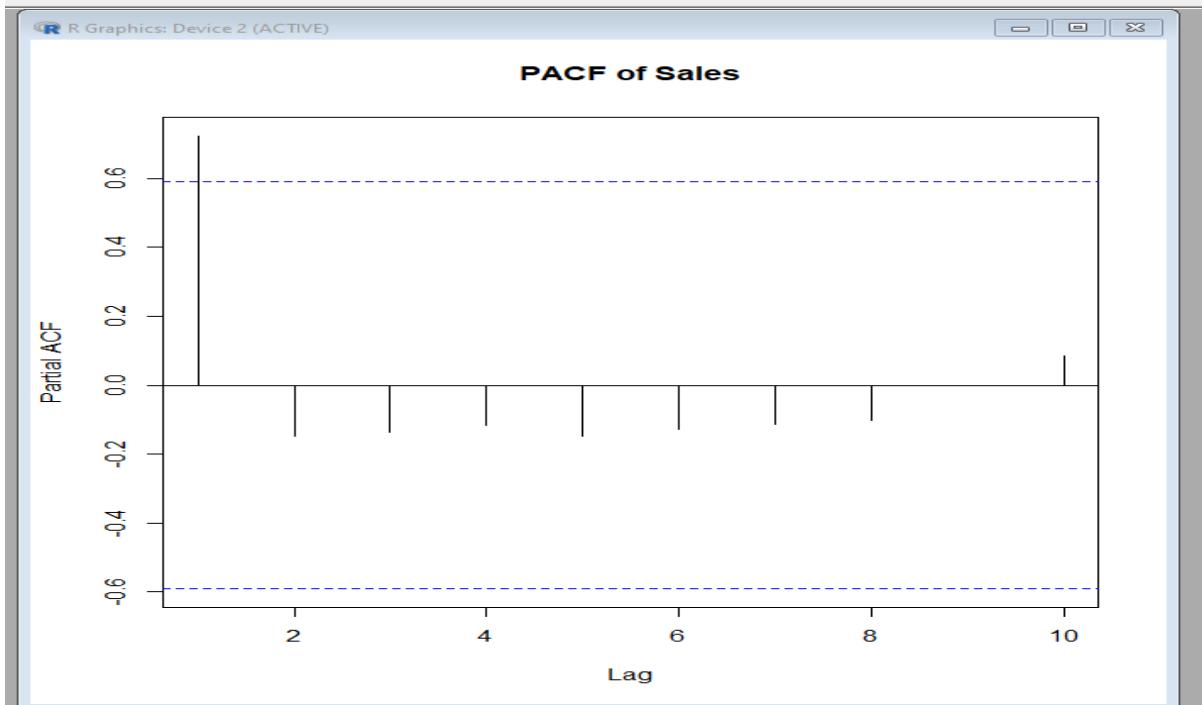
```
#check acf and pacf  
acf(sales_ts,main="ACF of Sales")
```

## Output:



**Code:**

```
pacf(sales_ts,main="PACF of Sales")
```

**Output:****Code:**

```
arma_model<-Arima(sales_ts,order=c(1,0,1))
summary(arma_model)
```

**Output:**

```
> summary(arma_model)
Series: sales_ts
ARIMA(1,0,1) with non-zero mean

Coefficients:
            ar1      ma1      mean
            0.9995  0.4353  3161.804
            s.e.  0.0055  0.3880  12442.605

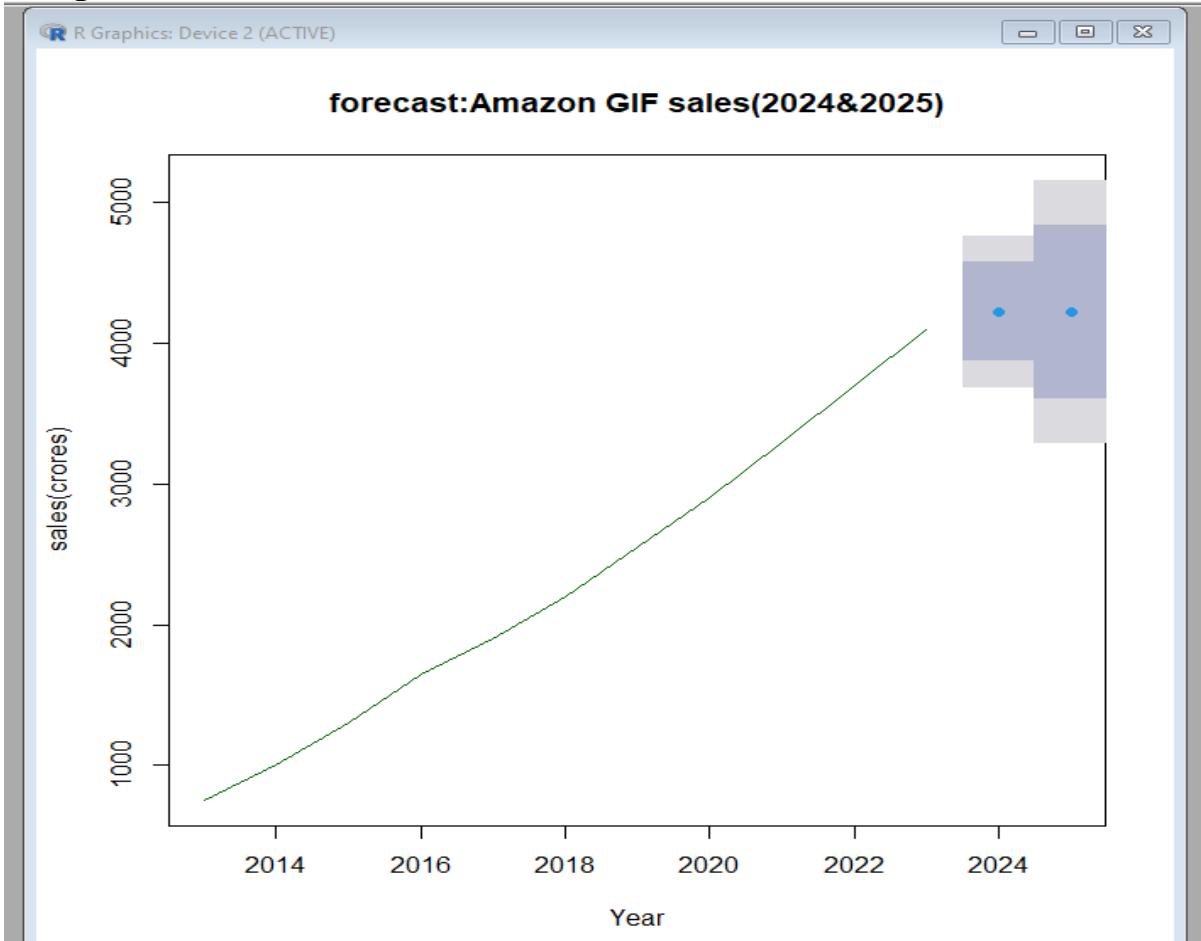
sigma^2 = 75024:  log likelihood = -79.48
AIC=166.95    AICc=173.62    BIC=168.54

Training set error measures:
          ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 213.668 233.5873 223.7415 9.653652 10.99679 0.6678851 0.02578363
```

**Code:**

```
forecast_result<-forecast(arma_model,h=2)
#plot the forecast
plot(forecast_result,main="forecast:Amazon GIF
sales(2024&2025)",xlab="Year",ylab="sales(crores)"
,col="darkgreen")
```

## Output:



## Code:

```
print(forecast_result)
```

## Output:

```
>  
> print(forecast_result)  
  Point Forecast    Lo 80     Hi 80    Lo 95     Hi 95  
2024      4222.182 3871.158 4573.206 3685.338 4759.027  
2025      4221.606 3607.715 4835.497 3282.741 5160.471  
> |
```

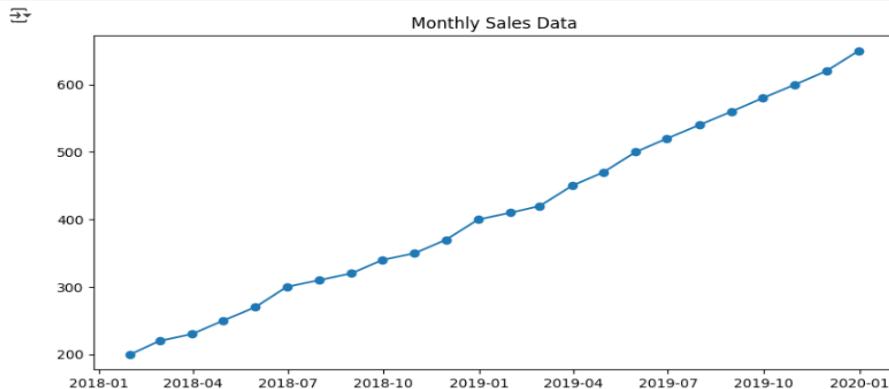
## PRACTICAL NO -7

**Aim:** Using Arima for forecasting

**Code:**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
from statsmodels.tsa.arima.model import ARIMA
#load dataset
#monthly sales
data=pd.DataFrame({'Month':pd.date_range(start='2018-01',periods=24,freq='ME'),'Sales':[200,220,230,250,270,300,310,320,340,350,370,400,410,420,450,470,500,520,540,560,580,600,620,650]})
data.set_index('Month',inplace=True)
#plot data
plt.figure(figsize=(10,5))
plt.plot(data['Sales'],marker='o')
plt.title("Monthly Sales Data")
plt.show()
```

**Output:**



**Code:**

```
#Check Stationary(ADF test)
def check_stationarity(series):
    result=adfuller(series)
    print("ADF Statistic:", result[0])
    print("p-value:",result[1])
    if result[1]<=0.05:
        print("Stationary")
    else:
        print("Non-Stationary")

check_stationarity(data['Sales'])
```

## Output:

```
✓ 0s check_stationarity(data['Sales'])
```

```
→ ADF Statistic: 1.3551050643521871  
p-value: 0.9969003347519994  
Non-Stationary
```

## Code:

```
#Difference to stationarity  
data_diff=data['Sales'].diff().dropna()  
check_stationarity(data_diff)
```

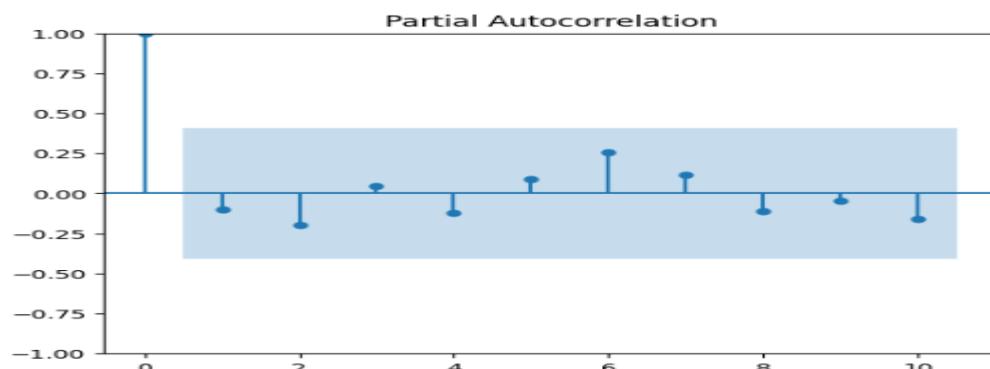
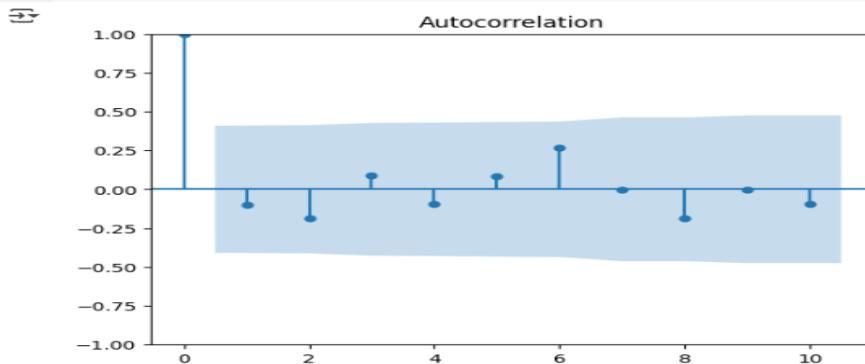
## Output:

```
→ ADF Statistic: -3.8658943604089964  
p-value: 0.0023006870324156995  
Stationary
```

## Code:

```
plot_acf(data_diff, lags=10)    #ARIMA(1,1,1) from acf and pacf  
plot_pacf(data_diff, lags=10)  
plt.show()
```

## Output:



## Code:

```
#fit ARIMA Model
model=ARIMA(data['Sales'],order=(1,1,1))
model_fit=model.fit()
print(model_fit.summary())
```

## Output:

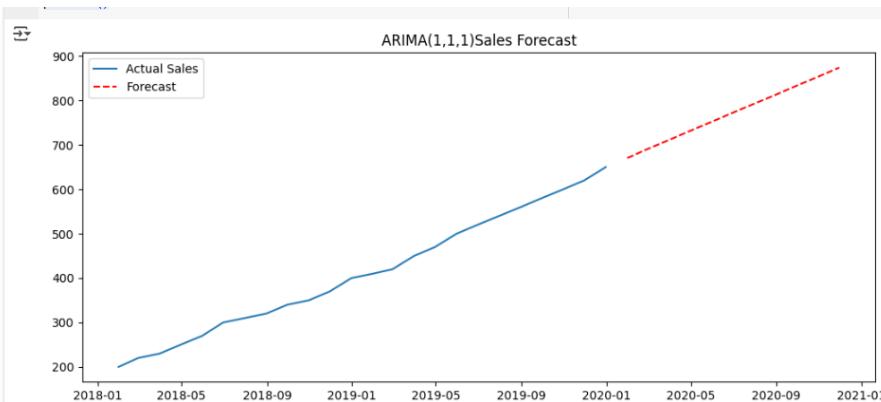
```
print(model_fit.summary())
SARIMAX Results
=====
Dep. Variable: Sales No. Observations: 24
Model: ARIMA(1, 1, 1) Log Likelihood -80.102
Date: Wed, 06 Aug 2025 AIC 166.203
Time: 18:37:38 BIC 169.610
Sample: 01-31-2018 HQIC 167.060
- 12-31-2019
Covariance Type: opg
=====
coef std err z P>|z| [0.025 0.975]
ar.L1 0.9995 0.005 203.702 0.000 0.990 1.009
ma.L1 -0.9137 0.355 -2.571 0.010 -1.610 -0.217
sigma2 52.6164 23.772 2.213 0.027 6.025 99.208
=====
Ljung-Box (L1) (Q): 0.86 Jarque-Bera (JB): 0.96
Prob(Q): 0.35 Prob(JB): 0.62
Heteroskedasticity (H): 0.60 Skew: 0.06
Prob(H) (two-sided): 0.49 Kurtosis: 2.01
=====
```

## Code:

```
#forecast future Sales
forecast = model_fit.forecast(steps=12)
future_index=pd.date_range(start=data.index[-1],periods=12,freq='ME')
forecast_df=pd.DataFrame({'Forecast':forecast},index=future_index)

plt.figure(figsize=(12,5))
plt.plot(data['Sales'],label='Actual Sales')
plt.plot(forecast_df['Forecast'],label='Forecast',color='red',linestyle='--')
plt.title("ARIMA(1,1,1) Sales Forecast")
plt.legend()
plt.show()
```

## Output:



## Code:

```
#forecast Airline Passengers
url="https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv"
df=pd.read_csv(url, parse_dates=['Month'], index_col='Month')
df.head()
```

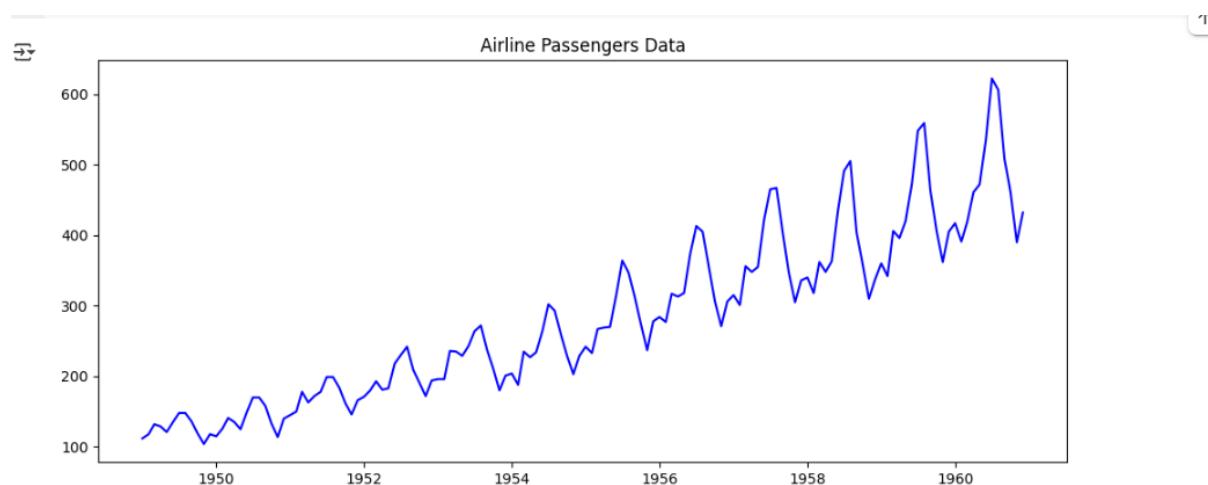
## Output:



## Code:

```
#plot Series
plt.figure(figsize=(12,5))
plt.plot(df['Passengers'],color='blue')
plt.title("Airline Passengers Data")
plt.show()
```

## Output:



## Code:

```
#fit Arima model
model=ARIMA(df['Passengers'],order=(2,1,2))
model_fit=model.fit()
print(model_fit.summary())
```

## Output:

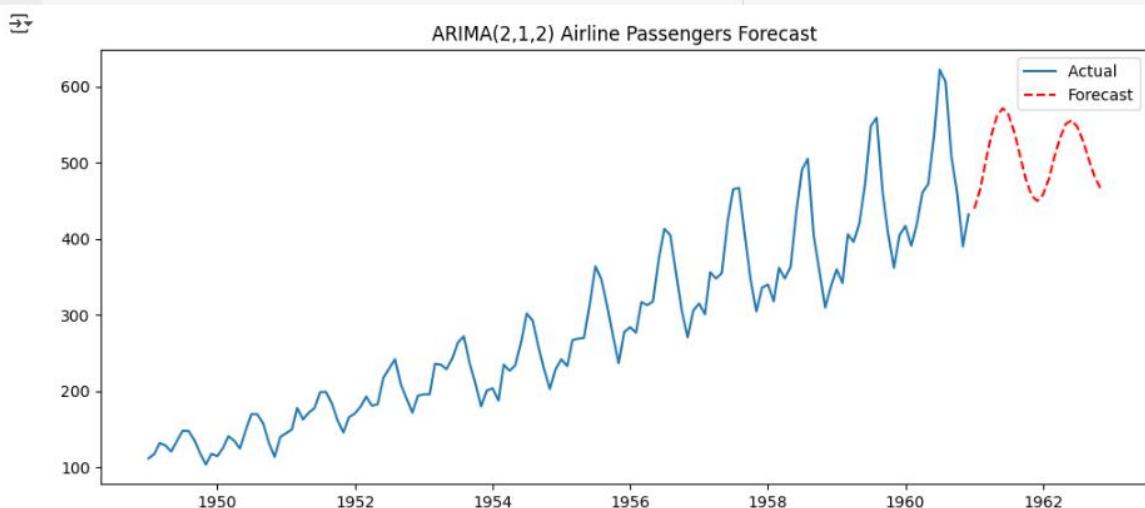
```
self._init_dates(dates, freq)
                                SARIMAX Results
=====
Dep. Variable:          Passengers    No. Observations:                  144
Model:                 ARIMA(2, 1, 2)    Log Likelihood:                -671.673
Date:                 Wed, 06 Aug 2025   AIC:                         1353.347
Time:                     18:37:39      BIC:                         1368.161
Sample:                01-01-1949    HQIC:                        1359.366
                           - 12-01-1960
Covariance Type:            opg
=====
            coef    std err        z     P>|z|      [0.025     0.975]
-----
ar.L1      1.6850    0.020    83.060    0.000      1.645     1.725
ar.L2     -0.9548    0.017   -55.420    0.000     -0.989    -0.921
ma.L1     -1.8432    0.124   -14.814    0.000     -2.087    -1.599
ma.L2      0.9953    0.135     7.382    0.000      0.731     1.260
sigma2    665.9600  114.029     5.840    0.000    442.467    889.453
=====
Ljung-Box (L1) (Q):        0.30    Jarque-Bera (JB):       1.84
```

## Code:

```
#forecast
forecast = model_fit.forecast(steps=24)
future_index=pd.date_range(start=df.index[-1], periods=24, freq='MS')
forecast_df=pd.DataFrame({'Forecast':forecast}, index=future_index)

plt.figure(figsize=(12,5))
plt.plot(df['Passengers'],label='Actual ')
plt.plot(forecast_df['Forecast'],label='Forecast',color='red',linestyle='--')
plt.title("ARIMA(2,1,2) Airline Passengers Forecast")
plt.legend()
plt.show()
```

## Output:

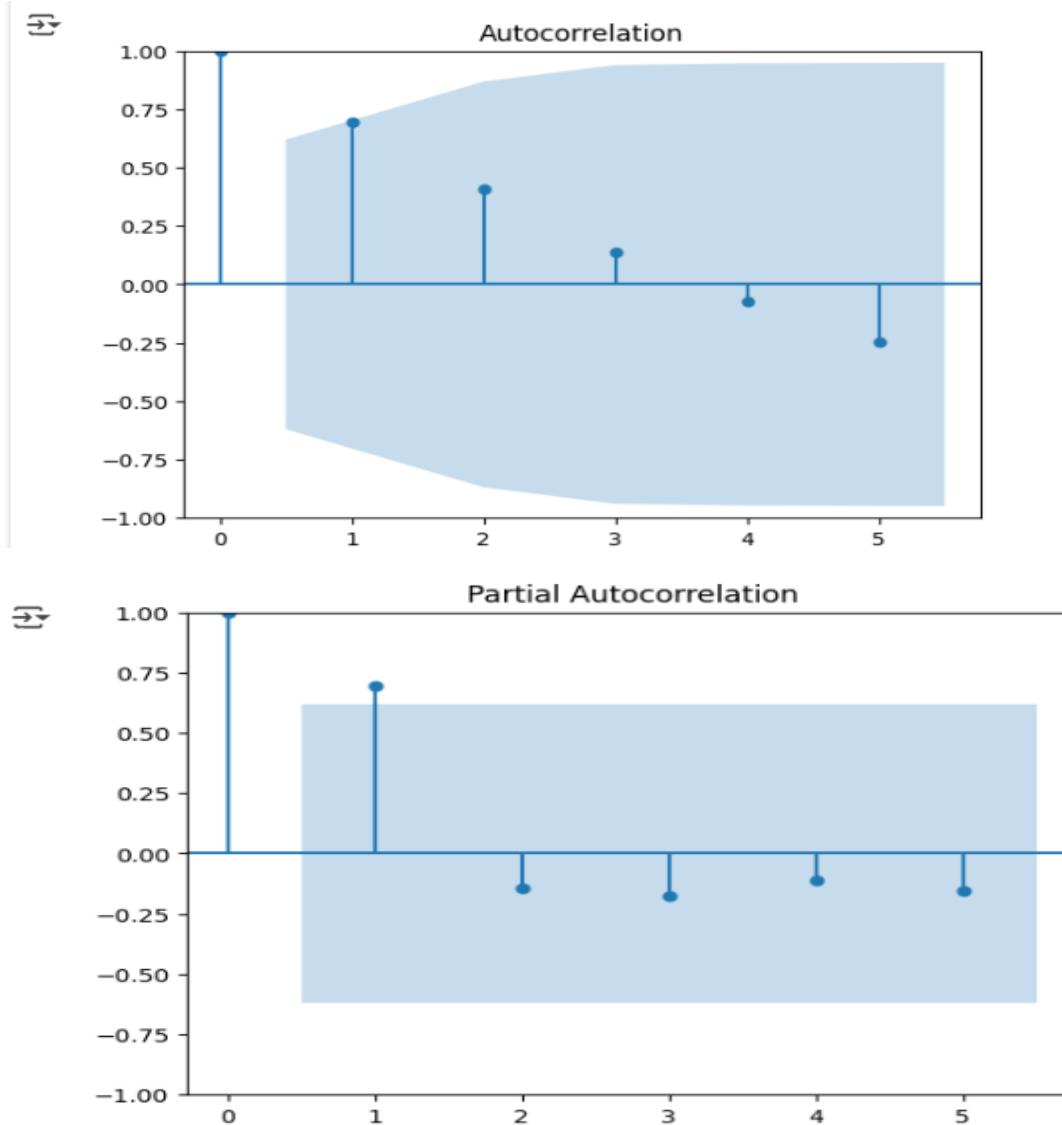


## #ARMA

### Code:

```
#ARMA model forecast (2,0)
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
#load series
data=[100,104,107,111,115,118,120,124,127,130]
series=pd.Series(data)
#ACF and PACF plots
plot_acf(series, lags=5)
plt.show()
plot_pacf(series, lags=5)
plt.show()
```

### Output:



## Code:

```
#fit ARMA(2,0) model
model=ARIMA(series,order=(2,0,0))
model_fit=model.fit()
print(model_fit.summary())
```

## Output:

```
SARIMAX Results
=====
Dep. Variable:      y    No. Observations:      10
Model:             ARIMA(2, 0, 0)   Log Likelihood: -18.670
Date:       Wed, 06 Aug 2025   AIC:                 45.340
Time:           18:37:39   BIC:                 46.550
Sample:          0 - 10   HQIC:                44.012
Covariance Type: opg
=====
coef      std err      z      P>|z|      [0.025      0.975]
-----
const    112.5162    39.524     2.847     0.004     35.050    189.983
ar.L1     1.9128     0.262     7.291     0.000     1.399     2.427
ar.L2    -0.9369     0.233    -4.025     0.000    -1.393    -0.481
sigma2    1.1114     0.832     1.336     0.182    -0.520     2.742
=====
Ljung-Box (L1) (Q):        4.22    Jarque-Bera (JB):      1.03
Prob(Q):                  0.04    Prob(JB):            0.60
Heteroskedasticity (H):    1.83    Skew:                 0.69
Prob(H) (two-sided):      0.63    Kurtosis:            2.25
=====
Warnings:
[...]
```

## Code:

```
#forecast next 3 values
forecast=model_fit.forecast(steps=3)
print("\nforecast for next 3 days:",forecast.tolist())
```

## Output:

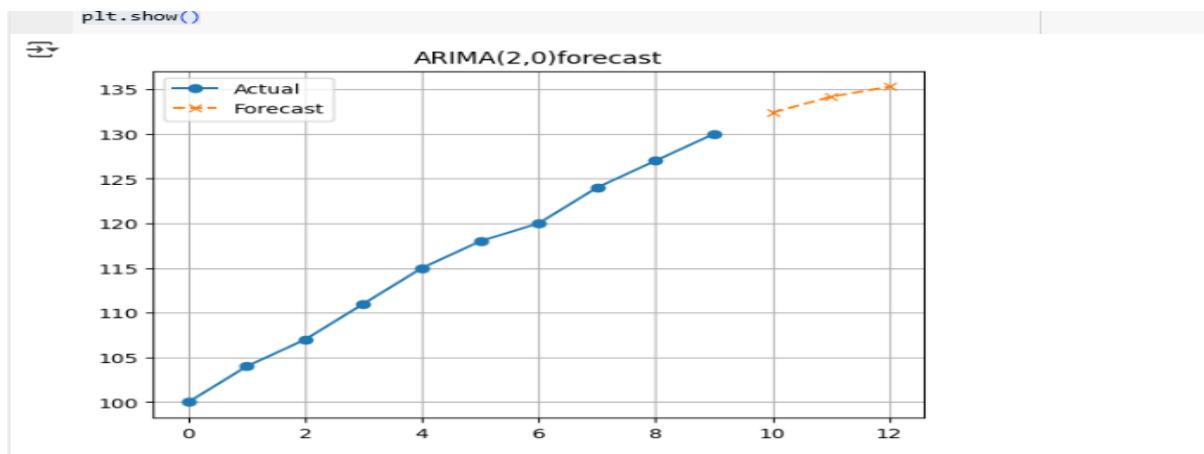
```
#forecast next 3 values
forecast=model_fit.forecast(steps=3)
print("\nforecast for next 3 days:",forecast.tolist())
```

```
forecast for next 3 days: [132.3893208087086, 134.1488641347984, 135.27591924000663]
```

## Code:

```
#plot
plt.plot(series,label="Actual",marker='o')
plt.plot(range(len(series),len(series)+3),forecast,label="Forecast",marker='x',linestyle='--')
plt.title("ARIMA(2,0) forecast")
plt.legend()
plt.grid()
plt.show()
```

## Output:



## Code:

```
pip install pandas numpy matplotlib statsmodels pmdarima

#import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
#from pmdarima import auto_arima
#create
np.random.seed(42)
months=pd.date_range(start='2018-01-01',periods=60,freq='ME')
sales=200+np.arange(60)*2+np.random.normal(0,10,60) #upward trend +
noise
data=pd.DataFrame({ 'Month':months, 'Sales':sales})
data.set_index('Month',inplace=True)
print(data.head())
```

## Output:

```
print(data.head())

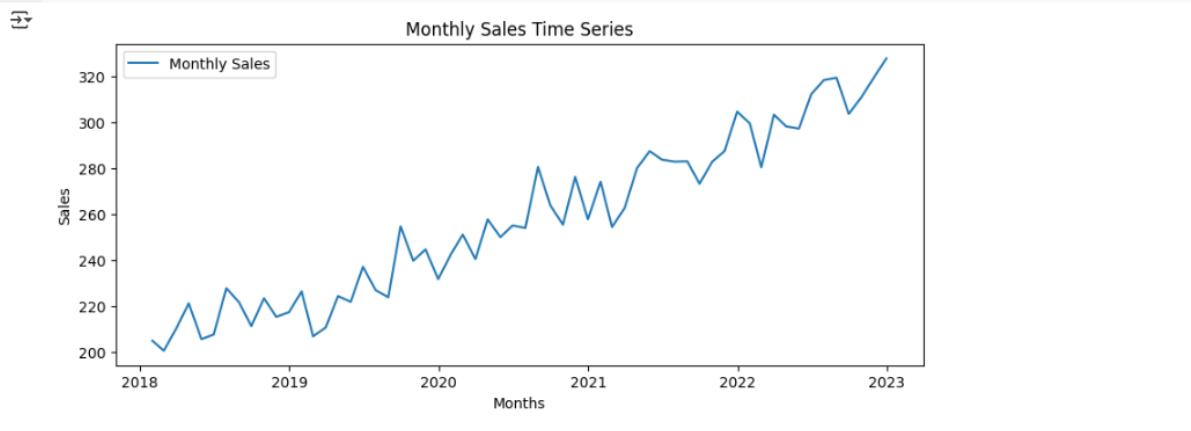
```

Month	Sales
2018-01-31	204.967142
2018-02-28	200.617357
2018-03-31	210.476885
2018-04-30	221.230299
2018-05-31	205.658466

## Code:

```
#visualize the data
plt.figure(figsize=(10,4))
plt.plot(data['Sales'],label='Monthly Sales')
plt.title("Monthly Sales Time Series")
plt.xlabel("Months")
plt.ylabel("Sales")
plt.legend()
plt.show()
```

## Output:



## Code:

```
#check stationarity
result=adfuller(data['Sales'])
print("ADF Statistic:",result[0])
print("p-value:",result[1])
if result[1]<=0.05:
    print("Data is Stationary")
else:
    print("Data is Non-Stationary")
model=ARIMA(data['Sales'],order=(1,1,1))
model_fit=model.fit()
print(model_fit.summary())
model_fit.plot_diagnostics(figsize=(12,6))
plt.show()
```

## Output:

```
ADF Statistic: 1.2914496798697157
p-value: 0.9965606304178265
Data is Non-Stationary
```

```

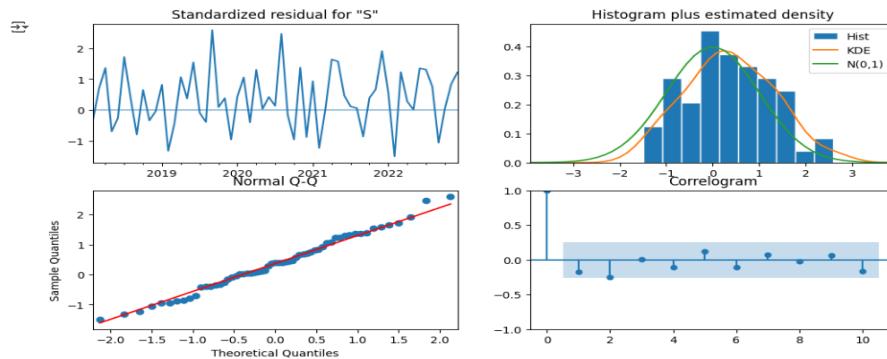
SARIMAX Results
=====
Dep. Variable: Sales No. Observations: 60
Model: ARIMA(1, 1, 1) Log Likelihood: -225.403
Date: Wed, 06 Aug 2025 AIC: 456.805
Time: 18:37:50 BIC: 463.038
Sample: 01-31-2018 HQIC: 459.238
- 12-31-2022
Covariance Type: opg
=====

coef std err z P>|z| [0.025 0.975]
ar.L1 -0.0771 0.293 -0.264 0.792 -0.650 0.496
ma.L1 -0.4760 0.260 -1.828 0.068 -0.986 0.034
sigma2 121.1655 24.103 5.027 0.000 73.924 168.407
=====

Ljung-Box (L1) (Q): 1.83 Jarque-Bera (JB): 0.67
Prob(Q): 0.18 Prob(JB): 0.72
Heteroskedasticity (H): 0.95 Skew: 0.14
Prob(H) (two-sided): 0.90 Kurtosis: 2.56
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```



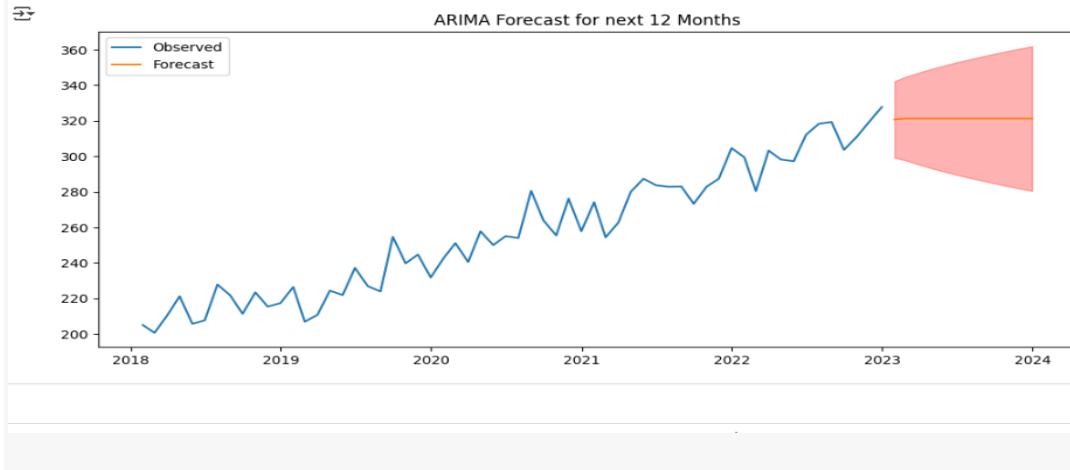
## Code:

```

forecast = model_fit.get_forecast(steps=12)
forecast_df=forecast.summary_frame()
#plot
plt.figure(figsize=(12,5))
plt.plot(data['Sales'],label='Observed')
plt.plot(forecast_df['mean'],label='Forecast')
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],color='red',alpha=0.3)
plt.title("ARIMA Forecast for next 12 Months")
plt.legend()
plt.show()

```

## Output:



## #R-studio code

### Code:

```
#forecast monthly sales usingg arima model  
install.packages("forecast")  
install.packages("tseries")  
library(forecast)  
library(tseries)  
#create or load forecast  
sales<-  
c(200,220,230,250,270,300,310,320,340,350,370,400,410,420,450,4  
70,500,520,540,560,580,600,620,650)  
#convert numeric vector to time series  
sales_ts<-ts(sales,start=c(2022,1),frequency=12)
```

### Output:

```
> #forecast monthly sales usingg arima model  
> install.packages("forecast")  
Installing package into 'C:/Users/Admin/AppData/Local/R/win-library/4.5'  
(as 'lib' is unspecified)  
--- Please select a CRAN mirror for use in this session ---  
Warning: failed to download mirror file (Cannot open URL 'https://cran.r-project.org/CRAN_mirrors.csv'); using local file 'C:/PROGRA~1/R/R-45~1.1/doc/CRAN_mirrors.csv'  
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/forecast_8.24.0.zip'  
Content type 'application/zip' length 1920158 bytes (1.8 MB)  
downloaded 1.8 MB  
  
package 'forecast' successfully unpacked and MD5 sums checked  
  
The downloaded binary packages are in  
  C:/Users/Admin/AppData/Local/Temp/RtmpQxs6dM/downloaded_packages  
Warning message:  
In download.file(url, destfile = f, quiet = TRUE) :  
  URL 'https://cran.r-project.org/CRAN_mirrors.csv': status was 'SSL connect error'  
> install.packages("tseries")  
Installing package into 'C:/Users/Admin/AppData/Local/R/win-library/4.5'  
(as 'lib' is unspecified)  
trying URL 'https://cloud.r-project.org/bin/windows/contrib/4.5/tseries_0.10-58.zip'  
Content type 'application/zip' length 387008 bytes (377 KB)  
downloaded 377 KB  
  
package 'tseries' successfully unpacked and MD5 sums checked  
  
The downloaded binary packages are in  
  C:/Users/Admin/AppData/Local/Temp/RtmpQxs6dM/downloaded_packages  
> library(forecast)  
Registered S3 method overwritten by 'quantmod':  
  method            from  
  as.zoo.data.frame zoo  
> library(tseries)  
  
'tseries' version: 0.10-58  
  
'tseries' is a package for time series analysis and computational  
  finance.  
  
See 'library(help="tseries")' for details.  
> #create or load forecast  
> sales<-c(200,220,230,250,270,300,310,320,340,350,370,400,410,420,450,470,500,520,540,560,580,600,620,650)  
> #convert numeric vector to time series  
> sales_ts<-ts(sales,start=c(2022,1),frequency=12)
```

### Code:

```
#display the first value  
print(sales_ts)
```

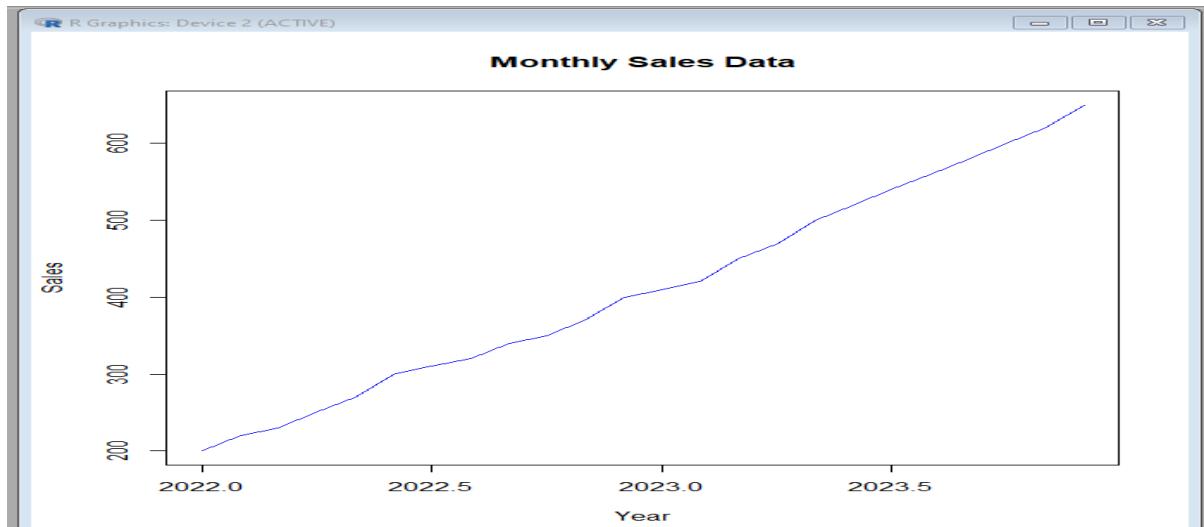
## Output:

```
> #display the first value  
> print(sales_ts)  
   Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec  
2022 200 220 230 250 270 300 310 320 340 350 370 400  
2023 410 420 450 470 500 520 540 560 580 600 620 650
```

## Code:

```
#plot the series  
  
plot(sales_ts,main="Monthly Sales  
Data",xlab="Year",ylab="Sales",col="blue")
```

## Output:



## Code:

```
#adf  
  
adf.test(sales_ts)
```

## Output:

```
> #adf  
> adf.test(sales_ts)  
  
Augmented Dickey-Fuller Test  
  
data: sales_ts  
Dickey-Fuller = -0.72851, Lag order = 2, p-value = 0.956  
alternative hypothesis: stationary  
.
```

## Code:

```
best_model<-auto.arima(sales_ts)  
print(best_model)
```

## Output:

```
> best_model<-auto.arima(sales_ts)  
> print(best_model)  
Series: sales_ts  
ARIMA(0,1,0) with drift  
  
Coefficients:  
      drift  
      19.5652  
s.e.   1.4392  
  
sigma^2 = 49.8: log likelihood = -77.07  
AIC=158.13  AICc=158.73  BIC=160.41
```

## Code:

```
#forecast for next 12 months  
  
forecast_sales<-forecast(best_model,h=12)  
print(forecast_sales)
```

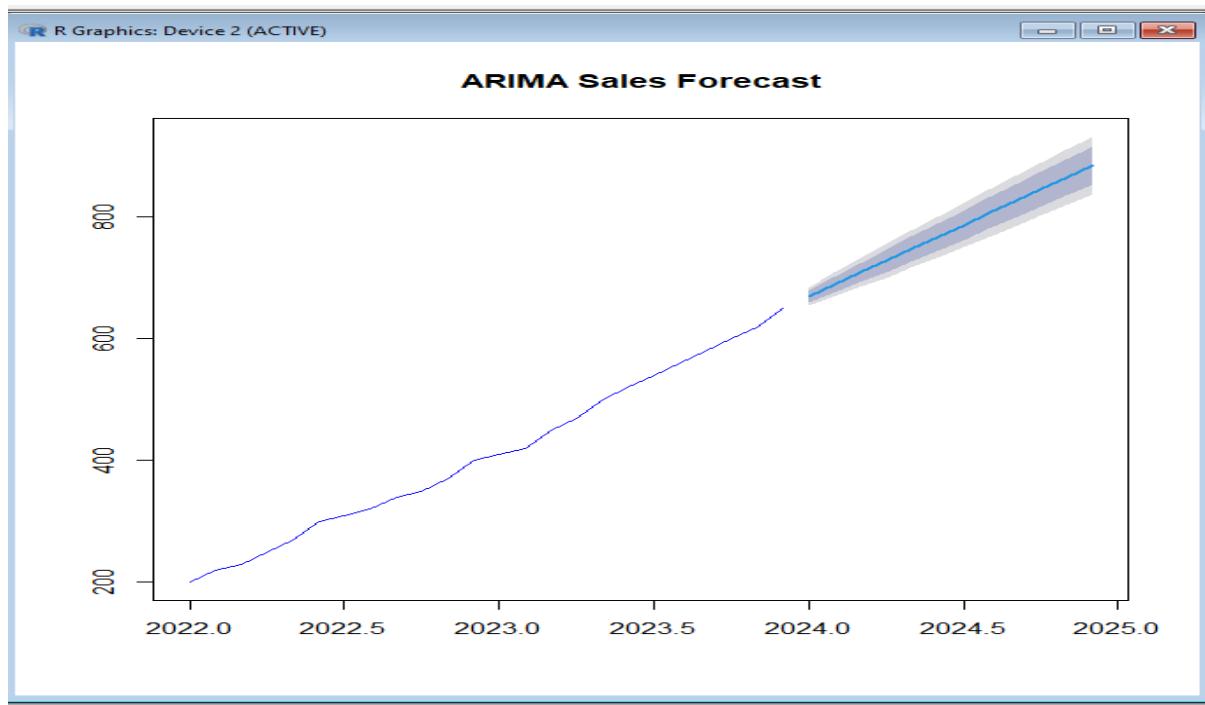
## Output:

```
-----  
> #forecast for next 12 months  
> forecast_sales<-forecast(best_model,h=12)  
> print(forecast_sales)  
    Point Forecast     Lo 80     Hi 80     Lo 95     Hi 95  
Jan 2024      669.5652 660.5211 678.6094 655.7334 683.3970  
Feb 2024      689.1304 676.3401 701.9208 669.5693 708.6916  
Mar 2024      708.6957 693.0307 724.3606 684.7382 732.6531  
Apr 2024      728.2609 710.1726 746.3492 700.5972 755.9245  
May 2024      747.8261 727.6028 768.0494 716.8972 778.7550  
Jun 2024      767.3913 745.2378 789.5448 733.5104 801.2722  
Jul 2024      786.9565 763.0280 810.8851 750.3609 823.5521  
Aug 2024      806.5217 780.9410 832.1024 767.3994 845.6441  
Sep 2024      826.0870 798.9545 853.2194 784.5915 867.5824  
Oct 2024      845.6522 817.0521 874.2523 801.9121 889.3923  
Nov 2024      865.2174 835.2214 895.2134 819.3424 911.0924  
Dec 2024      884.7826 853.4528 916.1124 836.8678 932.6975
```

## Code:

```
#plot forecast  
  
plot(forecast_sales,main="ARIMA Sales Forecast",col="blue")
```

## Output:



## Code:

```
#Arima on real dataset(Airline Passenger)
library(forecast)
library(tseries)
#load dataset
url<-
https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv
df<-read.csv(url)
head(df)
```

## Output:

```
> df<-read.csv(url)
> head(df)
  Month Passengers
1 1949-01        112
2 1949-02        118
3 1949-03        132
4 1949-04        129
5 1949-05        121
6 1949-06        135
> |
```

## Code:

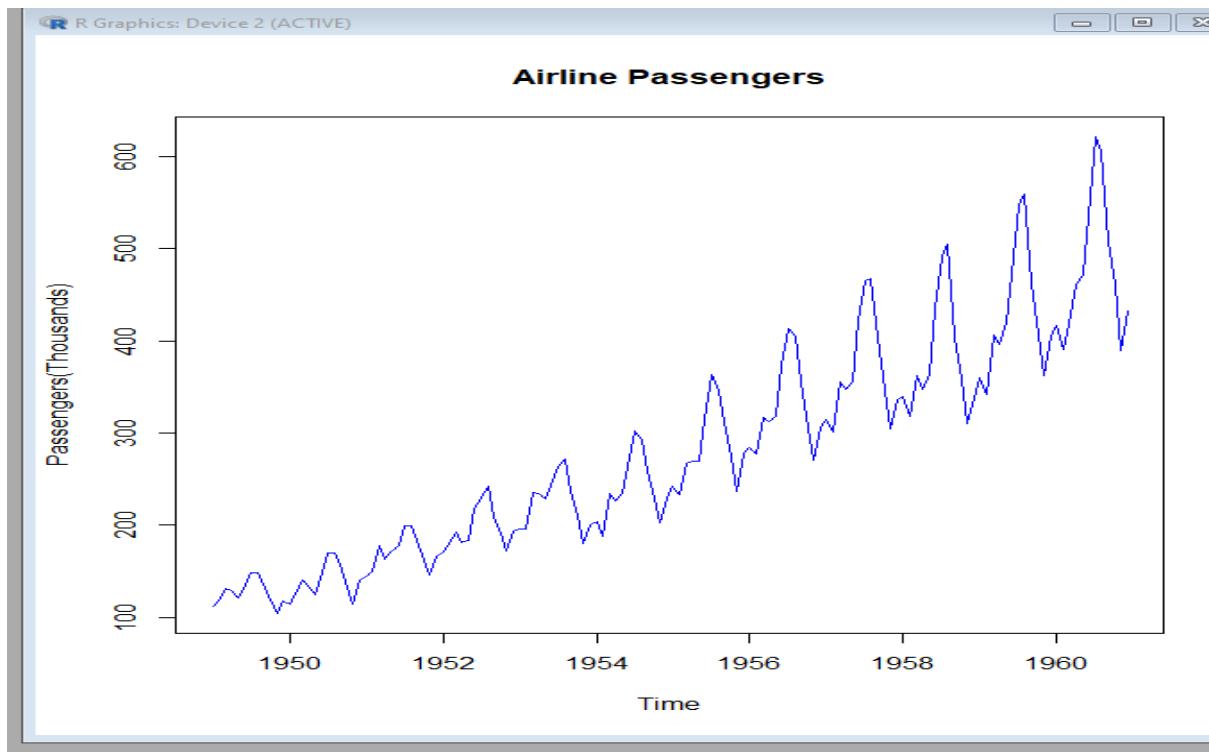
```
#step 2 convert to time series
```

```
air_ts<-ts(df$Passengers,start=c(1949,1),frequency=12)
```

```
#plot the data
```

```
plot(air_ts,main="Airline  
Passengers",ylab="Passengers(Thousands)",col="blue")
```

## Output:



## Code:

```
#check stationarity
```

```
adf.test(air_ts)
```

## Output:

```
> #check stationarity
> adf.test(air_ts)

Augmented Dickey-Fuller Test

data: air_ts
Dickey-Fuller = -7.3186, Lag order = 5, p-value = 0.01
alternative hypothesis: stationary

Warning message:
In adf.test(air_ts) : p-value smaller than printed p-value
~ |
```

## Code:

```
#fit arima model  
  
model_air<-auto.arima(air_ts)  
  
summary(model_air)
```

## Output:

```
> #fit arima model  
>  
> model_air<-auto.arima(air_ts)  
> summary(model_air)  
Series: air_ts  
ARIMA(2,1,1) (0,1,0) [12]  
  
Coefficients:  
      ar1      ar2      ma1  
    0.5960   0.2143  -0.9819  
  s.e.  0.0888   0.0880   0.0292  
  
sigma^2 = 132.3:  log likelihood = -504.92  
AIC=1017.85  AICc=1018.17  BIC=1029.35  
  
Training set error measures:  
      ME     RMSE     MAE     MPE     MAPE     MASE     ACFL  
Training set 1.3423 10.84619 7.86754 0.420698 2.800458 0.245628 -0.00124847  
> |
```

## Code:

```
#forecast next 24 months  
  
forecast_air<-forecast(model_air,h=24)  
  
print(forecast_air)
```

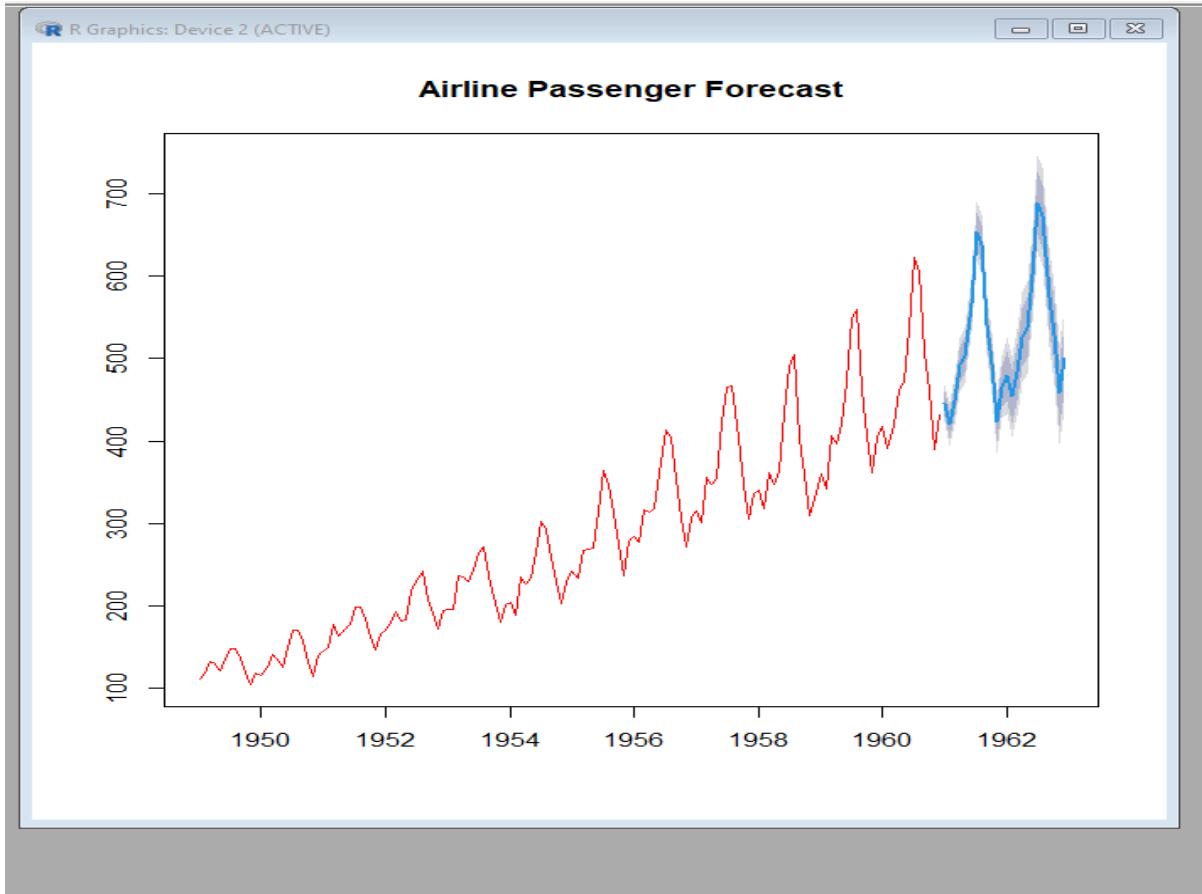
## Output:

```
> #forecast next 24 months  
> forecast_air<-forecast(model_air,h=24)  
> print(forecast_air)  
      Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95  
Jan 1961    445.6349 430.8903 460.3795 423.0851 468.1847  
Feb 1961    420.3950 403.0907 437.6993 393.9304 446.8596  
Mar 1961    449.1983 429.7726 468.6240 419.4892 478.9074  
Apr 1961    491.8399 471.0270 512.6529 460.0092 523.6707  
May 1961    503.3945 481.5559 525.2330 469.9953 536.7937  
Jun 1961    566.8624 544.2637 589.4612 532.3007 601.4242  
Jul 1961    654.2602 631.0820 677.4383 618.8122 689.7081  
Aug 1961    638.5975 614.9704 662.2246 602.4630 674.7320  
Sep 1961    540.8837 516.9028 564.8647 504.2081 577.5594  
Oct 1961    494.1266 469.8624 518.3909 457.0177 531.2356  
Nov 1961    423.3327 398.8381 447.8273 385.8715 460.7939  
Dec 1961    465.5076 440.8229 490.1923 427.7556 503.2596  
Jan 1962    479.2908 448.9986 509.5831 432.9629 525.6188  
Feb 1962    454.1768 421.7184 486.6353 404.5359 503.8178  
Mar 1962    483.0870 448.7343 517.4396 430.5491 535.6248  
Apr 1962    525.8193 490.1122 561.5263 471.2101 580.4284  
May 1962    537.4507 500.6863 574.2151 481.2244 593.6770  
Jun 1962    600.9839 563.3924 638.5754 543.4927 658.4752  
Jul 1962    688.4370 650.1834 726.6907 629.9331 746.9410  
Aug 1962    672.8213 634.0292 711.6134 613.4940 732.1487  
Sep 1962    575.1475 535.9102 614.3847 515.1393 635.1557  
Oct 1962    528.4242 488.8131 568.0352 467.8443 589.0040  
Nov 1962    457.6590 417.7293 497.5886 396.5918 518.7261  
Dec 1962    499.8582 459.6529 540.0634 438.3695 561.3468
```

**Code:**

```
plot(forecast_air,main="Airline Passenger Forecast",col="red")
```

**Output:**

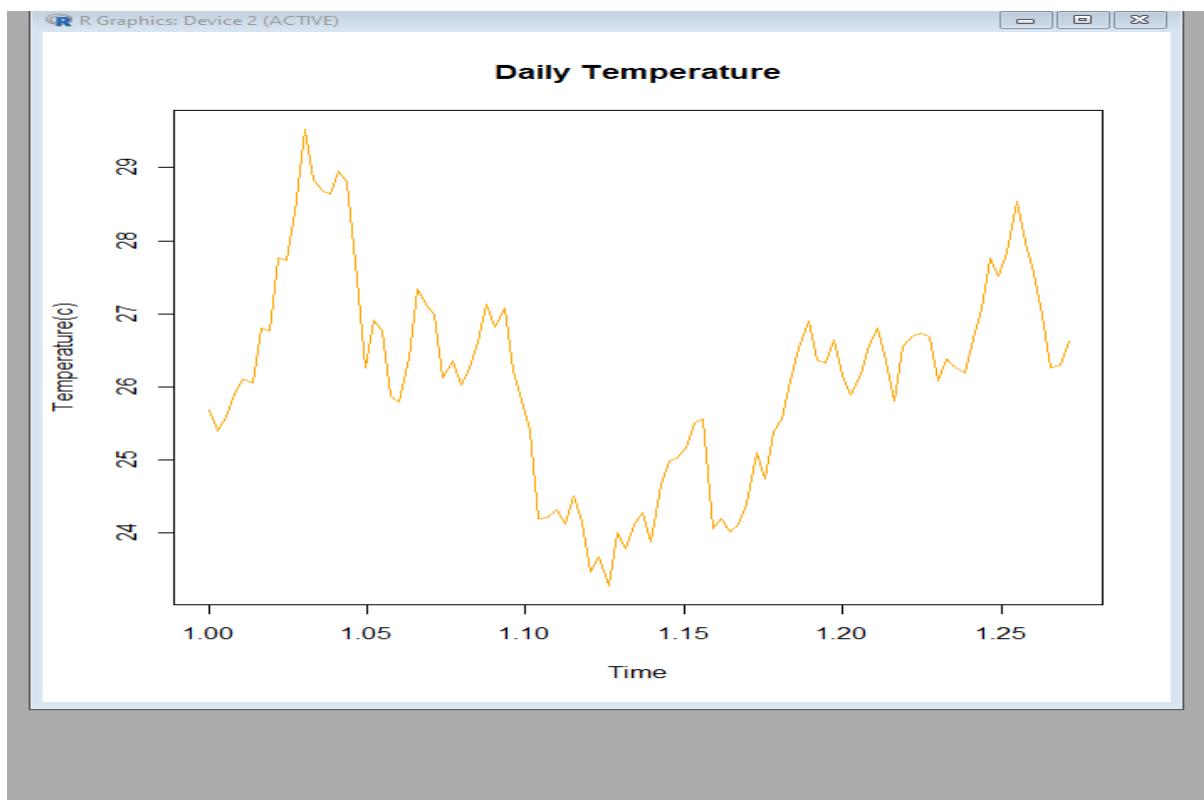


**Code:**

```
#example 3  
#Arima on daily temperature(synthetic data)  
#create dataset  
set.seed(42)  
days<-seq(as.Date("2022-01-01"),by="day",length.out=100)  
temperature<-25+cumsum(rnorm(100,0,0.5))#random noise  
#create time series  
temp_ts<-ts(temperature,frequency=365)
```

```
#plot  
plot(temp_ts,main="Daily  
Temperature",ylab="Temperature(c)",col="orange")
```

### Output:



### Code:

```
#fit arima model  
model_temp<-auto.arima(temp_ts)  
summary(model_temp)
```

### Output:

```
> model_temp<-auto.arima(temp_ts)  
> summary(model_temp)  
Series: temp_ts  
ARIMA(0,1,0)  
  
sigma^2 = 0.2666: log likelihood = -75.04  
AIC=152.08 AICc=152.12 BIC=154.68  
  
Training set error measures:  
ME RMSE MAE MPE MAPE MASE ACF1  
Training set 0.00965947 0.5137773 0.4064194 0.01774295 1.554414 NaN 0.0641011
```

## Code:

```
#forecast for next 30 days  
forecast_temp<-forecast(model_temp,h=30)  
print(forecast_temp)
```

## Output:

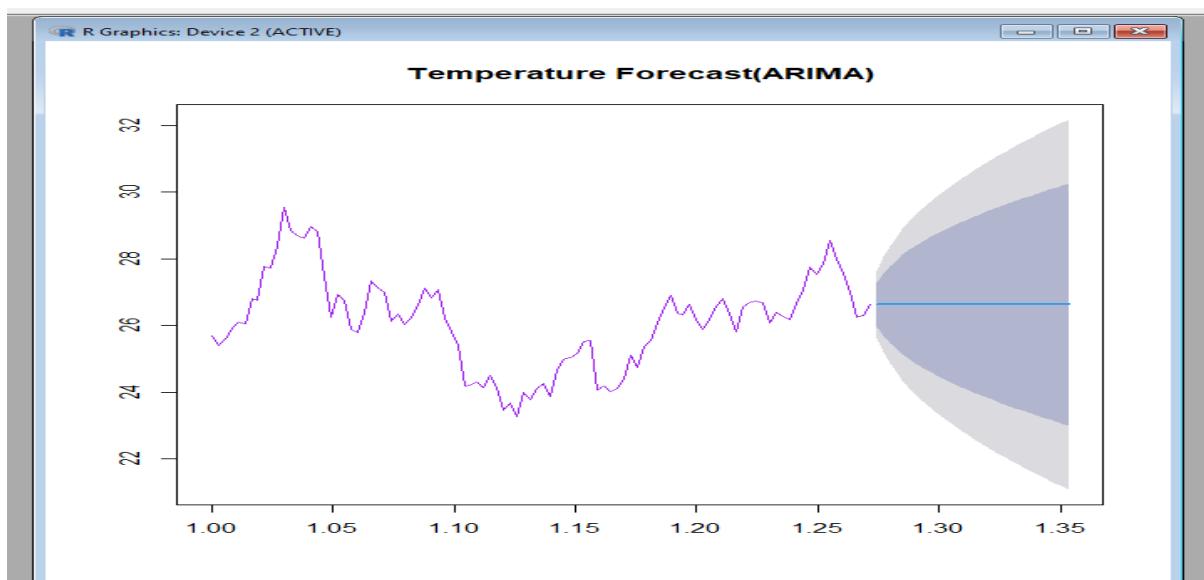
```
> #forecast for next 30 days  
> forecast_temp<-forecast(model_temp,h=30)  
> print(forecast_temp)  
   Point Forecast Lo 80 Hi 80 Lo 95 Hi 95  
1.273973 26.62574 25.96399 27.28749 25.61368 27.63780  
1.276712 26.62574 25.68989 27.56160 25.19447 28.05701  
1.279452 26.62574 25.47956 27.77192 24.87280 28.37868  
1.282192 26.62574 25.30224 27.94924 24.60162 28.64986  
1.284932 26.62574 25.14602 28.10546 24.36271 28.88877  
1.287671 26.62574 25.00479 28.24669 24.14672 29.10477  
1.290411 26.62574 24.87492 28.37656 23.94809 29.30339  
1.293151 26.62574 24.75403 28.49745 23.76321 29.48827  
1.295890 26.62574 24.64049 28.61099 23.55957 29.66191  
1.298630 26.62574 24.53311 28.71838 23.42533 29.82615  
1.301370 26.62574 24.43097 28.82051 23.26912 29.98236  
1.304110 26.62574 24.33337 28.91811 23.11987 30.13161  
1.306849 26.62574 24.23977 29.01171 22.97671 30.27477  
1.309589 26.62574 24.14970 29.10178 22.83897 30.41252  
1.312329 26.62574 24.06280 29.18868 22.70606 30.54542  
1.315068 26.62574 23.97874 29.27274 22.57751 30.67397  
1.317808 26.62574 23.89728 29.35420 22.45292 30.79856  
1.320548 26.62574 23.81818 29.43330 22.33194 30.91954  
1.323288 26.62574 23.74124 29.51024 22.21428 31.03720  
1.326027 26.62574 23.66631 29.58517 22.08868 31.15180  
1.328767 26.62574 23.59323 29.65826 21.98781 31.26357  
1.331507 26.62574 23.52186 29.72962 21.87877 31.37271  
1.334247 26.62574 23.45210 29.79938 21.77208 31.47940  
1.336986 26.62574 23.38385 29.86764 21.66769 31.58379  
1.339726 26.62574 23.31700 29.93449 21.56545 31.68603  
1.342466 26.62574 23.25147 30.00001 21.46524 31.78624  
1.345205 26.62574 23.18719 30.06429 21.36693 31.88455  
1.347945 26.62574 23.12409 30.12739 21.27043 31.98105  
1.350685 26.62574 23.06211 30.18937 21.17564 32.07584  
1.353425 26.62574 23.00119 30.25029 21.08247 32.16901
```

## Code:

```
#plot
```

```
plot(forecast_temp,main="Temperature  
Forecast(ARIMA)",col="purple")
```

## Output:



## PRACTICAL N0-8

**Aim:** Using Sarima for forecasting

**Code:**

```
#SARIMA with clear seasonal pattern and upward trenf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
#generate data
#Generate data
np.random.seed(0)
date_range=pd.date_range(start='2020-01-01',periods=48,freq='M')
trend=np.linspace(100,200,48)
seasonality=20*np.sin(2*np.pi*date_range.month/12)
noise=np.random.normal(0,5,48)
data=trend+seasonality+noise

df=pd.DataFrame({'Sales':data},index=date_range)
df
```

**Output:**

```
▶ seasonality=20*np.sin(2*np.pi*date_range.month/12)
  noise=np.random.normal(0,5,48)
  data=trend+seasonality+noise

  df=pd.DataFrame({'Sales':data},index=date_range)
  df

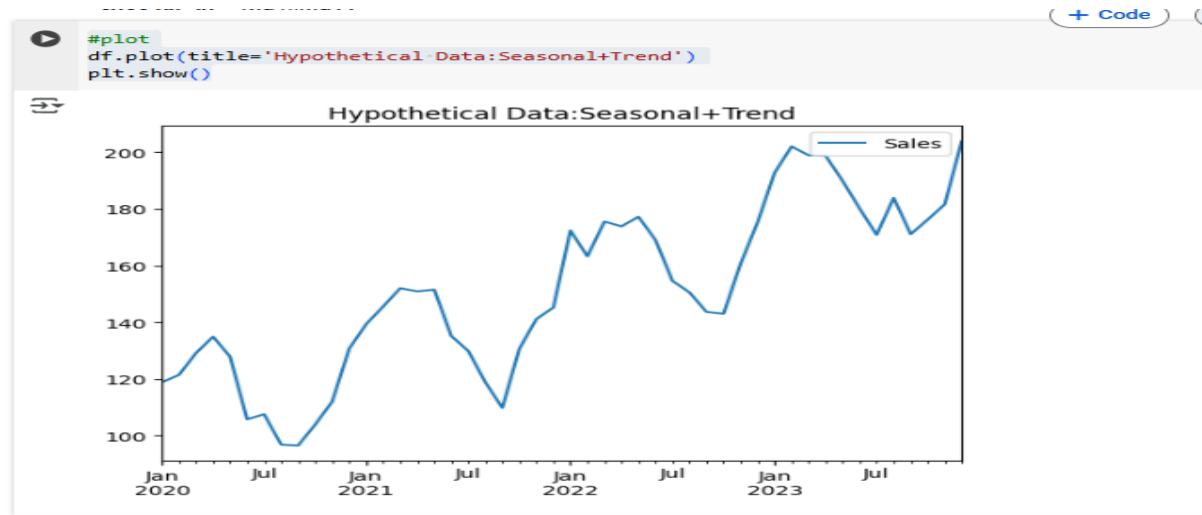
→ /tmp/ipython-input-955236531.py:4: FutureWarning: 'M' is deprecated and will be re
    date_range=pd.date_range(start='2020-01-01',periods=48,freq='M')

      Sales
      2020-01-31  118.820262
      2020-02-29  121.448954
      2020-03-31  129.149009
      2020-04-30  134.907953
      2020-05-31  127.848428
```

## Code:

```
#plot  
df.plot(title='Hypothetical Data:Seasonal+Trend')  
plt.show()
```

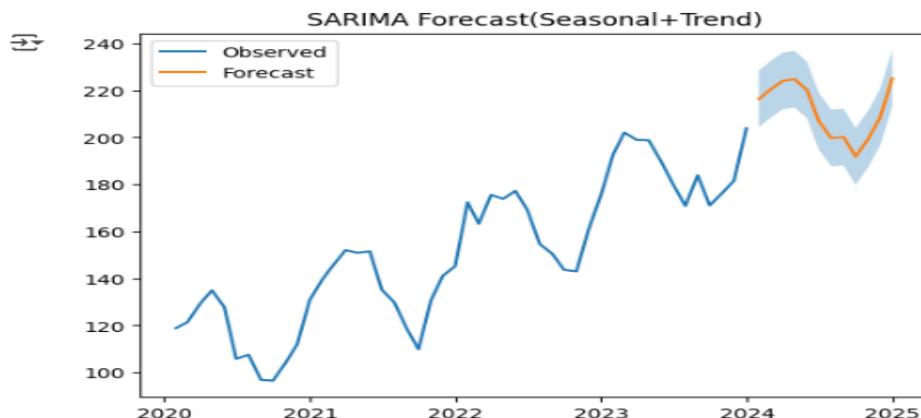
## Output:



## Code:

```
#fit sarima  
model=SARIMAX(df['Sales'],order=(1,1,1),seasonal_order=(1,1,1,12))  
result=model.fit()  
#forecast  
forecast=result.get_forecast(steps=12)  
forecast_df=forecast.summary_frame()  
plt.plot(df,label='Observed')  
plt.plot(forecast_df['mean'],label='Forecast')  
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],alpha=0.3)  
plt.title('SARIMA Forecast (Seasonal+Trend) ')  
plt.legend()  
plt.show()
```

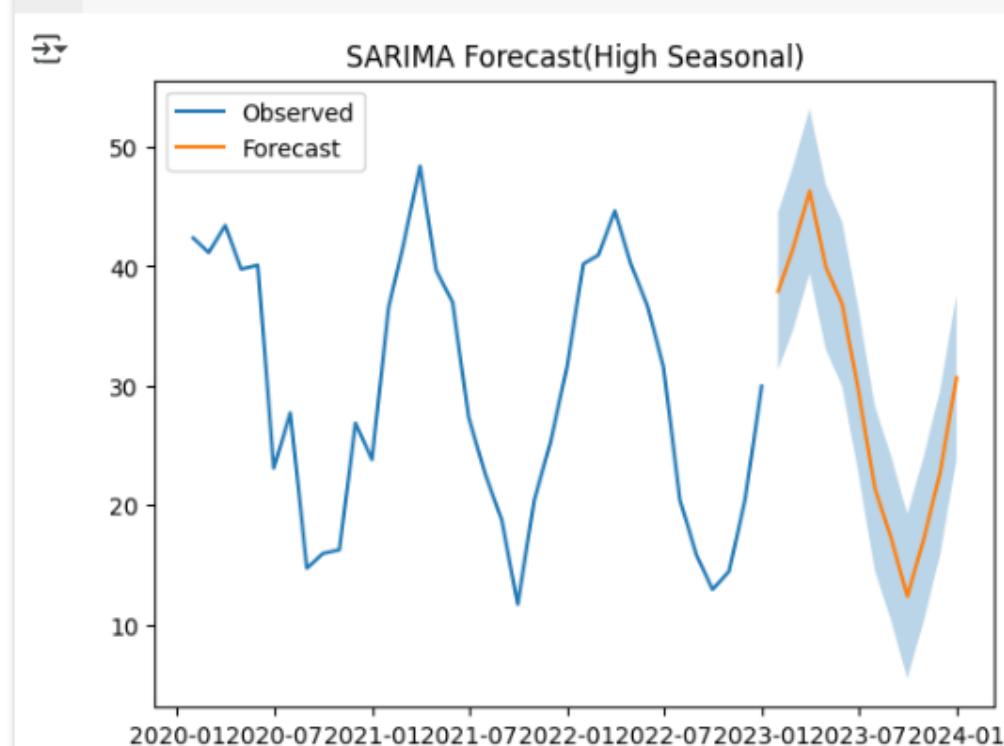
## Output:



## Code:

```
#example2 with high seasonality trend sysnthetic data no trend
np.random.seed(1)
dates=pd.date_range(start='2020-01-01',periods=36,freq='M')
seasonality=30+15*np.sin(2*np.pi*dates.month/12)
noise=np.random.normal(0,3,36)
values=seasonality+noise
df=pd.DataFrame({'Values':values},index=dates)
df
#fit Sarima
model=SARIMAX(df['Values'],order=(1,0,0),seasonal_order=(1,1,0,12))
res=model.fit()
#forecast
fcast=res.get_forecast(steps=12)
fcast_df=fcast.summary_frame()
plt.plot(df,label='Observed')
plt.plot(fcast_df['mean'],label='Forecast')
plt.fill_between(fcast_df.index,fcast_df['mean_ci_lower'],fcast_df['mean_ci_upper'],alpha=0.3)
plt.title('SARIMA Forecast(High Seasonal)')
plt.legend()
plt.show()
```

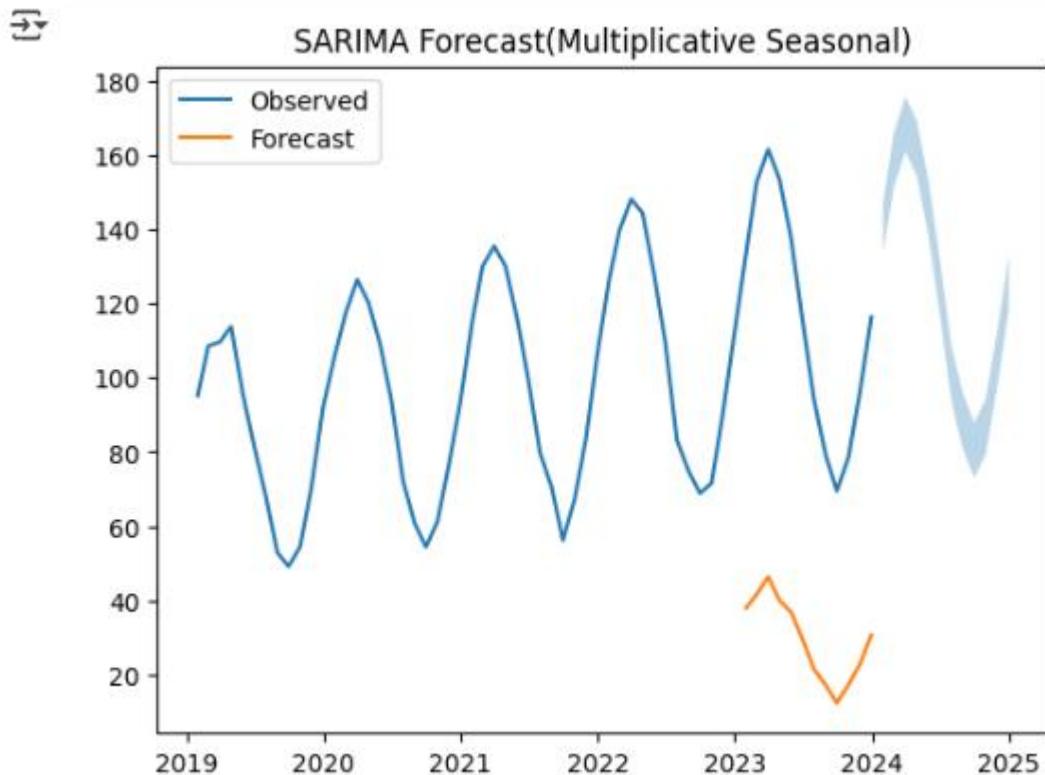
## Output:



## Code:

```
#example 3 Sarima with multiplicative seasonality and weak trend
np.random.seed(2)
months=pd.date_range(start='2019-01',periods=60,freq='M')
trend=np.linspace(80,120,60)
seasonal_factor=1+0.4*np.sin(2*np.pi*months.month/12)
noise=np.random.normal(0,2,60)
values=trend*seasonal_factor+noise
df=pd.DataFrame({'Value':values},index=months)
#fit sarima
model=SARIMAX(df['Value'],order=(1,1,1),seasonal_order=(1,1,1,12))
res=model.fit()
#forecast
pred=res.get_forecast(steps=12)
pred_df=pred.summary_frame()
#plot
plt.plot(df,label='Observed')
plt.plot(pred_df['mean'],label='Forecast')
plt.fill_between(pred_df.index,pred_df['mean_ci_lower'],pred_df['mean_ci_upper'],alpha=0.3)
plt.title('SARIMA Forecast (Multiplicative Seasonal)')
plt.legend()
plt.show()
```

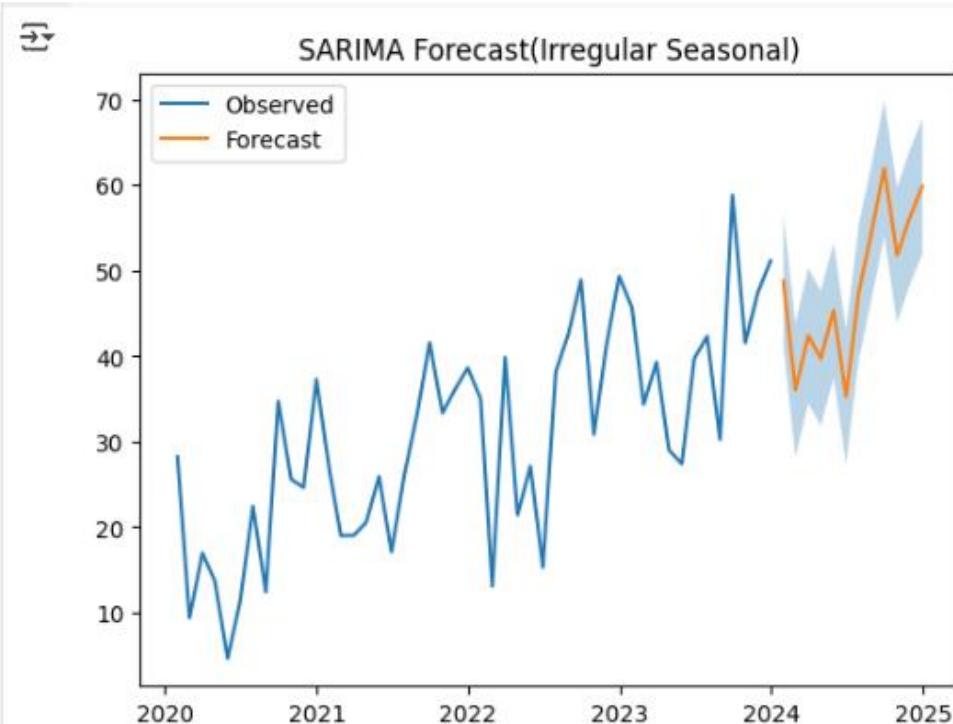
## Output:



## Code:

```
#example 4 SARIMA with irregular Seasonal Pattern
np.random.seed(3)
months=pd.date_range(start='2020-01',periods=48,freq='M')
pattern=np.random.choice([10,20,30],size=12,replace=True)
seasonality=np.tile(pattern,4)
trend=0.5*np.arange(48)
noise=np.random.normal(0,5,48)
values=trend+seasonality+noise
df=pd.DataFrame({'Value':values},index=months)
#fit sarima
model=SARIMAX(df['Value'],order=(1,1,1),seasonal_order=(2,1,0,12))
res=model.fit()
#forecast
future=res.get_forecast(steps=12)
future_df=future.summary_frame()
plt.plot(df,label='Observed')
plt.plot(future_df['mean'],label='Forecast')
plt.fill_between(future_df.index,future_df['mean_ci_lower'],future_df['mean_ci_upper'],alpha=0.3)
plt.title('SARIMA Forecast(Irregular Seasonal)')
plt.legend()
plt.show()
```

## Output:

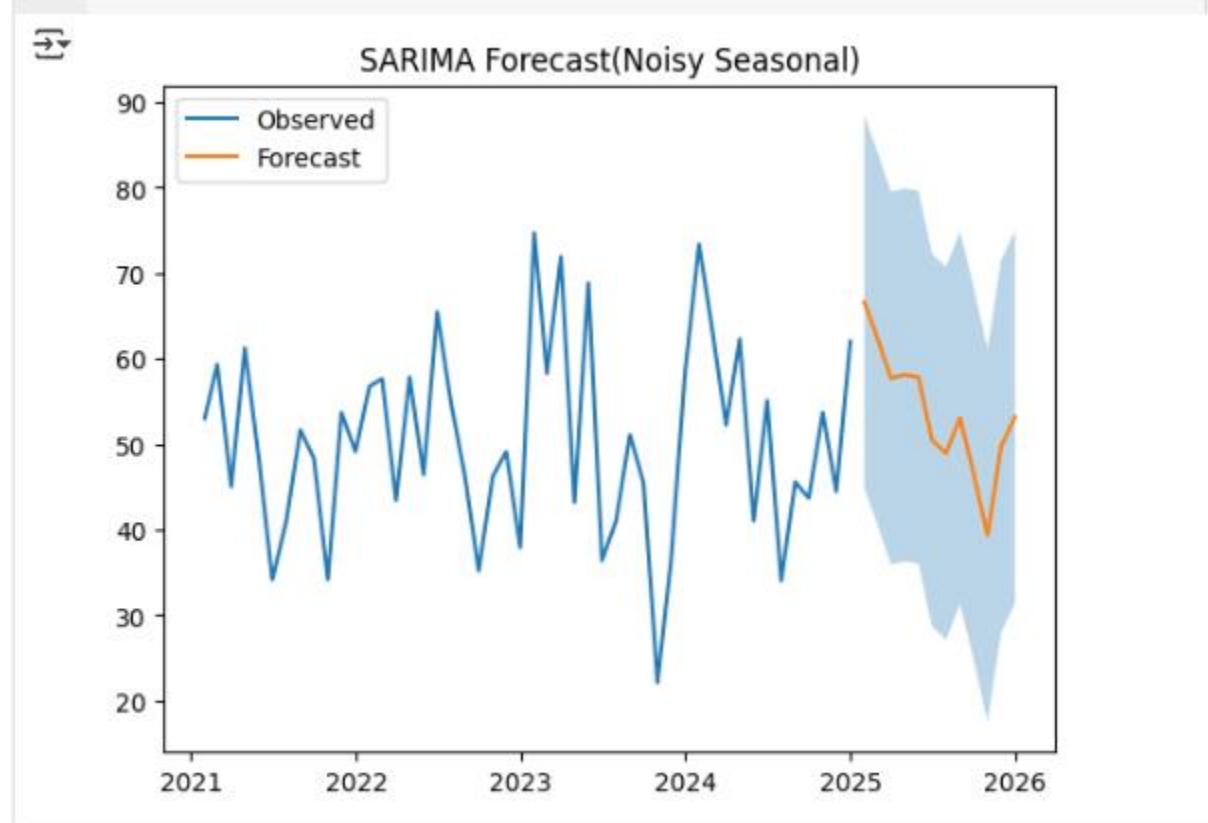


## Code:

```
#example 5: Sarima with weak seasonality and strong noise
np.random.seed(4)
months=pd.date_range(start='2021-01',periods=48,freq='M')
seasonal=5*np.sin(2*np.pi*months.month/12)
noise=np.random.normal(0,10,48)
values=50+seasonal+noise
df=pd.DataFrame({'Value':values},index=months)

#fit sarima
model=SARIMAX(df['Value'],order=(0,1,1),seasonal_order=(1,1,1,12))
res=model.fit()
#forecast
forecast=res.get_forecast(steps=12)
forecast_df=forecast.summary_frame()
plt.plot(df,label='Observed')
plt.plot(forecast_df['mean'],label='Forecast')
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],alpha=0.3)
plt.title('SARIMA Forecast(Noisy Seasonal)')
plt.legend()
plt.show()
```

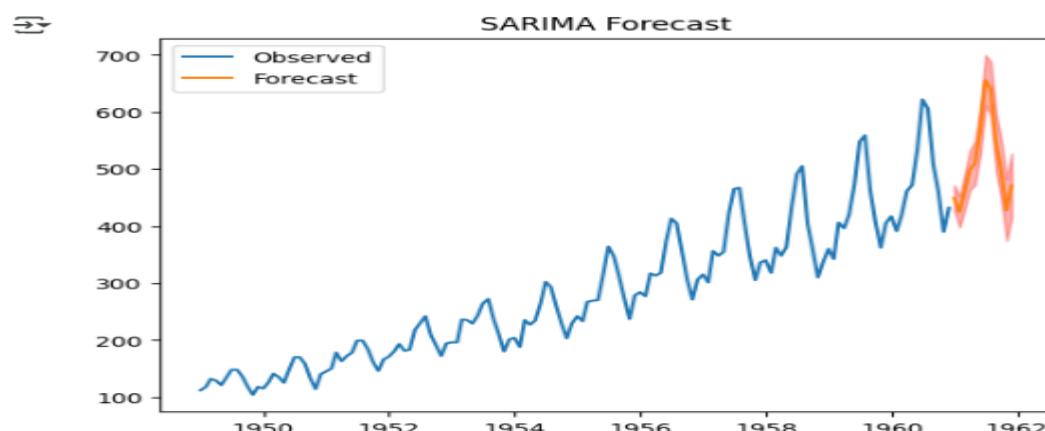
## Output:



## Code:

```
#example 6
#SARIMA on airline passengers dataset(seasonal monthly data)
import pandas as pd
from matplotlib import pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
#from pmdarima import auto_arima
from statsmodels.tsa.stattools import adfuller
#load dataset
url='https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv'
data=pd.read_csv(url,parse_dates=['Month'],index_col='Month')
data.columns=['Passengers']
#plot
data.plot(title='Airline Passengers')
plt.show()
#adf
result =adfuller(data['Passengers'])
print("ADF p-value:",result[1])
#fit Sarima
sarima=SARIMAX(data['Passengers'],order=(1,1,1),seasonal_order=(1,1,1,12))
sarima_fit=sarima.fit()
print(sarima_fit.summary())
#forecast
forecast=sarima_fit.get_forecast(steps=12)
forecast_df=forecast.summary_frame()
plt.plot(data,label='Observed')
plt.plot(forecast_df['mean'],label='Forecast')
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],color='r',alpha=0.3)
plt.title('SARIMA Forecast')
plt.legend()
plt.show()
```

## Output:

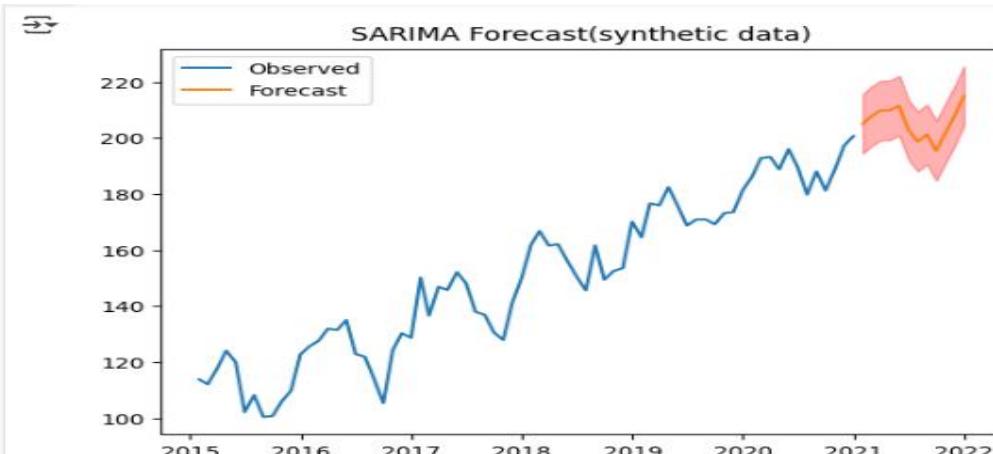


## Code:

```
#example 7 Sarima on synthetic Seasonal sales data
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
#create synthetic data
np.random.seed(0)
months=pd.date_range(start='2015-01',periods=72,freq='M')
seasonal=10*np.sin(2*np.pi*months.month/12)
trend=np.linspace(100,200,72)
noise=np.random.normal(0,5,72)
sales=trend+seasonal+noise
df=pd.DataFrame({'Sales':sales},index=months)
df
#plot
df.plot(title='Synthetic Seasonal Data')
plt.show()
#sarima model
model=SARIMAX(df['Sales'],order=(1,1,1),seasonal_order=(1,1,1,12))
result=model.fit()
#Forecast
#forecast
forecast=result.get_forecast(steps=12)
forecast_df=forecast.summary_frame()

#plot
plt.plot(df,label='Observed')
plt.plot(forecast_df['mean'],label='Forecast')
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],color='r',alpha=0.3)
plt.title('SARIMA Forecast(synthetic data)')
plt.legend()
plt.show()
```

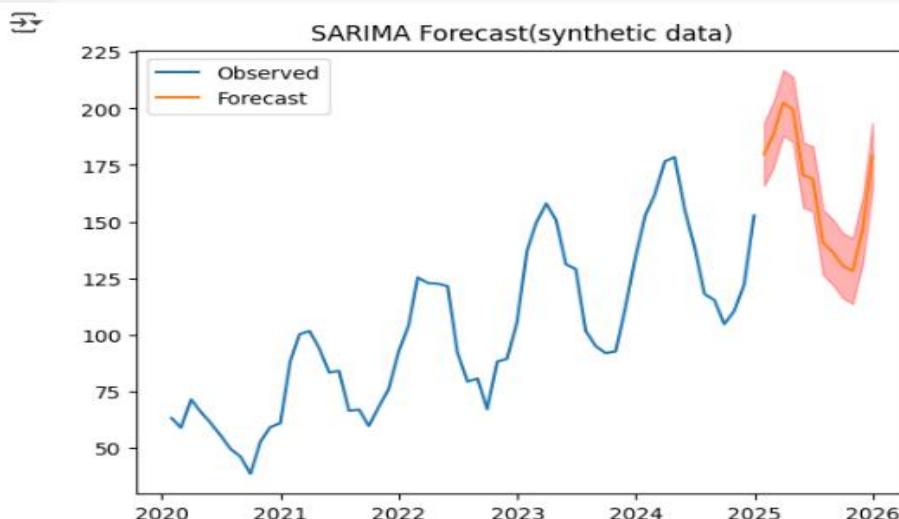
## Output:



## Code:

```
#sarima with multiplicative Seasonal component synthetic  
#multiplicative seasonality  
months=pd.date_range(start='2020-01',periods=60,freq='M')  
seasonal=1+0.3*np.sin(2*np.pi*months.month/12)  
trend=np.linspace(50,150,60)  
noise=np.random.normal(0,5,60)  
sales=trend*seasonal+noise  
df=pd.DataFrame({'Sales':sales},index=months)  
df  
  
#plot  
df.plot(title='Synthetic Multiplicative Seasonal Data')  
plt.show()  
#sarima model  
model=SARIMAX(df['Sales'],order=(1,1,1),seasonal_order=(1,1,1,12))  
result=model.fit()  
#forecast  
forecast=result.get_forecast(steps=12)  
forecast_df=forecast.summary_frame()  
#plot  
plt.plot(df,label='Observed')  
  
plt.plot(forecast_df['mean'],label='Forecast')  
plt.fill_between(forecast_df.index,forecast_df['mean_ci_lower'],forecast_df['mean_ci_upper'],color='r',alpha=0.3)  
plt.title('SARIMA Forecast(synthetic data)')  
plt.legend()  
plt.show()
```

## Output:



## PRACTICAL N0 -9

**Aim:** Aic, Bic & Fpe

**Code:**

```
from inspect import modulesbyfile
#example diwali sales aic,bic,fpe
#step1:generate dataset6
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
np.random.seed(42)
months=pd.date_range(start='2018-01-01',periods=60,freq='M')
#base trend=noise
sales=200+np.arange(60)*3+np.random.normal(0,15,60)
#add diwali effect:higher sales in oct-nov
for i,m in enumerate(months):
    if m.month in [10,11]:
        sales[i]+=100+np.random.randint(50,100)
data=pd.DataFrame({'Month':months,'Sales':sales})
data.set_index('Month',inplace=True)

#plot dataset
plt.figure(figsize=(12,4))
plt.plot(data['Sales'],marker='o',label="Monthly Sales")
plt.title("Synthetic Diwali Sales Data(2018-20522 )")
plt.xlabel("Month")
plt.ylabel("Sales")
plt.legend()
plt.show()

#unity function for aic, bic,fpe
from statsmodels.tsa.arima.model import ARIMA
def evaluate_arima(series,order):
    model=ARIMA(series,order=order)
    fit=model.fit()
    n_obs=fit.nobs
    sigma2=fit.mse
    k=fit.df_model+1
    fpe=sigma2*(n_obs+k)/(n_obs-k)
    return{"Order":order,"AIC":fit.aic,"BIC":fit.bic,"FPE":fpe}

#case1: ARIMA(1,1,1)
result1=evaluate_arima(data['Sales'],(1,1,1))
print("case1:",result1)

#case2: ARIMA(2,1,2)
result2=evaluate_arima(data['Sales'],(2,1,2))
```

```

print("case2:", result2)

#case3: ARIMA(3,1,2)
result3=evaluate_arima(data['Sales'], (3,1,2))
print("case3:", result3)

#case4: ARIMA(2,0,2)
result4=evaluate_arima(data['Sales'], (2,0,2))
print("case4:", result4)

#case5: ARIMA(1,1,2)
result5=evaluate_arima(data['Sales'], (1,1,2))
print("case5:", result5)

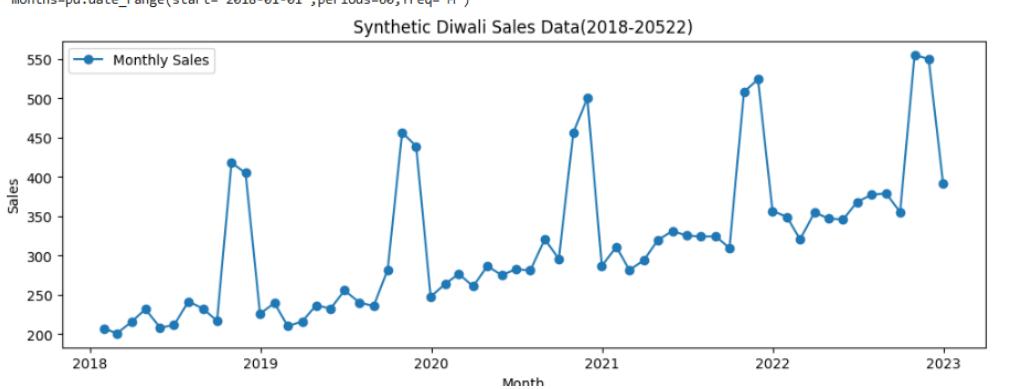
#compare all results
results=pd.DataFrame([result1,result2,result3,result4,result5])

print("\nComparsion of Model (AIC,BIC,FPE):")
print(results.sort_values(by="AIC"))

```

## Output:

 /tmp/ipython-input-3378351328.py:8: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.



 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 self.\_init\_dates(dates, freq)

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 self.\_init\_dates(dates, freq)

 /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

 self.\_init\_dates(dates, freq)

 case1: {'Order': (1, 1, 1), 'AIC': np.float64(675.7912864544597), 'BIC': np.float64(682.0238987861769), 'FPE': np.float64(6373.020392061551)}

 case2: {'Order': (2, 1, 2), 'AIC': np.float64(662.0576160724877), 'BIC': np.float64(672.4453032920163), 'FPE': np.float64(5194.986934089465)}

 case3: {'Order': (3, 1, 2), 'AIC': np.float64(664.0541013405347), 'BIC': np.float64(676.519326003969), 'FPE': np.float64(5373.086869940933)}

 case4: {'Order': (2, 0, 2), 'AIC': np.float64(675.1588481844228), 'BIC': np.float64(687.7249155577554), 'FPE': np.float64(4679.675364491733)}

 case5: {'Order': (1, 1, 2), 'AIC': np.float64(660.066182924016), 'BIC': np.float64(668.3763326996389), 'FPE': np.float64(5021.906342064552)}

Comparsion of Model (AIC,BIC,FPE):

Order	AIC	BIC	FPE
4 (1, 1, 2)	660.066183	668.376333	5021.906342
1 (2, 1, 2)	662.057616	672.445303	5194.986934
2 (3, 1, 2)	664.054101	676.519326	5373.086869
3 (2, 0, 2)	675.158848	687.724916	4679.675364
0 (1, 1, 1)	675.791286	682.023899	6373.020392

/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency ME will be used.

self.\_init\_dates(dates, freq)

## Code:

```
#example2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
np.random.seed(42)
#years from 2010 to 2022
years=pd.date_range(start='2010-12-01',end='2022-12-01',freq="12M")
#simulated automobile sales with upward trend+noise
sales=5000+(np.arange(len(years))*300)+np.random.normal(0,500,len(years))
data=pd.DataFrame ({ "Year":years, "Auto_Sales":sales})
data.set_index("Year",inplace=True)
#plot the dataset
plt.figure(figsize=(10,4))
plt.plot(data.index,data["Auto_Sales"],marker='o',label="Automobile
Sales(Dec)")
plt.title("Automobile Sales Data(2010-2022)")
plt.xlabel("Year")
plt.ylabel("Sales")
plt.legend()
plt.show()

#define function
from statsmodels.tsa.arima.model import ARIMA
def evaluate_arima(series,order):
    model=ARIMA(series,order=order)
    fit=model.fit()
    n_obs=fit.nobs
    sigma2=fit.mse
    k=fit.df_model+1
    fpe=sigma2*(n_obs+k)/(n_obs-k)
    return{"Order":order,"AIC":fit.aic,"BIC":fit.bic,"FPE":fpe}
#case1: ARIMA(1,1,1)
result1=evaluate_arima(data['Auto_Sales'],(1,1,1))
print("case1:",result1)

#case2: ARIMA(2,1,1)
result2=evaluate_arima(data['Auto_Sales'],(2,1,1))
print("case2:",result2)

#case3: ARIMA(2,1,2)
result3=evaluate_arima(data['Auto_Sales'],(2,1,2))
print("case3:",result3)

#case4: ARIMA(3,1,2)
result4=evaluate_arima(data['Auto_Sales'],(3,1,2))
print("case4:",result4)
```

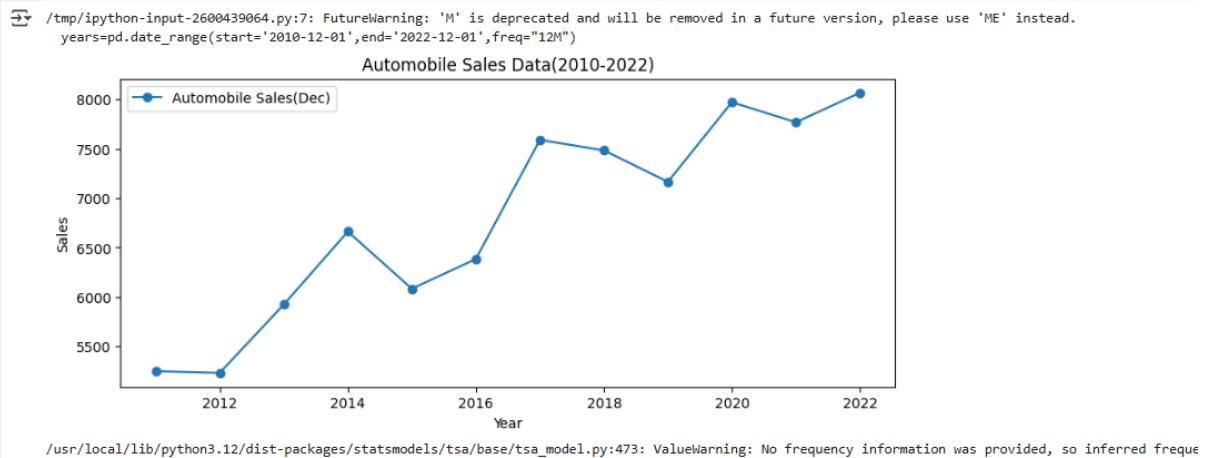
```

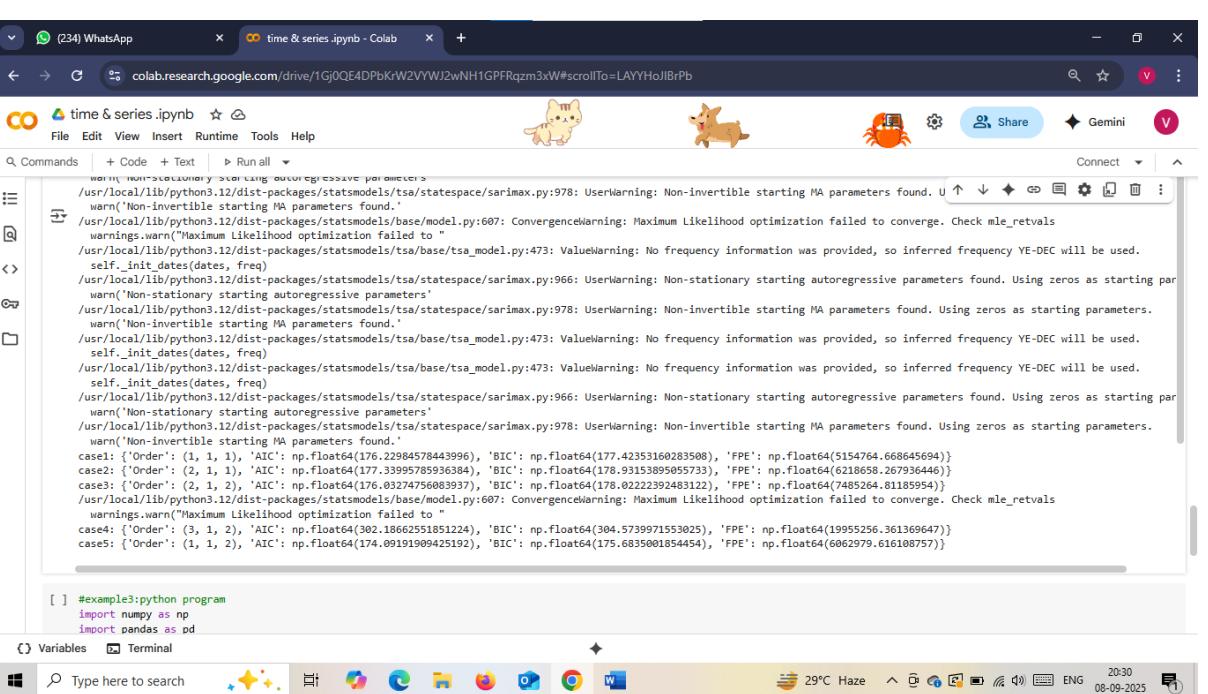
#case5: ARIMA(1,1,2)
result5=evaluate_arima(data['Auto_Sales'], (1,1,2))
print("case5:", result5)
#compare all results
results=pd.DataFrame([result1,result2,result3,result4,result5])

print("\nComparsion of Model (AIC,BIC,FPE):")
print(results.sort_values(by="AIC"))

```

## Output:





```

[234] WhatsApp x time & series.ipynb - Colab + 
← → ⌂ colab.research.google.com/drive/1Gj0QE4DPbKrW2vYVJ2wNH1GPFRqzm3xW#scrollTo=LAYHojlBrPb
File Edit View Insert Runtime Tools Help Gemini V
Q Commands + Code + Text ▶ Run all ▾
warning: non-stationary starting autoregressive parameters
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. U ↑ ↓ ◆ ↵ ☰ Share Connect ▾ ▾
warn('Non-invertible starting MA parameters found.')
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_revals
warnings.warn("Maximum Likelihood optimization failed to "
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YE-DEC will be used.
self._init_dates(freq)
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting par
warn('Non-stationary starting autoregressive parameters'
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
warn('Non-invertible starting MA parameters found.')
<> /usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YE-DEC will be used.
self._init_dates(freq)
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YE-DEC will be used.
self._init_dates(freq)
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting par
warn('Non-stationary starting autoregressive parameters'
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
warn('Non-invertible starting MA parameters found.')
case1: {'Order': (1, 1, 1), 'AIC': np.float64(176.22984578443996), 'BIC': np.float64(177.42353160283508), 'FPE': np.float64(5154764.668645694)}
case2: {'Order': (2, 1, 1), 'AIC': np.float64(177.33995785936384), 'BIC': np.float64(178.93153895055733), 'FPE': np.float64(6218658.267936446)}
case3: {'Order': (2, 1, 2), 'AIC': np.float64(176.03274756083937), 'BIC': np.float64(178.02222392483122), 'FPE': np.float64(7485264.81185954)}
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_revals
warnings.warn("Maximum Likelihood optimization failed to "
case4: {'Order': (3, 1, 2), 'AIC': np.float64(302.18662551851224), 'BIC': np.float64(304.5739971553025), 'FPE': np.float64(19955256.361369647)}
case5: {'Order': (1, 1, 2), 'AIC': np.float64(174.09191909425192), 'BIC': np.float64(175.6835001854454), 'FPE': np.float64(6062979.616108757)}

[ ] #example3:python program
import numpy as np
import pandas as pd

```

## Code:

```
#example3:python program
import numpy as np
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA
#synthetic data
np.random.seed(0)
n = 80
data=pd.Series(50+0.5*np.arange(n)+np.random.normal(0,2,n))
#fit arima
model=ARIMA(data,order=(2,1,2))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)
print("AIC",fit.aic)
print("BIC",fit.bic)
print("FPE",fpe)
```

## Output:

```
    print("BIC",fit.bic)
    print("FPE",fpe)

→ AIC 370.58181675910527
      BIC 382.4290560214404
      FPE 48.06658289120959
```

## Code:

```
#Synthetic time series
import numpy as np
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA
#generate synthetic trend data
np.random.seed(42)
n = 80
series = pd.Series(100+0.5*np.arange(n)+np.random.normal(0,2,n))
#fit arima(2,1,2)
model=ARIMA(series,order=(2,1,2))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)
```

```
print("Synthetic Example")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)
```

## Output:

```
↳ Synthetic Example
AIC: 368.4220658620451
BIC: 380.2693051243802
FPE: 154.42050581263337
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: U
  warn('Non-stationary starting autoregressive parameters'
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: U
  warn('Non-invertible starting MA parameters found.'
```

## Code:

```
import statsmodels.api as sm

data=pd.read_csv('AirPassengers.csv')
data['Month']=pd.to_datetime(data['Month'])
data.set_index('Month',inplace=True)

#fit arima
model=ARIMA(data['#Passengers'],order=(2,1,1))
fit=model.fit()

#compute fpe
n_obs=fit.nobs

print("\nAirline Passengers Example")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)
```

## Output:

```
↳ Airline Passengers Example
AIC: 1378.3383195987703
BIC: 1390.18969811981
FPE: 154.42050581263337
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py
  self._init_dates(dates, freq)
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py
  self._init_dates(dates, freq)
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/base/tsa_model.py
  self._init_dates(dates, freq)
```

## Code:

```
#Alternate method for mutliple orders
orders=[(1,1,1),(2,1,1),(3,1,2)]
results=[]

for order in orders:
    model=ARIMA(data['#Passengers'],order=order)
    fit=model.fit()
    n_obs=fit.nobs
    sigma2=fit.mse
    k=fit.df_model+1
    fpe=sigma2*(n_obs+k)/(n_obs-k)
    results.append((order,fit.aic,fit.bic,fpe))

import pandas as pd
df=pd.DataFrame(results,columns=["Order(p,q,d)", "AIC", "BIC", "FPE"])
print(df)
```

## Output:

```
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax
  warn('Non-stationary starting autoregressive parameters')
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax
  warn('Non-invertible starting MA parameters found.')
      Order(p,q,d)          AIC          BIC          FPE
0      (1, 1, 1)  1394.682505  1403.571039  1102.526488
1      (2, 1, 1)  1378.338320  1390.189698   992.972382
2      (3, 1, 2)  1377.086307  1394.863374   986.024100
```

## Code:

```
#example 4
#generate data
np.random.seed(0)
n=100
steps=np.random.normal(0,1,n)
series=pd.Series(np.cumsum(steps))
#fit arima
model=ARIMA(series,order=(1,1,1))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k)/(n_obs-k)
print("Synthetic Example")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)
```

## Output:

```
▶ Synthetic Example
AIC: 284.16336449892185
BIC: 291.9487240493256
FPE: 1.0770019210042068
```

## Code:

```
#Example 5

np.random.seed(10)
n=120
time=np.arange(n)
series=50+10*np.sin(2*np.pi*time/12)+np.random.normal(0,2,n)
series=pd.Series(series)

#fit arima
model=ARIMA(series,order=(1,1,1))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)

print("\n Synthetic stock prices Example 5")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)

model=ARIMA(series,order=(1,1,2))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)

print("\n Synthetic stock prices Example 5")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)

model=ARIMA(series,order=(2,1,2))
fit=model.fit()
#compute FPE
```

```

n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)

print("\n Synthetic stock prices Example 5")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)

model=ARIMA(series,order=(3,1,2))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)

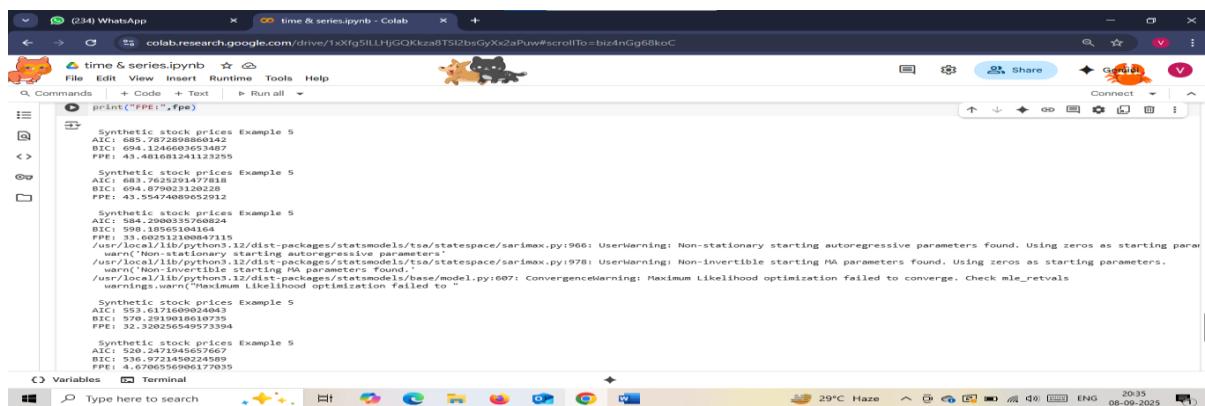
print("\n Synthetic stock prices Example 5")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)

model=ARIMA(series,order=(2,0,2))
fit=model.fit()
#compute FPE
n_obs=fit.nobs
sigma2=fit.mse
k=fit.df_model+1
fpe=sigma2*(n_obs+k) / (n_obs-k)

print("\n Synthetic stock prices Example 5")
print("AIC:",fit.aic)
print("BIC:",fit.bic)
print("FPE:",fpe)

```

## Output:



The screenshot shows a Google Colab notebook titled "time & series.ipynb". The code cell contains the same Python script as above. The output pane displays the results of the print statements and the ARIMA model fitting process. The terminal pane at the bottom shows standard output from the Python interpreter, including UserWarning messages about non-stationary and non-invertible parameters.

```

Synthetic stock prices Example 5
AIC: 685.7872898860142
BIC: 694.1246603653482
FPE: 43.5624889052912

Synthetic stock prices Example 5
AIC: 683.7625291477818
BIC: 692.1020205202912
FPE: 43.5624889052912

Synthetic stock prices Example 5
AIC: 584.2980335760822
BIC: 593.6360208051715
FPE: 33.602512108084715

/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
/usr/local/lib/python3.12/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
/usr/local/lib/python3.12/dist-packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retrvals
warnings.warn("Maximum Likelihood optimization failed to converge. Check mle_retrvals

Synthetic stock prices Example 5
AIC: 512.1216603653482
BIC: 520.291945657667
FPE: 33.602512108084715

Synthetic stock prices Example 5
AIC: 520.2472145657667
BIC: 536.9721456224589
FPE: 33.602512108084715

```

## PRACTICAL N0 –10

### Aim: ARCH & GARCH

#### Code:

```
pip install yfinance
pip install arch
#example1ARCH(1):Model on S & P 500 Returns
#step 1
import yfinance as yf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from arch import arch_model

#step2
#load S&P 500 daily data
data=yf.download("^GSPC",start="2020-01-01",end="2025-01-01")
returns=100*data['Close'].pct_change().dropna()
#step 3:Fit arch(1)model
model=arch_model(returns,vol='ARCH',p=1)
results=model.fit(disp='off')

print("Examole 1 - ARCH(1) on S&P 500 Returns")
print(results.summary())
```

#### Output:

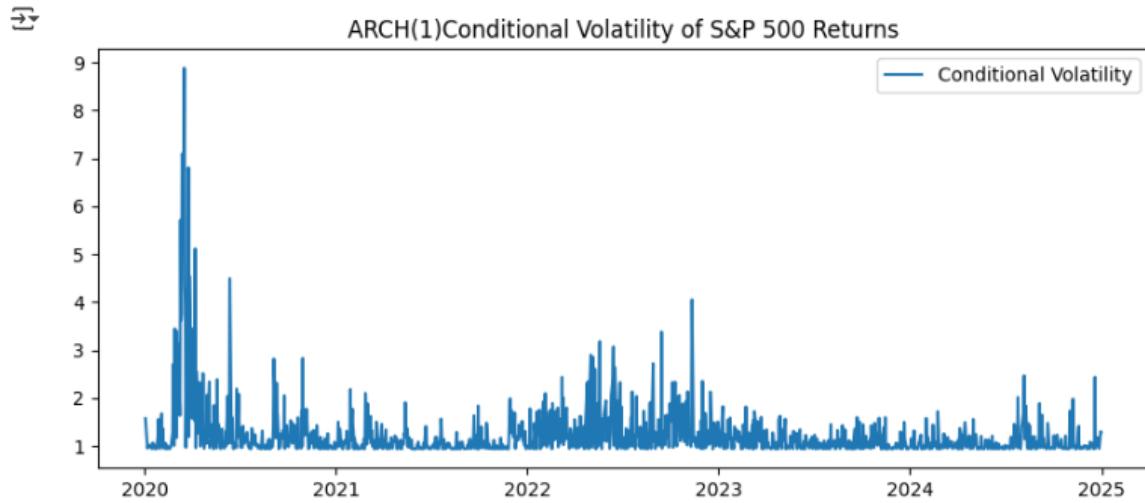
```
→ Examole 1 - ARCH(1) on S&P 500 Returns
      Constant Mean - ARCH Model Results
=====
Dep. Variable:          ^GSPC    R-squared:           0.000
Mean Model:            Constant Mean   Adj. R-squared:       0.000
Vol Model:              ARCH     Log-Likelihood:   -1974.89
Distribution:           Normal    AIC:                 3955.79
Method:                Maximum Likelihood   BIC:                 3971.19
                           No. Observations:      1257
Date:                  Sat, Sep 13 2025   Df Residuals:        1256
Time:                  10:21:59     Df Model:                   1
                           Mean Model
=====
            coef    std err        t     P>|t|  95.0% Conf. Int.
-----
mu        0.1347  3.299e-02     4.083  4.455e-05 [ 7.003e-02,  0.199]
Volatility Model
=====
            coef    std err        t     P>|t|  95.0% Conf. Int.
-----
omega     0.8862  8.363e-02     10.597 3.091e-26 [  0.722,  1.050]
alpha[1]   0.5308      0.123      4.319  1.567e-05 [  0.290,  0.772]
=====

Covariance estimator: robust
```

## Code:

```
#step 5: plot conditional volatility
plt.figure(figsize=(10,4))
plt.plot(results.conditional_volatility,label='Conditional Volatility')
plt.title('ARCH(1) Conditional Volatility of S&P 500 Returns')
plt.legend()
plt.show()
```

## Output:



## Code:

```
#example2:GARCH(1,1) model on S&P 500 returns
#step 1: fit garch(1,1) model
model=arch_model(returns,vol='GARCH',p=1,q=1)
results=model.fit(disp='off')

print("\nExample 2 - GARCH(1,1) on S&P 500 Returns")
print(results.summary())
#step 2: plot volatility
plt.figure(figsize=(10,4))
plt.plot(results.conditional_volatility,color='red')
plt.title('GARCH(1,1) Conditional Volatility of S&P 500 Returns')
plt.show()
```

## Output:



Example 2 - GARCH(1,1) on S&P 500 Returns

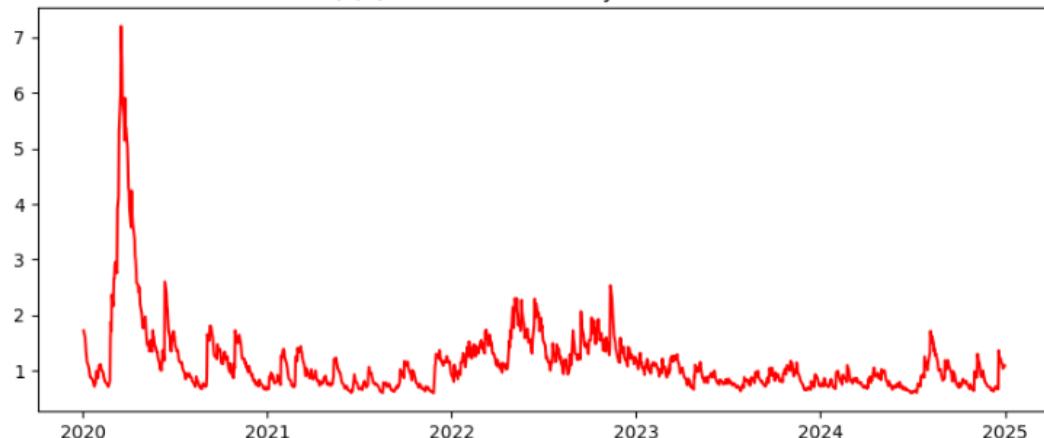
Constant Mean - GARCH Model Results

Dep. Variable:	^GSPC	R-squared:	0.000
Mean Model:	Constant Mean	Adj. R-squared:	0.000
Vol Model:	GARCH	Log-Likelihood:	-1832.95
Distribution:	Normal	AIC:	3673.91
Method:	Maximum Likelihood	BIC:	3694.45
Date:	Sat, Sep 13 2025	No. Observations:	1257
Time:	10:29:36	Df Residuals:	1256
		Df Model:	1
		Mean Model	
mu	0.0869	2.554e-02	3.404 6.651e-04 [3.687e-02, 0.137]
omega	0.0540	1.945e-02	2.775 5.515e-03 [1.586e-02, 9.210e-02]
alpha[1]	0.1560	3.458e-02	4.512 6.423e-06 [8.826e-02, 0.224]
beta[1]	0.8073	3.939e-02	20.494 2.431e-93 [ 0.730, 0.885]

Covariance estimator: robust



GARCH(1,1) Conditional Volatility of S&P 500 Returns



## Code:

```
#Example 3: GARCH(1,1) on gold prices
#Step 1: Download gold prices data(GLD ETF)
gold=yf.download('GLD',start='2020-01-01',end='2025-01-01')
gold_returns=100*gold['Close'].pct_change().dropna()

#Step 2: Fit GARCH(1,1) model
model=arch_model(gold_returns,vol='GARCH',p=1,q=1)
results=model.fit(disp='off')

print("\nExample 3 - GARCH(1,1) on Gold Returns")
```

```

print(results.summary())
#step 3
plt.figure(figsize=(10,4))
plt.plot(results.conditional_volatility,color='orange')
plt.title('GARCH(1,1) Conditional Volatility Gold')
plt.legend()
plt.show()

```

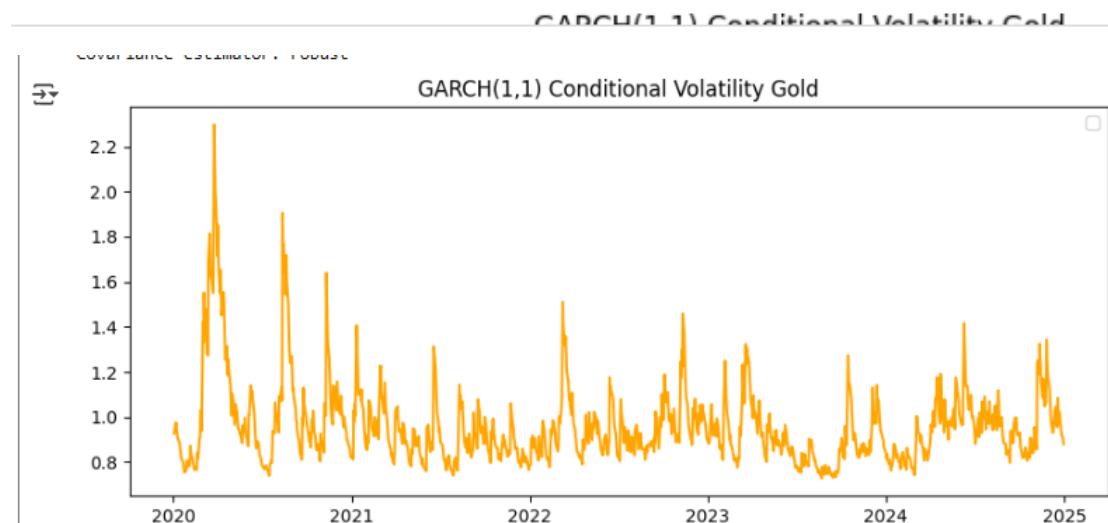
## Output:

```

Example 3 - GARCH(1,1) on Gold Returns
Constant Mean - GARCH Model Results
=====
Dep. Variable: GLD R-squared: 0.000
Mean Model: Constant Mean Adj. R-squared: 0.000
Vol Model: GARCH Log-Likelihood: -1714.30
Distribution: Normal AIC: 3436.61
Method: Maximum Likelihood BIC: 3457.15
Date: Sun, Sep 14 2025 No. Observations: 1257
Time: 12:52:46 Df Residuals: 1256
Mean Model Df Model: 1
=====
coef std err t P>|t| 95.0% Conf. Int.
-----
mu 0.0305 2.636e-02 1.158 0.247 [-2.113e-02, 8.222e-02]
Volatility Model
=====
coef std err t P>|t| 95.0% Conf. Int.
-----
omega 0.0817 2.423e-02 3.372 7.450e-04 [3.423e-02, 0.129]
alpha[1] 0.0890 2.563e-02 3.474 5.135e-04 [3.880e-02, 0.139]
beta[1] 0.8255 3.686e-02 22.399 4.048e-111 [ 0.753, 0.898]
=====

Covariance estimator: robust

```



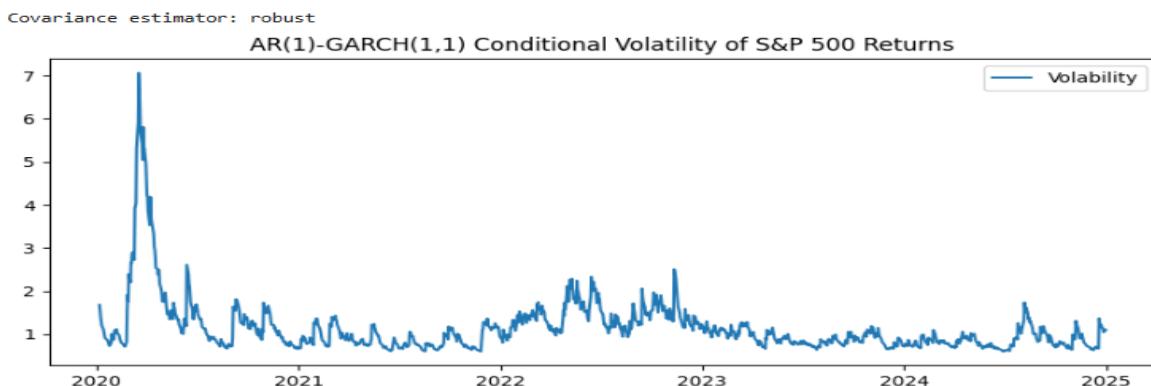
## Code:

```
#Example 4: GARCH with AR(1) Mean (ARIMA + GARCH)
#step1 :fit AR(1)-GARCH(1,1)
model=arch_model(returns,mean='AR',lags=1,vol='GARCH',p=1,q=1)
results=model.fit(disp='off')

print("\nExample 4 - AR(1)-GARCH(1,1) on S&P 500 Returns")
print(results.summary())
#step 2:
plt.figure(figsize=(10,4))
plt.plot(results.conditional_volatility,label='Volatility')
plt.title('AR(1)-GARCH(1,1) Conditional Volatility of S&P 500 Returns')
plt.legend()
plt.show()
```

## Output:

```
Example 4 - AR(1)-GARCH(1,1) on S&P 500 Returns
AR - GARCH Model Results
=====
Dep. Variable:          ^GSPC    R-squared:         0.011
Mean Model:             AR     Adj. R-squared:      0.010
Vol Model:              GARCH   Log-Likelihood:   -1831.10
Distribution:            Normal   AIC:            3672.20
Method:                 Maximum Likelihood   BIC:            3697.88
                           No. Observations:      1256
Date:       Sun, Sep 14 2025 Df Residuals:        1254
Time:           12:52:46   Df Model:                   2
               Mean Model
=====
                  coef    std err        t     P>|t|    95.0% Conf. Int.
-----
Const        0.0914  2.610e-02      3.502  4.611e-04  [4.025e-02,  0.143]
^GSPC[1]     -0.0344  3.160e-02     -1.090    0.276 [-9.637e-02,2.749e-02]
                           Volatility Model
=====
                  coef    std err        t     P>|t|    95.0% Conf. Int.
-----
omega        0.0540  1.935e-02      2.792  5.236e-03 [1.611e-02,9.198e-02]
alpha[1]      0.1560  3.508e-02      4.447  8.716e-06  [8.724e-02,  0.225]
beta[1]       0.8073  3.959e-02     20.391  1.998e-92   [ 0.730,  0.885]
=====
```



## Code:

```
#example 5: comparing GARCH(1,1) vs GARCH(2,1)
#step 1:
models=[(1,1), (2,1)]
results=[]
#step2:
for p,q in models:
    model=arch_model(returns, vol='GARCH', p=p, q=q)
    fit = model.fit(disp='off')
    results.append({ "Model":f"GARCH({p},{q})",
                      "AIC":fit.aic,
                      "BIC":fit.bic,
                      "LogLik":fit.loglikelihood})
#step 3:
df=pd.DataFrame(results)
print("\nExample 5 - Model Comparsion")
print(df)
```

## Output:



	Model	AIC	BIC	LogLik
0	GARCH(1,1)	3673.907096	3694.453029	-1832.953548
1	GARCH(2,1)	3675.500570	3701.182986	-1832.750285

## Code:

```
#program 6:Simulate & fit ARCH(1)Model
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from arch import arch_model

#step 1: simulate synthetic data(volatility Clustering)
np.random.seed(42)
n= 500
eps = np.random.normal(size=n)
vol=np.zeros(n)
ret=np.zeros(n)

alpha0,alpha1 = 0.1,0.8 #ARCH(1) parameters

for t in range(1,n):
    vol[t]=np.sqrt(alpha0+alpha1*ret[t-1]**2)
    ret[t]=vol[t]*eps[t]
series=pd.Series(ret)
```

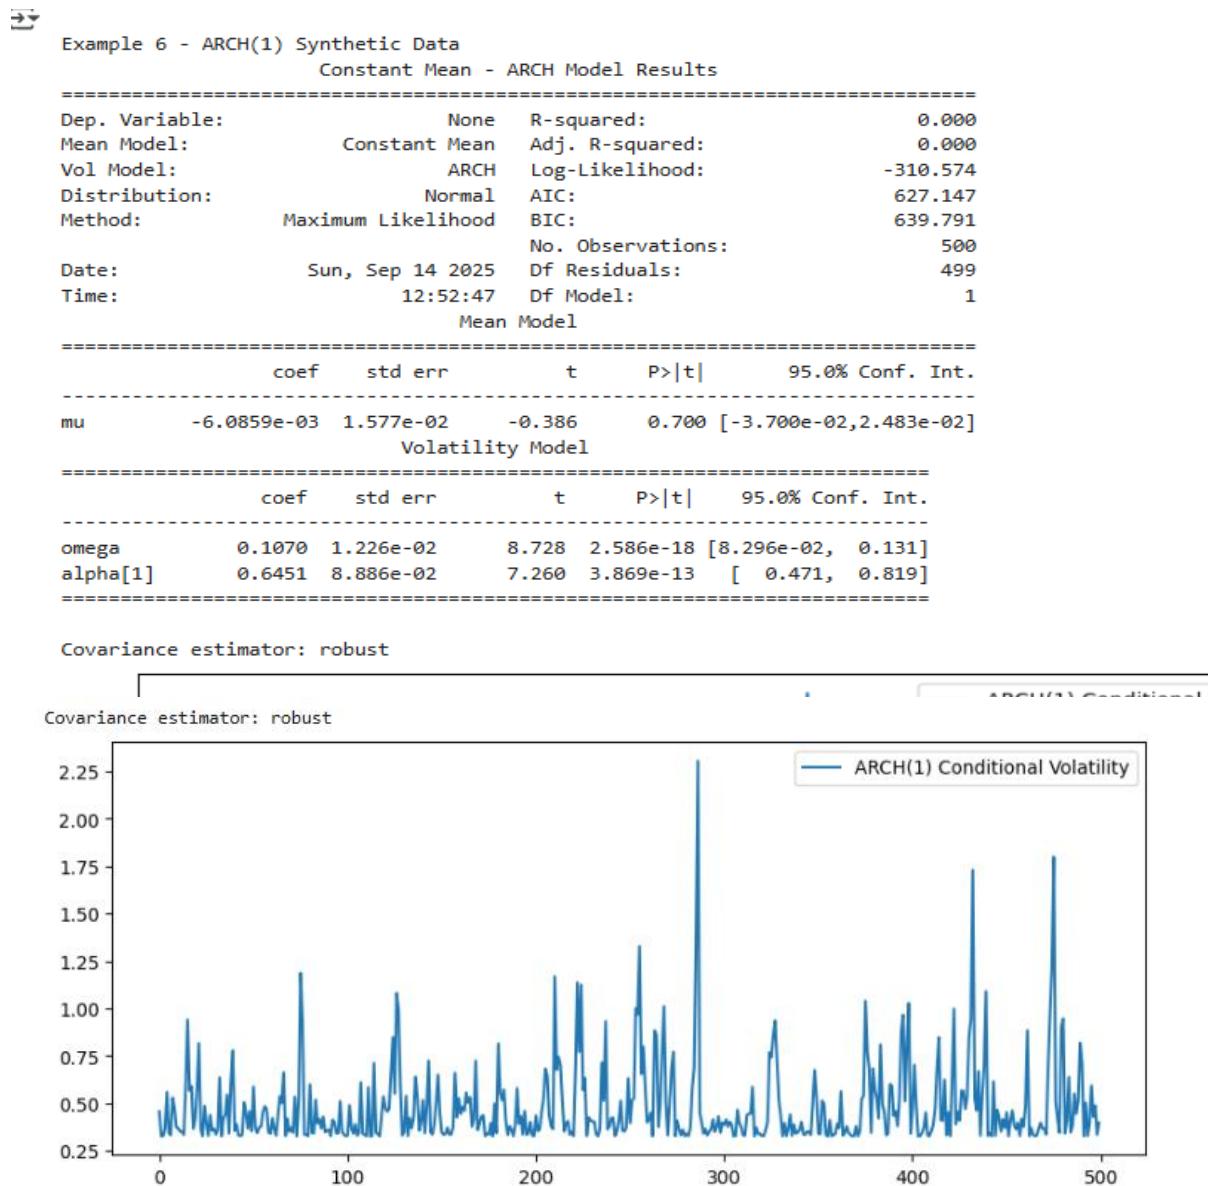
```

#step2:fit arch(1) model
arch_fit=arch_model(series,vol='ARCH',p=1).fit(disp='off')
print("\nExample 6 - ARCH(1) Synthetic Data")
print(arch_fit.summary())

#step 3: plot volatility
plt.figure(figsize=(10,4))
plt.plot(arch_fit.conditional_volatility,label='ARCH(1) Conditional Volatility')
plt.legend()
plt.show()

```

## Output:



## Code:

```
#program 7:Simulate & fit GARCH(1,1) model
#step 1:
np.random.seed(123)
n= 500
eps = np.random.normal(size=n)
ret=np.zeros(n)
vol=np.zeros(n)

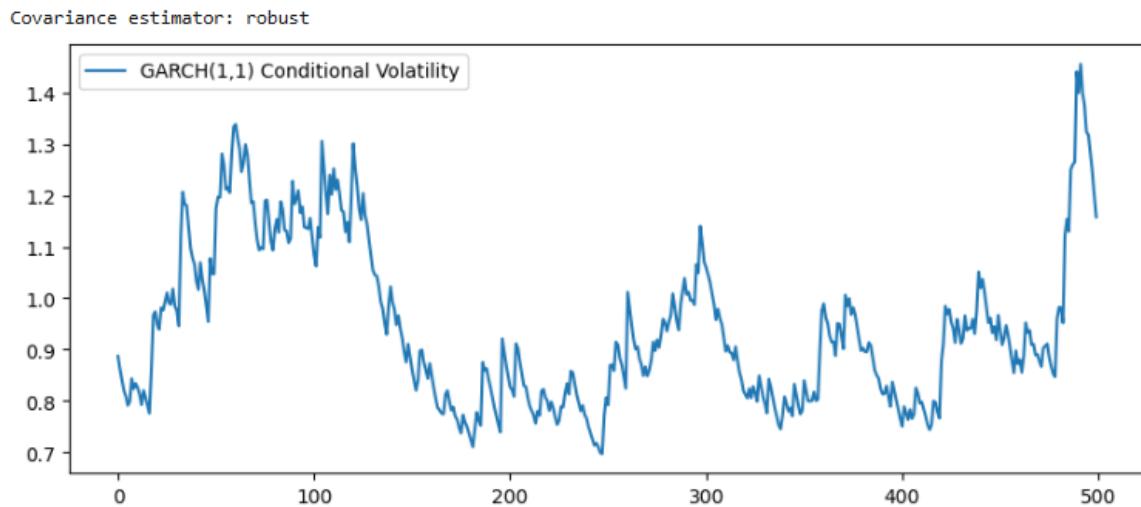
omega,alpha,beta=0.05,0.1,0.85 #GARCH(1,1)
for t in range(1,n):
    vol[t]=np.sqrt(omega+alpha*ret[t-1]**2+beta*vol[t-1]**2)
    ret[t]=vol[t]*eps[t]
series = pd.Series(ret)

#step2 fit garch(1,1)
garch_fit=arch_model(series,vol='GARCH',p=1,q=1).fit(disp='off')
print("\nExample 7 - GARCH(1,1) Synthetic Data")
print(garch_fit.summary())

plt.figure(figsize=(10,4))
plt.plot(garch_fit.conditional_volatility,label='GARCH(1,1) Conditional Volatility')
plt.legend()
plt.show()
```

## Output:

```
Example 7 - GARCH(1,1) Synthetic Data
Constant Mean - GARCH Model Results
=====
Dep. Variable: None R-squared: 0.000
Mean Model: Constant Mean Adj. R-squared: 0.000
Vol Model: GARCH Log-Likelihood: -672.793
Distribution: Normal AIC: 1353.59
Method: Maximum Likelihood BIC: 1370.44
No. Observations: 500
Date: Sun, Sep 14 2025 Df Residuals: 499
Time: 12:52:48 Df Model: 1
Mean Model
=====
      coef  std err      t  P>|t|  95.0% Conf. Int.
-----
mu     -0.0400  4.004e-02   -1.000    0.317 [-0.118, 3.845e-02]
Volatility Model
=====
      coef  std err      t  P>|t|  95.0% Conf. Int.
-----
omega  0.0295  1.539e-02    1.917  5.528e-02 [-6.656e-04, 5.965e-02]
alpha[1] 0.0620  1.687e-02    3.677  2.361e-04 [2.897e-02, 9.511e-02]
beta[1]  0.9068  2.421e-02   37.451  5.853e-307 [ 0.859,  0.954]
=====
```



## Code:

```
#example 8 Compare arch(1) vs garch(1,1)
#step1: fit arch(1)
arch_fit=arch_model(series,vol='ARCH',p=1).fit(disp='off')
#step2: compare aic/bic
print("\nProgram 3- Model Comparsion on Same Synthetic Data")
print("ARCH(1)-AIC:",arch_fit.aic,"BIC:",arch_fit.bic)
print("GARCH(1,1)-AIC:",garch_fit.aic,"BIC:",garch_fit.bic)
```

## Output:



Program 3- Model Comparsion on Same Synthetic Data  
 ARCH(1)-AIC: 1375.1455766213467 BIC: 1387.7894009166132  
 GARCH(1,1)-AIC: 1353.5850430802452 BIC: 1370.443475473934

## Code:

```
#program9:AR(1)-GARCH(1,1) (Mean+volatility)
np.random.seed(99)
n= 500
eps = np.random.normal(size=n)
ret=np.zeros(n)
vol=np.zeros(n)

phi=0.3 # AR(1) parameter
omega,alpha,beta=0.05,0.1,0.85 # GARCH(1,1) parameters: omega, alpha_1,
beta_1
for t in range(1,n):
    vol[t]=np.sqrt(omega+alpha*ret[t-1]**2+beta*vol[t-1]**2)
    ret[t]=phi*ret[t-1]+vol[t]*eps[t]

series=pd.Series(ret)
```

```
#step2: fit ar(1) - garch(1,1)
ar_garch_fit=arch_model(series,mean='AR',lags=1,vol='GARCH',p=1,q=1).fit(disp='off')
print("\nExample 8 - AR(1)-GARCH(1,1) on Synthetic Data")
print(ar_garch_fit.summary())
```

## Output:

```
Example 8 - AR(1)-GARCH(1,1) on Synthetic Data
AR - GARCH Model Results
Dep. Variable: None R-squared: 0.082
Mean Model: AR Adj. R-squared: 0.080
Vol Model: GARCH Log-Likelihood: -679.645
Distribution: Normal AIC: 1369.29
Method: Maximum Likelihood BIC: 1390.35
No. Observations: 499
Date: Sun, Sep 14 2025 Df Residuals: 497
Time: 12:52:48 Df Model: 2
Mean Model
=====
      coef    std err        t   P>|t|    95.0% Conf. Int.
-----
Const    0.0228  3.958e-02     0.575    0.565 [-5.480e-02,  0.100]
None[1]  0.2812  4.366e-02     6.440  1.192e-10    [ 0.196,  0.367]
          Volatility Model
=====
      coef    std err        t   P>|t|    95.0% Conf. Int.
-----
omega   0.0623  2.864e-02     2.173  2.976e-02 [6.111e-03,  0.118]
alpha[1] 0.1163  3.305e-02     3.519  4.327e-04 [5.153e-02,  0.181]
beta[1]  0.8226  5.141e-02    16.001 1.248e-57    [ 0.722,  0.923]
=====
```

## Code:

```
#program 1:Innovation from ar(1)model(Sunspots)
import matplotlib.pyplot as plt
from statsmodels.datasets import sunspots
from statsmodels.tsa.arima.model import ARIMA
#LOAD data
data = sunspots.load_pandas().data['SUNACTIVITY']
#fit ar(1)
model=ARIMA(data,order=(1,0,0)).fit()
innovations=model.resid

print("==Program 1 Output ==")
print(model.summary())
print("\nFirst 10 innovations:\n",innovations.head(10))

#plot
plt.figure(figsize=(10,4))
plt.plot(innovations)
plt.axhline(0,linestyle='--',color='black')
```

```
plt.title('Sunspots Innovation')
plt.show()
```

## Output:

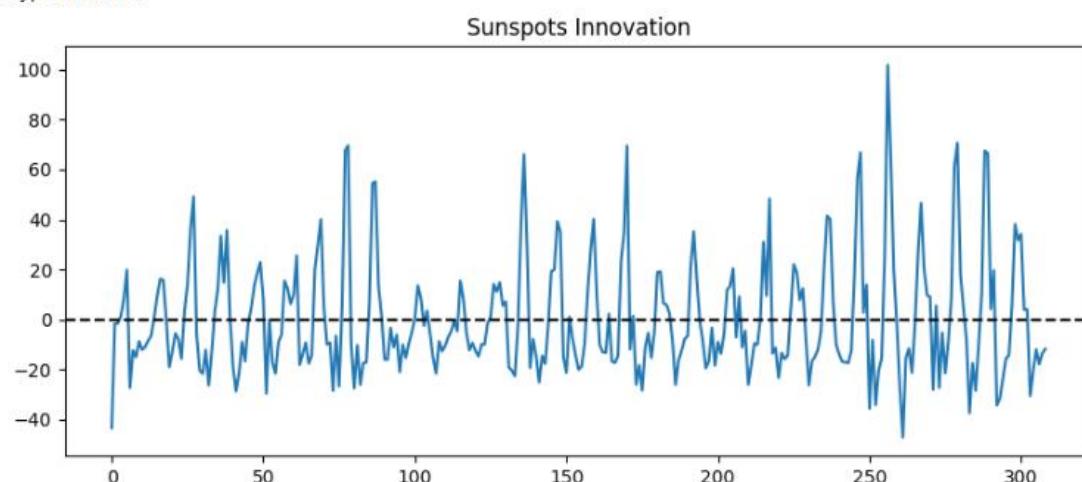
```
==Program 1 Output ==
SARIMAX Results
=====
Dep. Variable: SUNACTIVITY No. Observations: 309
Model: ARIMA(1, 0, 0) Log Likelihood: -1406.585
Date: Sun, 14 Sep 2025 AIC: 2819.169
Time: 12:52:50 BIC: 2830.369
Sample: 0 HQIC: 2823.647
                           - 309
Covariance Type: opg
=====
            coef    std err      z   P>|z|    [0.025    0.975]
-----
const    48.3963    9.949    4.865    0.000    28.897    67.895
ar.L1     0.8244    0.036   23.011    0.000     0.754    0.895
sigma2   524.5565   39.103   13.415    0.000   447.916   601.197
=====
Ljung-Box (L1) (Q): 100.01 Jarque-Bera (JB): 135.37
Prob(Q):          0.00 Prob(JB):          0.00
Heteroskedasticity (H): 1.81 Skew:           1.29
Prob(H) (two-sided): 0.00 Kurtosis:        4.96
=====
```

### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

### First 10 innovations:

```
1  -1.619376
2  -1.565917
3   1.311966
4   8.541001
5  19.823496
6 -27.313821
7 -12.405539
8 -14.985728
9 -8.741493
dtype: float64
```



## Code:

```
!pip install arch
import numpy as np
import pandas as pd
import yfinance as yf
from statsmodels.tsa.arima.model import ARIMA
from arch import arch_model
from typing_extensions import dataclass_transform
import matplotlib.pyplot as plt
#Step1: load data
ticker="AAPL"
data=yf.download(ticker,start="2015-01-01",end="2023-12-31",progress=False)
data=data['Close'].dropna()

#step2: compute return
returns=np.log(data).diff().dropna()

#step 3: fit arima model on return
arima_model_obj = ARIMA(returns, order=(1, 0, 0))
arima_res = arima_model_obj.fit()

#step 4: fit arch model on arima residuals
arch_model_obj=arch_model(arima_res.resid,vol="ARCH",p=1,o=0,q=0,dist='normal')
arch_res=arch_model_obj.fit(update_freq=5)
print(arch_res.summary())

#step 5:forecast
horizon = 5
arima_fore = arima_res.get_forecast(steps=horizon)
vol_fore = arch_res.forecast(horizon=horizon)
print("Mean forecast:\n",arima_fore.predicted_mean.tail(horizon))
print("Vol forecast:\n",vol_fore.variance.tail(horizon))

#Step6: plot
fig,ax=plt.subplots(2,1,figsize=(10,6),sharex=True)
ax[0].plot(arima_res.resid,label='ARIMA residuals')
ax[0].legend()
ax[1].plot(np.sqrt(arch_res.conditional_volatility),label='Conditional volatility (ARCH)')
ax[1].legend()
plt.show()
```

## Output:

```

Function evaluations: 45
Gradient evaluations: 6
Constant Mean - ARCH Model Results
=====  

Dep. Variable: None R-squared: 0.000
Mean Model: Constant Mean Adj. R-squared: 0.000
Vol Model: ARCH Log-Likelihood: 5926.25
Distribution: Normal AIC: -11846.5
Method: Maximum Likelihood BIC: -11829.3
Date: Mon, Sep 15 2025 No. Observations: 2263
Time: 11:26:19 Df Residuals: 2262
Df Model: 1  

Mean Model
=====  

      coef  std err      t    P>|t|  95.0% Conf. Int.  

-----  

mu      5.3491e-04  3.955e-04   1.352    0.176 [-2.403e-04, 1.310e-03]  

Volatility Model
=====  

      coef  std err      t    P>|t|  95.0% Conf. Int.  

-----  

omega   2.5262e-04  1.759e-05   14.364  8.735e-47 [2.181e-04, 2.871e-04]  

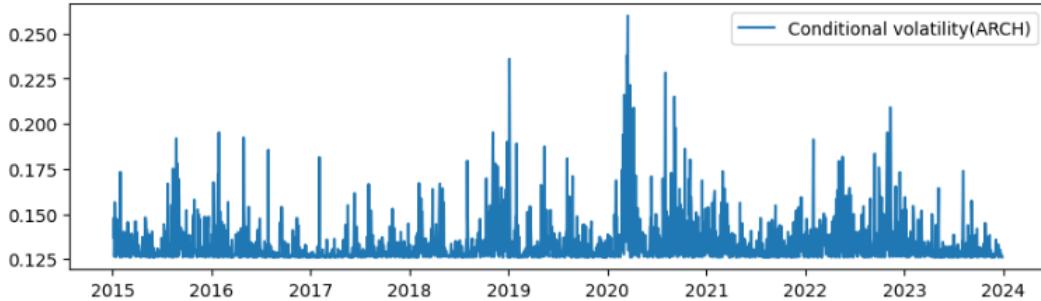
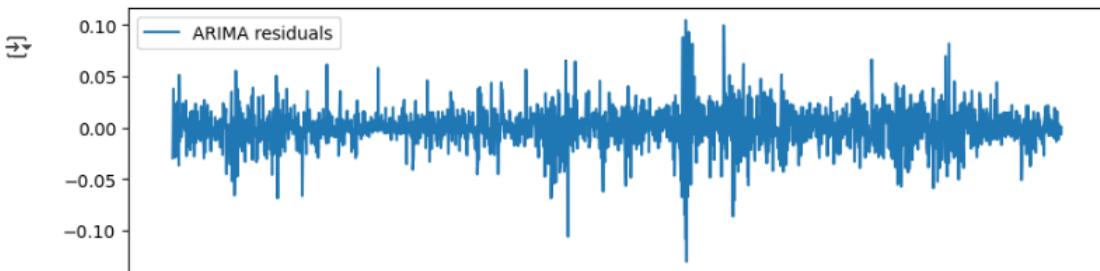
alpha[1]  0.2512   5.939e-02    4.229  2.344e-05   [ 0.135,  0.368]  

=====
```

Covariance estimator: robust

Mean forecast:

2263 0.001385



## Code:

```
#example 2: arima+arch on exchange rates of us dollars
import pandas as pd
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from arch import arch_model
import matplotlib.pyplot as plt
```

```

#step 1: Load some exchange rate data
data=pd.read_csv('EURUSD_daily.csv')
data['Date'] = pd.to_datetime(data['Date'])
data = data.set_index('Date')
close=data['Close'].dropna()

#step 2: compute return
returns=np.log(close).diff().dropna()

#step 3: stationary check
adf=adfuller(returns)
print("ADF p-value:",adf[1])

#step 4: fit arima on return
model_arima=ARIMA(returns,order=(1,0,1))
res_arima=model_arima.fit()
print(res_arima.summary())

#step 5:ARCH(1) on residuals
am=arch_model(res_arima.resid,vol='ARCH',p=1,o=0,q=0,dist='t')
#use t distribution
res_arch=am.fit(update_freq=5)
print(res_arch.summary())
#Step6: forecast
steps=10
mean_fore=res_arima.get_forecast(steps=steps)
vol_fore=res_arch.forecast(horizon=steps)
print("Forecasted return:\n",mean_fore.predicted_mean.tail(steps))
print("Forecasted variance:\n",vol_fore.variance.tail(steps))

```

## Output:

The screenshot shows a Google Colab notebook titled "time series.ipynb". The code cell above displays the results of an ADF test and an ARIMA model fit. The output window shows the SARIMAX Results table, which includes parameters like Dep. Variable, Model, Date, Time, Sample, and Covariance Type. It also lists the coefficients (coef), standard errors (std err), z-scores (z), and p-values (P>|z|) for each term. Below the table, diagnostic statistics like Ljung-Box (Q), Jarque-Bera (JB), and Heteroskedasticity (H) are provided. The bottom of the output window shows warning messages related to covariance matrix calculation and optimization termination.

```

ADF p-value: 1.5520618526580473e-11
SARIMAX Results
-----
Dep. Variable: Close No. Observations: 22
Model: ARIMA(1, 0, 1) Log Likelihood: 93.561
Date: Mon, 15 Sep 2025 AIC: -179.122
Time: 11:51:13 BIC: -174.758
Sample: 0 HQIC: -178.094
- 22
Covariance Type: opg
-----
            coef  std err      z   P>|z|    [0.025    0.975]
-----
const     -0.0002    0.000   -0.557    0.577   -0.001    0.000
ar.L1     -0.2330    0.410   -0.568    0.570   -1.037    0.571
ma.L1     -0.4439    0.420   -1.057    0.290   -1.267    0.379
sigma2    1.157e-05  5.38e-06  2.149    0.032   1.02e-06  2.21e-05
-----
Ljung-Box (11) (Q): 0.10 Jarque-Bera (JB): 0.58
Prob(Q): 0.76 Prob(JB): 0.75
Heteroskedasticity (H): 0.66 Skew: 0.31
Prob(H) (two-sided): 0.59 Kurtosis: 2.49
-----
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
Iteration: 5, Func. Count: 54, Neg. LLF: 249.42119615184248
Optimization terminated successfully (Exit mode 0)
Current function value: -93.1432694090515
Iterations: 11
Function evaluations: 76
Gradient evaluations: 7

```

```

Time: 11:51:13 Df Model: 1
-----
          coef std err   t   P>|t|  95.0% Conf. Int.
mu      3.0821e-05 1.771e-03 1.741e-02    0.986 [-3.448e-03, 3.501e-03]
----- Mean Model
----- Variability Model
-----          coef std err   t   P>|t|  95.0% Conf. Int.
omega  1.0500e-05 1.476e-08 711.480     0.000 [1.047e-05, 1.053e-05]
alpha[1] 0.1000       0.762   0.131     0.896  [-1.393, 1.593]
----- Distribution
-----          coef std err   t   P>|t|  95.0% Conf. Int.
nu      12.0000      5.559   2.159  3.088e-02 [ 1.104, 22.896]
-----
Covariance estimator: robust
Forecasted return:
22 -0.0002857
23  0.000426
24 -0.0003059
25 -0.000161
26 -0.000202
27 -0.000192
28 -0.000195
29 -0.000194
30 -0.000194
31 -0.000194
Name: predicted_mean, dtype: float64
Forecasted variance:

```

## Code:

```

#program2 innovation from arma(2,1) (real GDP)
from statsmodels.datasets import macrodata
from statsmodels.tsa.arima.model import ARIMA

#load gdp data
macro=macrodata.load_pandas().data
gdp=macro['realgdp']

#fit arma(2,1)
model=ARIMA(gdp,order=(2,0,1)).fit()
innovations=model.resid

print("\n==== Program 2 Output ====")
print(model.summary())
print("\nFirst 10 innovations:\n",innovations.head(10))

```

## Output:

```

SARIMAX Results
-----
Dep. Variable: realgdp No. Observations: 203
Model: ARIMA(2, 0, 1) Log Likelihood: -1105.633
Date: Sat, 20 Sep 2020 AIC: 2211.266
Ljung-Box (Q): 10132.26 BIC: 2243.833
Sample: 0 HQIC: 2233.969
Length: 203 Covariance Type: opg
-----
          coef std err   z   P>|z|  [0.025  0.975]
const  7221.2123 4111.469   1.75   0.085  -9.014  1.54e+04
ar.L1   1.9395  0.035  55.487   0.000   1.871   2.008
ar.L2  -0.9397  0.035  -26.691   0.000   -1.009  -0.871
ma.L1  -0.6039  0.067  -8.998   0.000   -0.735  -0.472
sig2_res 3082.065 257.614  12.065   0.000  2567.866  3597.695
-----
Ljung-Box (L1) (Q): 0.15 Jarque-Bera (JB): 8.77
Prob(Q): 0.69 Prob(JB): 0.01
Heteroskedasticity (H): 2.46 Prob(H): 0.09
Prob(H) (two-sided): 0.00 Kurtosis: 4.00
-----
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
First 10 Innovations:
0 610.863279
1 -67.769750
2 -47.691538
3 -12.489874
4 11.734000
5 10.821710
6 13.271300
7 10.821710
8 13.271300
9 11.734000

```

## Code:

```
#Program 3: innovations from ARIMA(1,1,1) (CPI inflation)

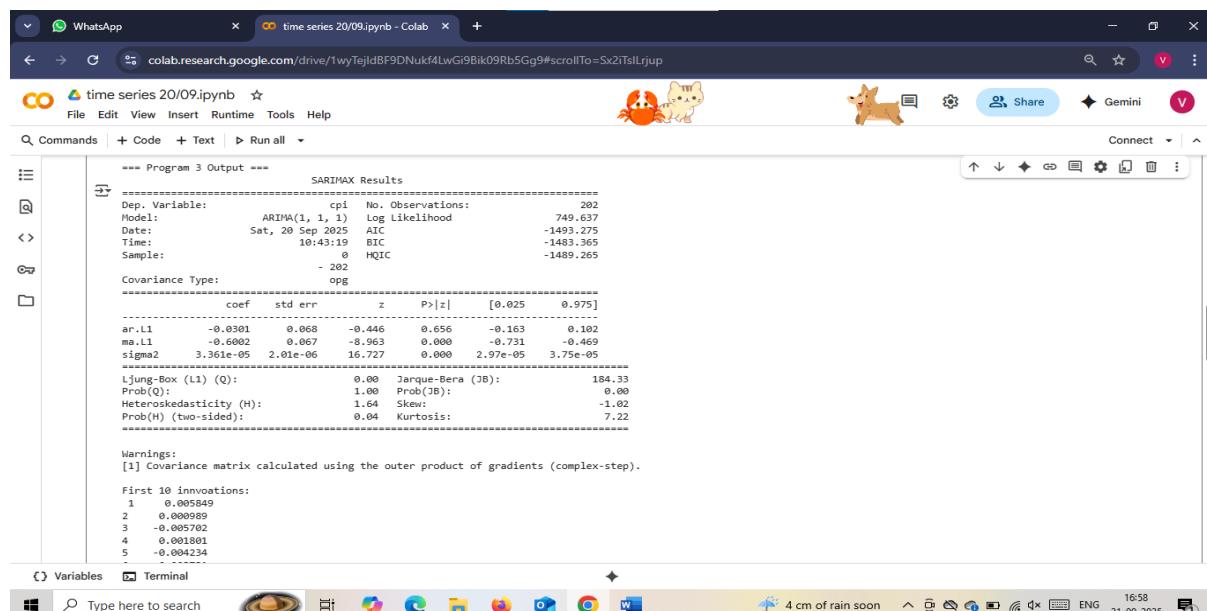
import numpy as np
#compute quarterly inflation
cpi=macro['cpi']
inflation=np.log(cpi).diff().dropna()

model=ARIMA(inflation,order=(1,1,1)).fit()
innovations=model.resid

print("\n\n==== Program 3 Output ===")
print(model.summary())

print("\nFirst 10 innovations:\n",innovations.head(10))
```

## Output:



```
*** Program 3 Output ***
SARIMAX Results
=====
Dep. Variable:          cpi    No. Observations:             202
Model:                 ARIMA(1, 1, 1)    Log Likelihood:         749.637
Date:         Sat, 20 Sep 2025   AIC:                  -1493.275
Time:            10:43:19     BIC:                  -1485.365
Sample:               0 - 202   HQIC:                  -1489.265
Covariance Type:      opg

coef    std err        z      P>|z|      [0.025      0.975]
-----+
ar.L1   -0.0301    0.068   -0.446     0.656     -0.163     0.102
ma.L1   -0.6902    0.067   -8.963     0.000     -0.731     -0.469
sigma2  3.361e-05  2.01e-06  16.727     0.000    2.97e-05   3.75e-05

Ljung-Box (L1) (Q):      0.00    Jarque-Bera (JB):       184.33
Prob(Q):                   1.00    Prob(JB):                   0.00
Heteroskedasticity (H):  1.64    Skew:                   -1.02
Prob(H) (two-sided):     0.04    Kurtosis:                  7.22

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

First 10 innovations:
1  0.0005849
2  0.000589
3  -0.0005702
4  0.001801
5  -0.004234
```

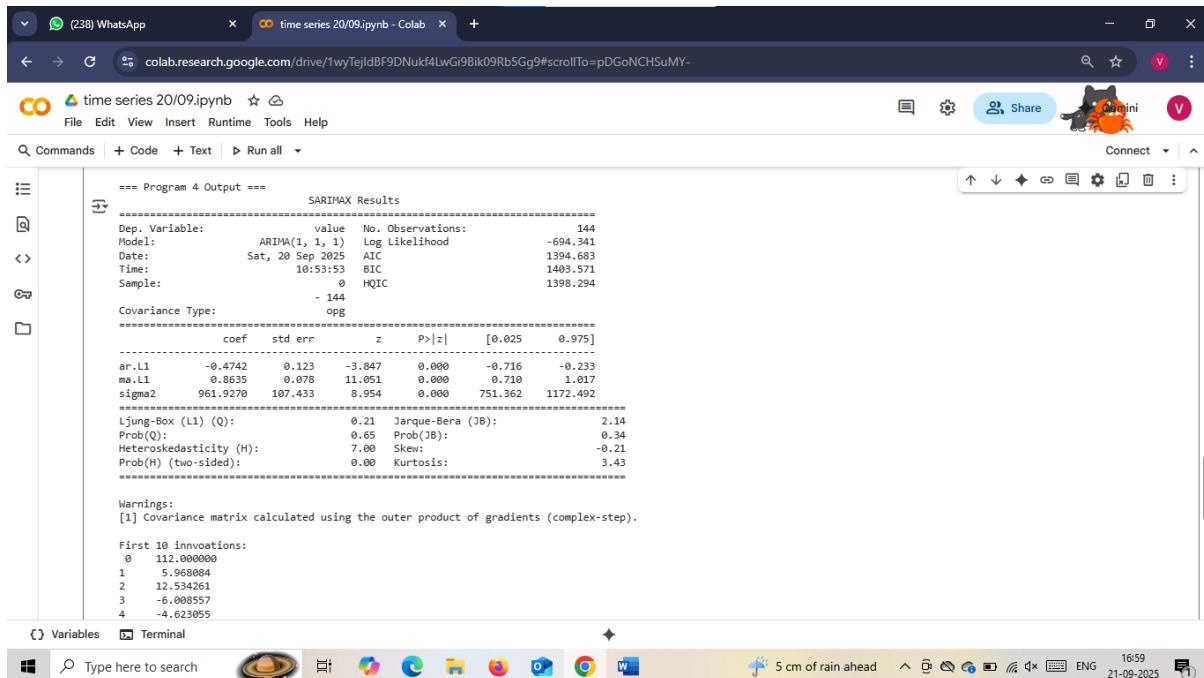
## Code:

```
from statsmodels.datasets import get_rdataset
from statsmodels.tsa.arima.model import ARIMA
#load airline data
airline_data = get_rdataset("AirPassengers").data['value']
#fit arima(1,1,1)
model = ARIMA(airline_data, order=(1, 1, 1)).fit()
innovations=model.resid

print("\n\n==== Program 4 Output ===")
print(model.summary())

print("\nFirst 10 innovations:\n",innovations.head(10))
```

## Output:



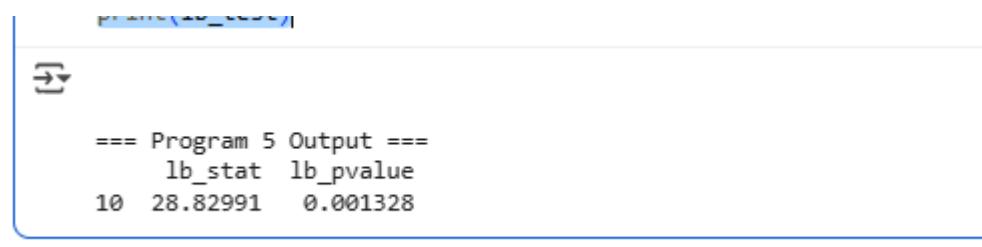
The screenshot shows a Google Colab notebook titled "time series 20/09.ipynb". The code cell displays the output of a SARIMAX model fit. The output includes the model summary, parameter estimates, and diagnostic statistics.

```
== Program 4 Output ==
SARIMAX Results
=====
Dep. Variable: value No. Observations: 144
Model: ARIMA(1, 1, 1) Log Likelihood: -694.341
Date: Sat, 20 Sep 2025 AIC: 1394.683
Time: 10:53:53 BIC: 1403.571
Sample: 0 HQIC: 1398.294
Covariance Type: opg
=====
            coef    std err      z   P>|z|      [0.025      0.975]
-----
ar.L1     -0.4742   0.123   -3.847   0.000    -0.716    -0.233
ma.L1      0.8635   0.078   11.051   0.000     0.710     1.017
sigma2    961.9270  107.433   8.954   0.000    751.362   1172.492
=====
Ljung-Box (L1) (Q): 0.21 Jarque-Bera (JB): 2.14
Prob(Q): 0.65 Prob(JB): 0.34
Heteroskedasticity (H): 7.00 Skew: -0.21
Prob(H) (two-sided): 0.00 Kurtosis: 3.43
=====
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
First 10 innovations:
0 112.000000
1 5.968084
2 12.534261
3 -6.008557
4 -4.623055
```

## Code:

```
#program 5 :Ljung box test on innovations
from statsmodels.stats.diagnostic import acorr_ljungbox
lb_test = acorr_ljungbox(innovations, lags=[10], return_df=True)
print("\n\n==== Program 5 Output ===")
print(lb_test)
```

## Output:



The screenshot shows the output of the Ljung Box test on innovations. The output displays the test statistic and p-value for lags 10.

```
==== Program 5 Output ===
lb_stat lb_pvalue
10 28.82991 0.001328
```